

A Survey of On-Policy Distillation for Large Language Models

Mingyang Song & Mao Zheng

Large Language Model Department

Tencent, China

nickmysong@tencent.com

 [nick7nlp/Awesome-LLM-On-Policy-Distillation](https://github.com/nick7nlp/Awesome-LLM-On-Policy-Distillation)

Abstract

As Large Language Models (LLMs) grow in both capability and cost, knowledge distillation has become a standard technique for transferring frontier capabilities into smaller, deployable students. The prevailing recipe, off-policy imitation of teacher-generated text, carries a structural weakness: the student trains on flawless teacher prefixes but must generate from its own imperfect outputs at inference time, and the resulting exposure bias compounds along the trajectory. On-Policy Distillation (OPD) addresses this mismatch by having the student generate its own sequences and the teacher provide feedback on those sequences, turning distillation from single-pass imitation into an iterative correction process. The approach has moved rapidly from research into production, yet the literature remains scattered across the knowledge distillation, reinforcement learning, and imitation learning communities without a unified treatment. This survey provides such a treatment. We formalize OPD as f -divergence minimization over student-sampled trajectories, organize the field along three design axes (objective functions, signal sources, and training dynamics) and consolidate success conditions, recurring failure modes, and the formal connection between OPD and KL-constrained reinforcement learning. We close with open problems including distillation scaling laws, uncertainty-aware teacher feedback, and agent-level distillation.

1 Introduction

The capability of Large Language Models (LLMs) has advanced rapidly over the past two years, with recent systems reaching strong performance on reasoning, code generation, and multilingual benchmarks (Yang et al., 2025a; DeepSeek-AI et al., 2025; Gemma Team et al., 2024). This progress has come at a training cost that only a few organizations can sustain. Transferring the resulting capabilities from large frontier models into smaller deployable students has therefore become a core step in the modern LLM pipeline. Knowledge distillation, first formalized by Hinton et al. (2015) as a way for a student to inherit the soft output structure of a teacher, has broadened in scope over this period. What began as a model compression technique for transferring knowledge within the same architecture now serves as a general mechanism for moving capabilities across model scales and architectural families. The release of DeepSeek-R1 (DeepSeek-AI et al., 2025) made this shift concrete, distilling a 671B mixture-of-experts teacher into dense students spanning 1.5B to 70B parameters while preserving long chain-of-thought reasoning largely intact.

The dominant recipe behind these successes, however, has a known structural weakness. In almost all large-scale pipelines, distillation proceeds *off-policy*. The student is trained to match the teacher’s next-token distribution over a fixed corpus, typically the original pre-training data or traces generated in advance by the teacher, and every gradient step conditions on flawless teacher prefixes. At inference, this assumption no longer holds. The student generates autoregressively from its own partial outputs, and any divergence between the training states and the states it visits at deployment translates into compounding error along the trajectory. This distribution mismatch corresponds to the *compounding*

error identified in interactive imitation learning (Ross et al., 2011), where autoregressive generation amplifies the distribution mismatch quadratically in sequence length, producing an expected discrepancy of $O(\epsilon T^2)$ over a horizon T . The consequences are most pronounced on reasoning tasks, where a single early misstep can derail an entire proof or program. Gudibande et al. (2023) report that off-policy imitation of proprietary LLMs often yields students that reproduce surface style without acquiring the underlying reasoning competence.

Addressing this mismatch requires changing *where* the training data comes from rather than *what* is being matched. **On-Policy Distillation (OPD)** addresses this by shifting the sampling distribution from a static corpus to the student’s own evolving policy. The student proposes trajectories, populates the states it will visit at deployment, and receives teacher feedback on those states rather than on idealized expert prefixes. Depending on the available teacher access, this feedback ranges from full token-level distributions in white-box settings (Agarwal et al., 2024; Gu et al., 2024), through scalar rewards or pairwise preferences when only black-box APIs are available (Ye et al., 2025; Jiang et al., 2023), to contrastive signals the model derives from its own outputs in teacher-free settings (Zhao et al., 2026b). The theoretical payoff mirrors the classical DAgger result (Ross et al., 2011), where querying an expert on the learner’s own visited states reduces the $O(\epsilon T^2)$ compounding of off-policy imitation to $O(\epsilon T)$. Translated into autoregressive language modeling, this turns distillation from a single-pass matching problem into an iterative optimization loop in which the training distribution co-evolves with the model, allowing distributional gaps between training and deployment to be corrected during training rather than deferred to inference.

The body of work that has emerged from this reformulation can be organized around a small number of recurring design choices. Early methods focused on *how* to sample and *which* divergence to apply on student trajectories (Agarwal et al., 2024; Gu et al., 2024; Ko et al., 2024), establishing that an appropriate divergence depends on where in the sequence one is training and what kind of supervision the teacher can supply. A parallel line of work revealed that the resulting objective is formally equivalent to a KL-constrained form of reinforcement learning (Yang et al., 2026e), bridging two previously separate research communities. A third branch showed that the “teacher” need not be external at all, with methods such as OPSD (Zhao et al., 2026b) and SD-ZERO (He et al., 2026) showing that a model can bootstrap improvement from its own generations by exploiting asymmetries in its output distribution across contexts. These ideas have since moved from papers into large-scale training pipelines, with Qwen3 (Yang et al., 2025a), DeepSeek-V4 (DeepSeek-AI, 2026), Gemma 2 (Gemma Team et al., 2024), and MiMo-V2-Flash (Xiaomi LLM-Core Team et al., 2026) all adopting OPD as a core training ingredient. The adoption cuts across architectural lines, covering both mixture-of-experts and dense systems and spanning proprietary and open-weight release strategies, because all of these systems face the same structural problem once reasoning chains grow long enough for compounding error to dominate. DeepSeek-V4 went further than most, replacing its mixed RL stage with pure multi-teacher OPD for model consolidation.

GKD (Agarwal et al., 2024), alongside the concurrent MiniLLM (Gu et al., 2024), brought on-policy distillation into the LLM era in mid-2023, and in roughly three years the literature has expanded to over two hundred papers spanning divergence design, reward-guided optimization, self-play, multi-teacher debate, agentic trajectory distillation, and cross-modal transfer. The research focus has shifted accordingly. Where early work asked whether on-policy training helps at all, the current frontier asks how to make it efficient enough for trillion-parameter teacher inference, how to extend it beyond single-turn generation into multi-step agent trajectories, and how to unify it with the reinforcement learning pipelines that increasingly share its computational infrastructure. This trajectory parallels a broader shift in LLM development. As models become more capable and reasoning chains grow longer, the gap between off-policy training states and deployment-time states widens, which makes on-policy correction increasingly attractive for the next generation of reasoning-capable systems (Fu et al., 2026b; DeepSeek-AI, 2026).

Despite this rapid adoption, the research literature on OPD has yet to catch up with the practice. Methods approach the same problem through the separate lenses of knowledge

distillation, RLHF, and interactive imitation learning, each carrying different notations, benchmarks, and failure taxonomies from its parent community. Existing surveys of LLM distillation (Xu et al., 2024) generally retain the classical compression framing, treating off-policy and on-policy methods as interchangeable variants instead of as regimes with materially different theoretical properties and failure modes. No current treatment we are aware of offers a unified mathematical account of the on-policy dynamics that connect these threads, and a systematic comparison across the white-box, black-box, and teacher-free instantiations that researchers and engineers must now weigh against one another in practice remains missing.

This survey addresses that gap. To our knowledge, it is the first to treat on-policy distillation as a distinct paradigm rather than a variant of classical knowledge distillation, and it develops a unified account through four contributions that also structure the rest of the paper.

- **A unified theoretical framework.** We recast the transition from off-policy to on-policy distillation as a sequential decision-making problem and show that the core OPD algorithms are instances of f -divergence minimization over student-sampled trajectories (Section 2). This framework supplies a common analytical vocabulary for methods previously studied in isolation and clarifies their relationships in terms of divergence choice, argument ordering, and sampling mixture.
- **A design-centric taxonomy.** Rather than grouping methods by surface similarity, we organize them along three core design axes (Section 3), namely objective function design (Section 4), signal source architecture (Section 5), and training dynamics optimization (Section 6). The resulting one-method-one-category classification clarifies the relationships among methods in terms of the specific design choices each paper optimizes.
- **A theoretical account of failure.** We consolidate the empirical record into a coherent view of when and why OPD breaks down (Section 7), covering the flawed prefix trap, self-play saturation, diversity collapse, the calibration-capability gap, and emerging multi-turn agentic failure modes, and we describe formal conditions under which OPD improves over off-policy alternatives.
- **Cross-regime comparative analysis.** We survey the white-box, black-box, and self-distillation regimes within this common framework, summarize them in per-section comparison tables (Tables 1, 3–9, and 2), and document method selection factors (Section 3.3) that map deployment constraints to applicable methods and compute budgets.

The structure of the paper follows the same design-axis logic. Section 2 develops the mathematical preliminaries, formalizes autoregressive generation as a sequential decision problem, and derives the exposure bias that motivates on-policy training. Section 3 presents the method landscape and selection considerations. Sections 4–6 then traverse the three design axes from “what to optimize” through “where the signal comes from” to “how to stabilize training in practice,” with each section building on the choices surfaced in the previous one. Section 7 consolidates the theoretical perspectives, showing that the patterns observed across Sections 4–6 follow from a small number of unifying principles. Section 8 discusses industrial deployments, system-level optimization, and emerging domains, and Section 9 lays out the open problems that emerge from this synthesis.

Our scope is centered on the on-policy regime. We cover methods in which the student generates its own training data during distillation, and we include work from adjacent fields such as imitation learning, online RL, and preference optimization whenever it operates on student-generated data under teacher or verifier supervision. We exclude generic off-policy KD, compression techniques orthogonal to the training regime such as pruning and quantization, and inference-time methods that leave model weights unchanged.

2 Background and Unified Math

To understand why on-policy training is often preferable, we formalize the compounding geometry of autoregressive generation. We begin with classical KD and its limitations, then establish the unified f -divergence framework that subsumes modern OPD objectives.

Defining “on-policy.” A distillation method is *on-policy* if the training data for the student is sampled from the student’s own current policy p_θ at training time, rather than from a fixed external corpus $\mathcal{D}$ or from the teacher’s generation distribution p_T . Formally, on-policy training optimizes:

$$\min_{\theta} \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y \sim p_\theta(\cdot|x)} [\mathcal{L}(y, x; \theta, T)]$$

where $\mathcal{L}$ is the distillation loss (divergence, reward, or hybrid). The key distinction is that the *outer* expectation is over the student’s own generations, not a static dataset. This creates a non-stationary optimization landscape because as θ updates, the data distribution p_θ shifts, requiring fresh rollouts at each training step. The computational cost of this fresh generation is a central systems-level challenge of OPD (Section 6.3).

Notation. Throughout this survey, we use lowercase p for token-level conditional distributions ($p_T(\cdot|x, y_{<t})$ for teacher, $p_\theta(\cdot|x, y_{<t})$ for student) and uppercase P for sequence-level distributions ($P_T(y|x)$, $P_\theta(y|x)$). $\mathcal{D}$ denotes the training dataset, p_{data} the data distribution, and $|V|$ the vocabulary size. D_{KL} and D_f denote KL divergence and general f -divergence respectively. In the RL formalization, π denotes policies (e.g., π_{mix} for a mixture policy, π_{teach} for a sequence-level teacher policy in DPO-style derivations, π_{ref} for a reference policy) and $R(x, y)$ the outcome reward. When discussing individual methods, we follow each paper’s conventions where clarity demands.

2.1 Classical Knowledge Distillation

The original KD formulation (Hinton et al., 2015) trains the student to match the teacher’s temperature-softened output distribution:

$$p_T^{(\tau)}(y|x) = \frac{\exp(z_y/\tau)}{\sum_j \exp(z_j/\tau)}$$

where z_y are the teacher’s logits and $\tau > 1$ is a temperature parameter. Moderate temperatures ($\tau > 1$) expose the teacher’s “dark knowledge” about inter-class similarities that hard labels hide. In the limit $\tau \rightarrow \infty$, the distribution approaches uniform and this structure is lost. The gradient of the distillation loss with respect to the student’s logits z_i^S shows why temperature matters:

$$\frac{\partial \mathcal{L}_{\text{KD}}}{\partial z_i^S} = \frac{1}{\tau} (p_i^S - p_i^T)$$

In the high-temperature limit ($\tau \rightarrow \infty$), a Taylor expansion $\exp(z_i/\tau) \approx 1 + z_i/\tau$ together with the zero-mean assumption $\sum_i z_i^S = \sum_i z_i^T = 0$ reduces the gradient to $\frac{1}{|V|\tau^2} (z_i^S - z_i^T)$ (Hinton et al., 2015), so classical KD in this regime is equivalent to minimizing the mean squared error between raw logits. This encourages the student to replicate the teacher’s *entire* logit structure, transferring both primary predictive modes and nuanced sub-optimal probability masses (the “dark knowledge”). The $1/\tau^2$ factor explains why the soft-target term is commonly upweighted by τ^2 relative to the hard-label cross-entropy in the combined KD objective. In practice, LLM distillation operates at $\tau = 1$ or low-to-moderate temperatures (Zhong et al., 2024) because the large vocabulary ($|V| > 30,000$) already produces richly structured non-peak probabilities without explicit softening, and higher temperatures may amplify noise in the teacher’s poorly calibrated tail distribution.

For autoregressive LLMs, the token-level distillation loss factors as:

$$\mathcal{L}_{\text{Token-KD}} = \mathbb{E}_{x, y \sim \mathcal{D}} \left[\sum_{t=1}^{|y|} D_{\text{KL}}(p_T(\cdot|x, y_{<t}) \parallel p_\theta(\cdot|x, y_{<t})) \right] \quad (1)$$

This is computationally tractable because it relies on static, pre-computed prefixes from the dataset, but it optimizes next-token accuracy while assuming an error-free history $y_{<t}$. Kim & Rush (2016) extended this to Sequence-Level Knowledge Distillation (Seq-KD), which instead minimizes the global KL divergence over the entire sequence space:

$$\mathcal{L}_{\text{Seq-KD}} = D_{\text{KL}}(P_T(y|x) \parallel P_\theta(y|x)) = \sum_{y \in \mathcal{Y}} P_T(y|x) \log \frac{P_T(y|x)}{P_\theta(y|x)}$$

Since $|\mathcal{Y}| = |V|^T$ grows exponentially with length, exact computation is intractable. Seq-KD approximates $P_T(y|x)$ with a Dirac delta at the beam-search output $\hat{y} = \arg \max_y P_T(y|x)$, reducing to standard negative log-likelihood on teacher-generated sequences:

$$\mathcal{L}_{\text{Seq-KD}} \approx -\log P_\theta(\hat{y}|x)$$

This approximation works well when the teacher is highly confident but discards much of the information about the teacher’s uncertainty and alternative modes. The practical consequence is twofold. Seq-KD collapses the teacher’s distributional richness into a single point estimate, and it still trains on a *static* dataset disconnected from the student’s own generation behavior. These two limitations (loss of distributional signal and off-policy training) motivate the richer objectives below.

Subsequent work (Zhong et al., 2024) revisited these foundations for modern LLMs and suggested that classical assumptions (small teacher-student gap, shared vocabulary, similar architectural inductive biases) frequently fail when distilling across heterogeneous model families. The gap between classical assumptions and modern reality motivates the more general f -divergence framework below.

From classical KD to on-policy. The historical progression from Hinton’s classical KD to modern OPD amounts to a series of relaxations of four classical assumptions, namely (1) shared vocabulary, (2) i.i.d. data, (3) static teacher, and (4) off-policy training data. Seq-KD (Kim & Rush, 2016) relaxes assumption (2) by operating at the sequence level but retains the static data assumption. GKD (Agarwal et al., 2024) relaxes assumption (4) by training on student-generated sequences. DSKD (Zhang et al., 2025b) relaxes assumption (1) through dual-space alignment. G-OPD (Yang et al., 2026e) relaxes assumptions (2)–(4) simultaneously while adding RL-augmented reward extrapolation that pushes the student beyond the teacher’s frontier. Each relaxation targets a specific failure mode that intensifies with model scale. Vocabulary mismatch is often irrelevant for same-family distillation but severe for cross-family transfer (relaxation 1), while exposure bias is negligible for short sequences but highly problematic for multi-step reasoning (relaxation 4). This progression is not merely historical but informative, since the relevant assumptions in a given deployment scenario determine which level of relaxation is necessary, and unnecessary relaxation introduces avoidable complexity.

2.2 Off-Policy Exposure Bias

The token-level loss in Eq. (1) conditions on ground-truth prefixes $y_{<t}$ drawn from a static dataset. During inference, the student generates its own prefixes, creating a distribution mismatch between training and deployment that compounds over sequence length. We now formalize why this mismatch is catastrophic.

Standard knowledge distillation minimizes the KL divergence over states drawn from the dataset distribution $d_{\mathcal{D}}(s)$:

$$\mathcal{L}_{\text{Off-Policy}} = \mathbb{E}_{x, y \sim p_{\text{data}}} \left[\sum_{t=1}^{|y|} D_{\text{KL}}(p_T(\cdot|x, y_{<t}) \parallel p_\theta(\cdot|x, y_{<t})) \right]$$

At inference, the student acts according to its own policy, inducing a different state visitation distribution $d_{p_\theta}(s)$. By the DAgger theorem (Ross et al., 2011), if a policy mimics an expert with per-step error ϵ under the training distribution, the expected total discrepancy over a trajectory of length T under the learner’s own state visitation scales as $O(\epsilon T^2)$.

This quadratic compounding has severe practical consequences for reasoning tasks. Even under the simplest independent-error assumption, a mathematical proof requiring 10 reasoning steps with per-step accuracy 95% yields only $(0.95)^{10} \approx 60\%$ trajectory-level correctness. The DAgger bound predicts something strictly worse. Because off-policy training rarely exposes the student to its own error states, each mistake pushes the model into regions where its next-step accuracy degrades further (distributional shift), producing $O(\epsilon T^2)$ total error rather than the $O(\epsilon T)$ of independent compounding. On-policy training offers a way to mitigate this self-reinforcing dynamic by exposing the student to its own error states

during training, so that it can learn recovery strategies that keep distributional drift from amplifying local mistakes into trajectory-level collapse.

Remark on the DAgger bound in LLMs. While the theoretical reduction from $O(\epsilon T^2)$ to $O(\epsilon T)$ is appealing, applying the DAgger bound to LLMs demands careful qualification. The original theorem assumes an *interactive expert* that emits optimal actions in any state. In white-box OPD, the “expert” emits a next-token distribution $p_T(y_t | \hat{y}_{<t})$ conditioned on the student’s prefix $\hat{y}_{<t}$. If the student hallucinates a severely out-of-distribution prefix, the teacher’s conditional distribution may itself become poorly calibrated, because the teacher was rarely trained on such inputs. Forcing the student to match this noisy distribution violates the core assumption of interactive imitation learning, potentially destabilizing training rather than recovering the $O(\epsilon T)$ bound. Jeong (2026) offer direct empirical evidence that naive OPD without stable teacher dynamics suffers catastrophic instability. Unconditional token-level matching with periodic hard-copy teacher resets leads to abrupt output collapse (KL divergence dropping from 2.637 to 0.343 at a single teacher-reset event), demonstrating that unconditional token-level matching alone is insufficient for stable OPD. This observation motivates the adaptive trust mechanisms surveyed in Section 6.1 and the stable teacher dynamics discussed in Section 7.2.

2.3 The Unified f -Divergence Framework

OPD resolves exposure bias by altering the training expectation to sample from the student’s own rollouts, or a mixture policy π_{mix} . The generalized on-policy objective decouples the sampling trajectory from the local matching metric:

$$\mathcal{L}_{\text{OPD}}(\theta) = \mathbb{E}_{y \sim \pi_{\text{mix}}} \left[\sum_{t=1}^{|y|} D_f(p_T(\cdot | x, y_{<t}), p_\theta(\cdot | x, y_{<t})) \right]$$

Here, D_f represents a divergence from the f -divergence family (Wen et al., 2023). Formally, given two distributions P and Q , an f -divergence is defined as:

$$D_f(P \parallel Q) = \mathbb{E}_{y \sim Q} \left[f\left(\frac{P(y)}{Q(y)}\right) \right]$$

where $f : (0, \infty) \rightarrow \mathbb{R}$ is a convex generator with $f(1) = 0$. The choice of f determines the implicit weighting of likelihood ratios $p_T(y)/p_\theta(y)$:

- **Forward KL** ($f(u) = u \log u$): Mode-covering (zero-avoiding). The student is encouraged to allocate probability mass wherever the teacher does, often bridging distinct teacher modes and hallucinating in the inter-mode space.
- **Reverse KL** ($f(u) = -\log u$): Mode-seeking (zero-forcing). The student concentrates its mass on a subset of the teacher’s modes, avoiding regions where the teacher assigns near-zero probability but potentially ignoring secondary acceptable outputs.
- **JSD** ($f(u) = u \log u - (u + 1) \log \frac{u+1}{2}$): Symmetric, bounded, and smoothly interpolates between mode-covering and mode-seeking behavior.
- **α -divergence**: Parameterized family that continuously interpolates between Forward KL ($\alpha \rightarrow 1$) and Reverse KL ($\alpha \rightarrow 0$), allowing fine-grained control over the mode-seeking/covering tradeoff.

Concretely, these divergences each suit different task geometries. For tasks with a *unique* correct answer (mathematical proofs, code correctness), Reverse KL’s mode-seeking behavior concentrates the student on an appropriate solution path. For tasks with *many* acceptable outputs (creative writing, open-ended QA), Forward KL’s mode-covering behavior preserves output diversity. JSD offers a compromise but lacks theoretical optimality properties for either extreme. All these divergences are computable as expectations under the student’s own policy (since $D_f(P_T \parallel P_\theta) = \mathbb{E}_{y \sim p_\theta} [f(p_T(y)/p_\theta(y))]$), making them directly amenable to on-policy optimization. The sampling distribution p_θ is the policy being optimized, so gradients flow through the reparameterized samples without importance weights or off-policy corrections that would add additional variance.

This framework subsumes many modern OPD objectives. The choice of f determines the geometric behavior, while the choice of π_{mix} controls the degree of on-policy exploration. We illustrate by mapping the three foundational methods into this space.

GKD (Agarwal et al., 2024) defines $\pi_{\text{mix}} = \lambda p_\theta + (1 - \lambda)p_{\text{data}}$ as an explicit interpolation, and is agnostic to D_f , empirically testing Forward KL, Reverse KL, and JSD. Setting $\lambda \rightarrow 1$ makes GKD purely on-policy, while $\lambda = 0$ reduces to standard off-policy KD. The λ parameter offers a smooth dial between the extremes, and moderate values ($\lambda \approx 0.5$) tend to combine the benefits of on-policy exposure with the stability of grounded, dataset-derived prefixes. GKD’s experiments report that on-policy sampling ($\lambda = 1$) outperforms off-policy across the divergence choices they tested. JSD performs best on translation tasks (e.g., WMT), where the output space has moderate diversity (neither a single correct answer nor fully open-ended generation) and the symmetric nature of JSD naturally balances mode-covering and mode-seeking pressures appropriate for this intermediate regime.

Empirically, all three divergences yield comparable results on summarization and instruction-following, supporting the exposure bias hypothesis. This suggests that the sampling policy (on-policy vs. off-policy) matters more than the specific divergence choice when task geometry does not strongly favor one extreme.

MiniLLM (Gu et al., 2024) selects Reverse KL as $D_f = D_{\text{KL}}(p_\theta \parallel p_T)$ and employs $\pi_{\text{mix}} = (1 - \alpha)p_\theta + \alpha p_T$ with $\alpha = 0.2$ for stability. Because Reverse KL places the student in both the expectation and the log-ratio, MiniLLM reformulates optimization via REINFORCE, treating $\log(p_T/p_\theta)$ as a per-step reward.

DistiLLM (Ko et al., 2024) computes its token-level loss on student-generated prefixes but uses an adaptive scheduler that conservatively uses student-generated outputs based on the validation loss. It also maintains a replay buffer to improve sample efficiency (an “adaptive off-policy” strategy). Its core contribution is a skew KLD loss using a mixture distribution $\tilde{p} = \alpha p + (1 - \alpha)q$, computing $D_{\text{KL}}(p \parallel \tilde{p})$ (SKL) or $D_{\text{KL}}(q \parallel \tilde{p})$ (SRKL). This avoids zero-division instability when $p_\theta(y) \approx 0$ or $p_T(y) \approx 0$, because $\tilde{p}(y)$ is bounded away from zero, improving tractability and stability without policy gradients.

Comparative synthesis. These three methods expose a core tradeoff triangle in OPD design. GKD prioritizes *generality* (divergence-agnostic) and *simplicity* (no REINFORCE variance) at the cost of limited theoretical guidance on divergence selection. MiniLLM commits to Reverse KL for its mode-seeking precision but inherits REINFORCE’s high variance, requiring reward baselines, length penalties, and careful α tuning. DistiLLM avoids both the divergence selection problem and MiniLLM’s variance problem by engineering a mixture target that is numerically stable by construction, but at the cost of adding a replay buffer (mixing on-policy and off-policy data) and an adaptive scheduler that introduces additional hyperparameters. The arc from GKD through MiniLLM to DistiLLM thus traces a shift from algorithmic simplicity toward computational and engineering sophistication, with later methods often addressing specific limitations of earlier approaches.

Viewed through this decomposition, modern OPD is a systematic exploration of three interacting design choices, trajectory sampling (π_{mix}), divergence generator (f), and argument ordering within the chosen divergence.

2.4 Distillation Scaling Laws

The unified f -divergence framework characterizes *what* to optimize, but leaves open *how* performance scales with computational investment. Unlike pre-training, where Hoffmann et al. (2022) established compute-optimal scaling laws (the Chinchilla laws), distillation creates a qualitatively different resource allocation problem. The budget is typically split between teacher inference, student rollouts, and gradient updates, with non-trivial interactions between these axes.

Busbridge et al. (2025) initiated the study of distillation-specific scaling and found that the teacher-student relationship obeys qualitatively different dynamics from pre-training alone. They find that a *capacity gap* emerges when the teacher becomes too powerful relative to the

student. The student’s loss follows a power law in teacher cross-entropy that transitions between two regimes. Initially, stronger teachers improve the student (denser gradient signal from better-calibrated next-token probabilities), but beyond a critical teacher capacity the student’s ability to absorb the signal saturates. This capacity-gap phenomenon connects directly to the divergence choice in Section 2.3, where Forward KL’s mode-covering pressure exacerbates the gap (the student must spread mass across all teacher modes, including those beyond its capacity), while Reverse KL’s mode-seeking behavior partially mitigates it (the student selectively matches its highest-confidence subset).

No comprehensive framework yet predicts how distillation quality scales jointly with teacher size N_T , student size N_S , data volume D , and on-policy rollout budget R . A natural extension of the Busbridge et al. (2025) scaling law might take the form $\text{Quality} \propto N_T^\alpha \cdot N_S^\beta \cdot D^\gamma \cdot R^\delta$, but pinning down the exponents and interaction terms for on-policy methods remains open (Section 9). The rollout budget R is unique to on-policy distillation, adding a new scaling axis absent from both pre-training and off-policy distillation. Each rollout step requires a full forward pass through the student, making compute cost scale linearly with R while the learning benefit likely exhibits diminishing returns.

3 Landscape and Method Selection

Before examining individual methods in detail, this section provides a structured overview of the field, covering the organizational taxonomy, a comparative table of representative methods, and a discussion of method selection considerations.

3.1 Method Landscape

Existing OPD methods can be organized along three axes that correspond to sequential design decisions in the training pipeline: (1) the objective function to optimize, (2) the source of the supervisory signal, and (3) the mechanism for stabilizing training dynamics. For the signal source axis, we adopt a two-level classification. The first level distinguishes *external teacher distillation*, where a separate and typically stronger model provides supervision, from *self-distillation*, where the model generates its own training signal. Within external teacher distillation, the second level separates methods by access level, with white-box methods that require the teacher’s full output distribution versus black-box methods that operate with only generated text or scalar scores. This nested structure keeps teacher identity (external versus self) and teacher access level (white-box versus black-box) on separate levels rather than treating them as peer categories, since the former is a methodological choice while the latter is a deployment constraint. Figure 1 illustrates this organization.

These three axes map onto independent terms of the unified OPD objective. The f -divergence generator f governs the *objective function* (Axis 1, Section 4), determining whether the student is mode-seeking, mode-covering, or adaptive. The identity and access level of the entity supplying $p_T(\cdot \mid x, y_{<t})$ governs the *signal source* (Axis 2, Section 5), whether an external teacher with full logit access, an external teacher with only text or score outputs, or the model’s own conditionally privileged distribution. The sampling policy and its scheduling govern the *training dynamics* (Axis 3, Section 6), including the interpolation coefficient λ , the curriculum over prompt difficulty, and token-level weighting. The first axis is largely decoupled from the other two, since a given divergence can be paired with most sampling schedules and with either teacher identity. The signal source and the objective are only partially decoupled, because black-box access removes token-level distributions and so restricts the feasible divergences to sequence-level surrogates. Treating the three as separate axes therefore organizes the design space without implying that every combination is realizable.

To illustrate how the taxonomy maps to concrete systems, GKD (Agarwal et al., 2024) selects forward KL (or JSD) as its divergence objective, relies on white-box teacher logits as its signal source, and interpolates between student-generated and ground-truth sequences (π_{mix}) to stabilize on-policy training dynamics. The DeepSeek-R1 distilled models (DeepSeek-AI et al., 2025), though off-policy (trained on static teacher traces rather than student rollouts),




Figure 1: Taxonomy of On-Policy Distillation methods organized along three design axes: (1) Objective function design (§4), (2) Signal source and teacher architecture (§5), and (3) Training dynamics and efficiency (§6). Each leaf lists a representative subset of methods, with the badge giving the total method count for that subsection (covering both the listed methods and additional papers discussed in the corresponding subsection prose). Methods are placed under their *primary contribution*, so a method that uses white-box logits but innovates on the objective (e.g., GKD, DistiLLM, EOPD, G-OPD) appears under §4 rather than §5.1.1, and the §5.1.1 leaf collects only methods whose primary contribution is the white-box signal interface itself. Methods that contribute to secondary dimensions (e.g., PAINT and PRISM contributing implicit curricula) are discussed in the relevant section prose rather than duplicated in the tree.

illustrate the same three-stage decomposition at industrial scale, adopting cross-entropy as the objective, 800K curated samples (predominantly reasoning trajectories) distilled from R1 as the signal source, and cosine learning-rate decay with difficulty-aware curricula for stability.

The three stages interact in both constraining and reinforcing ways. Some combinations are incompatible. Forward KL in its exact token-level form requires the teacher’s full output distribution, ruling out API-constrained settings where only generated text is available. Other combinations reinforce each other, as RL-augmented objectives naturally couple with external feedback sources (verifiers, reward models), while fixed divergences align with white-box logit access that permits exact gradient computation. An objective choice therefore constrains the viable signal sources and dynamics strategies.

Chronologically, the distribution of methods reflects shifting research emphasis. Early work (2023–2024) concentrated on the objective axis, debating forward KL, reverse KL, and JSD (Agarwal et al., 2024; Gu et al., 2024). By mid-2025, the focus shifted to signal architecture, in particular self-distillation methods that operate without an external teacher. The most recent work (late 2025–2026) addresses training dynamics, specifically the instabilities that on-policy sampling introduces but static-dataset distillation avoids. Industrial systems such as DeepSeek-V4 and Qwen3 increasingly combine all three axes, pairing adaptive or RL-augmented objectives with multi-teacher signals and curriculum-based stabilization (Yang et al., 2026e; 2025a).

3.2 Method Comparison Table

Classification methodology. Each method is assigned to exactly one primary category based on its *core contribution dimension*. A paper whose novelty lies in the loss function (e.g., G-OPD’s KL-constrained RL reformulation) belongs to “Objective.” A new signal architecture (e.g., DSKD’s dual-space projection) places a method under “Signal.” A training stability or efficiency contribution (e.g., TIP’s token weighting) maps to “Dynamics.” Methods that contribute to multiple dimensions are classified by their *most distinctive* contribution, reducing the categorization ambiguity that arises when a single method spans multiple dimensions (Xu et al., 2024).

Distributional observations. As of early 2026, self-distillation constitutes the largest and fastest-growing category, with the majority of entries published in 2025–2026. This distribution reflects a gradual shift from the foundational “which KL direction” question toward self-improvement and adaptive alternatives, and matches a broader trend in LLM research in which the model’s own rollouts increasingly serve as the primary data source, whether or not an external teacher is present.

Table 1: Comparison of key OPD methods categorized by their primary contribution.

Method	Category	Core Contribution	Signal Source	Loss Granularity
GKD (Agarwal et al., 2024)	Objective (Fixed)	Dagger for LLMs, interpolation	White-Box Logits	Token
DistiLLM (Ko et al., 2024)		Skewed KL mixtures	White-Box Logits	Token
MiniLLM (Gu et al., 2024)		Sequence-level Reverse KL via RL	White-Box Logits	Sequence
EOPD (Jin et al., 2026b)	Objective (Adaptive)	Entropy-aware FKL+RKL blend	White-Box Logits	Token
G-OPD (Yang et al., 2026e)	Objective (RL)	OPD equivalence to KL-RL	White-Box Logits	Sequence/Token
RLKD (Xu et al., 2025b)		KD augmented with reward model	Reward Model	Sequence
DSKD (Zhang et al., 2025b)	Signal (White-Box)	Cross-architecture alignment	White-Box Logits	Token
Lion (Jiang et al., 2023)	Signal (Black-Box)	Verbal feedback curriculum	Black-Box API	Sequence
GAD (Ye et al., 2025)		Adversarial distribution matching	Black-Box API	Sequence
OPSD (Zhao et al., 2026b)	Signal (Self)	Privileged Info (PI) conditioning	Self-Generated	Token
TIP (Xu et al., 2026c)	Dynamics (Weight)	Token importance via entropy & divergence	White-Box Logits	Token
PACED (Xu et al., 2026b)	Dynamics (Curriculum)	Beta-kernel curriculum pacing	White-Box Logits	Problem

Table 2: Representative experimental configurations of on-policy distillation methods across different application domains.

Method	Teacher	Student	Training Signal	Key Benchmarks
<i>Mathematical Reasoning</i>				
G-OPD (Yang et al., 2026e)	Qwen3-30B-A3B-Instruct	Qwen3-1.7B/4B	Dense logit + reward extrap.	AIME, HMMT, Code
OPSD (Zhao et al., 2026b)	Self (GT-conditioned)	Qwen3-1.7B/4B/8B-Inst	GT-F self-distillation	AIME, HMMT
PACED (Xu et al., 2026b)	Qwen3-14B	Qwen3-8B	Beta-kernel curriculum	MATH-500, AIME
KDRL (Xu et al., 2025a)	Skywork-OR1-Math-7B	R1-Distill-Qwen-1.5B / DeepScaleR	Joint KD + RL	MATH500, Minerva, AIME24
LUFFY (Yan et al., 2025)	DeepSeek-R1 (off-policy traces)	Qwen3-1.5B-Math-7B/1.5B	Mixed-policy GRPO	AIME, AMC, MATH-500
AOPD (Jia et al., 2026)	Qwen3-32B/8B (teacher)	Qwen3-8B/4B-Base	Asymmetric PG + FKL	AIME24/25, HMMT25
VOFD (Oh et al., 2026)	Qwen3-1.7B-Inst (Inst → Base)	Qwen3-1.7B/4B-Base, OLMo-3-7B	CV baseline + single-sample RKL	MATH-500, AIME24/25, Minerva
<i>Instruction Following & General</i>				
GKD (Agarwal et al., 2024)	TS-XL (3B)	TS-Small/Base	On-policy FKL/RKL/JSD	XSum, WMT
DistiLLM (Ko et al., 2024)	GPT-2 XL (1.5B)	GPT-2 (124M)	Skewed KL	ROUGE-L, Dolly
AlignDistil (Zhang et al., 2025a)	Self (contrastive DPO reward)	Qwen2/2.5-1.5B-Inst	Token-level RLHF distillation	AlpacaEval, MT-Bench, Arena-Hard
<i>Industrial Scale</i>				
Qwen3 (Yang et al., 2025a)	Qwen3-32B/238B-A22B	Qwen3-0.6B-14B	On-policy logit distillation	AIME, MATH, LiveCode
Gemma 2 (Gemma Team et al., 2024)	Gemma 2 27B	Gemma 2 9B / 3B	KD in pre-training	MMMLU, HellaSwag
MMo-V2-Flash (Xiaoqi LLM-Core Team et al., 2026)	Domain specialists	MMo-V2-Flash (309B/15B MoE)	Multi-teacher logit + reward	AIME, MATH, GPQA
KAT-Coder-V2 (Li et al., 2026b)	5 domain specialists	Unified agentic coder	Specialize-then-Unity OPD	SWE-Bench (79.6%)
Nemotron-Cascade 2 (Ving et al., 2025b)	Domain RL/SFT teachers	SIB MoE (5B active)	Cascade RL + domain OPD	IMO, IOI, ICPC
DeepSeek-V4 (DeepSeek-AI, 2026)	10+ domain experts (1.6T)	DeepSeek-V4-Pro (1.6T MoE)	Full-vocab multi-teacher RKL	LiveCodeBench, HMMT, IMOAnswerBench
ORBIT (Liang et al., 2026a)	Multi-mode RL experts	DeepSeek-Distill-Qwen-1.5B, Qwen3-4B, Openmath-Nemotron-7B	Multi-teacher mode fusion	AIME, GPQA, MMLU-Pro
<i>Multimodal & Domain-Specific</i>				
SCoLE (Lyu et al., 2025)	Qwen2.5-72B-Inst	Qwen2.5-7B/3B, Llama-3.1-8B	Earliest-error correction	12 Agent benchmarks
VOLD (Bousnelham et al., 2025)	Text-only teacher	VLM student	GRPO + token-level KL	MMU-Pro, MathVista, LogicVista
CORD (Hu et al., 2026)	Text-mode LALM	Audio-mode LALM	Cross-modal RKL + GRPO	Audio QA/reasoning
SKD (Yu et al., 2025c)	Gemma-7B-T, Qwen2-7B	Gemma-2B, Qwen2-0.5B	Block-verified on-policy KL	Translation, Dialogue
DP-OPD (Khadedem et al., 2026)	GPT-2 Large 774M	DistilGPT-2-82M	GKD + DP-SCD (student only)	Yelp, BigPatent
Veto (Jiang et al., 2026)	Qwen2-7B-T	Qwen2-0.5B-T	Geometric bridge KL	Reasoning, Generation
GLU-SO (Zhang et al., 2026a)	Self (visual PI)	Qwen3-VL-Inst-8B	Entropy-guided KL	ScreenSpot-Pro, CSWorld
MAD-OPD (Wang et al., 2026c)	Multi-teacher debate	Qwen3/3.5 (1.7B-14B)	Confidence-weighted JSD/RKL	Agentic, Code
MSD (Qin et al., 2026)	Self (English CoT PI)	Qwen2.5-7B, Llama-3-8B	DPSW cross-lingual distill	Multilingual Safety
NPD (Riang et al., 2026)	Teacher (parallel prefill)	openPanga-Embedded-1B	Asyrc, CE-KD + A-TPD	GSMA, MATH (6.1x speedup)
SimCT (Sun et al., 2026)	Qwen2.5-7B-Inst, Phi-4-mini	Phi-4-mini, Gemma-2-2B-T	Multi-token continuation KL	Math, Code (cross-tokenizer)
Prune-OPD (Yang et al., 2026g)	R1-Distill-Qwen-7B, Qwen3-4B	R1-Distill-Qwen-1.5B, Qwen3-1.7B	Top-k overlap drift pruning	AMC, AIME, HMMT (37-48% time saved)
SDQ (Zhong et al., 2026)	Qwen3-0.6B/1.7B	Qwen3-0.6B/1.7B	Step-level divergence reweighting	AIME25 (26.1%, 80.8B), SoCode
LiteGLU (Wu et al., 2026b)	Qwen3-VL-32B	2B-3B scale agents	Guided OPD + dual-level GRPO	ScreenSpot-Pro, CS-World, Lite-Bench
VED (Lin et al., 2026a)	Self (video-aware judge)	Qwen3 (1.7B-4B), Qwen3-VL-4B	Dir-mag, decoupled advantage + RL	Qwen3-Videos benchmarks
Uni-OPD (Hou et al., 2026)	Multi-teacher (LLM+MLLM)	NLPACT/NEU students	Dual-perspective OPD	Math, Code, MLLM (6 benchmarks)
TGPO (Liu et al., 2026e)	Teacher (directional PG)	Dense directional teacher RKL	NLPACT, NEU translation	ML2PCT, NEU translation
ROVOPD (Zou et al., 2026)	Code+SVIA teacher	Verifier-reward-weighted FKL	LLN2VA	SWE-Bench Verified
ProteinOPD (Zhang et al., 2026m)	Pretrained PLMs	OPD on student trajectories	Protein property design (8x speedup)	IMEB
HyperEyes (Li et al., 2026d)	External teacher	HyperEyes-30B	Dual-grained efficiency RL	Reasoning (CMU)
DAgger-LLM (Li et al., 2026a)	GPT-4o	Multi-rollback students	Success/failure group signals	5 in-domain + OOD benchmarks
MOFD (Yu et al., 2026b)	Peer-conditioned teacher	Qwen3 family (strong)	BIC-truncated OPD	
Local Collapse (Liu et al., 2026b)	Qwen3 family (strong)	Qwen3 family (weak)		
<i>Self-Distillation</i>				
CRSP (Sang et al., 2026)	Self (concise prompt)	Qwen3-8B/14B	Per-token RKL	AIME, MATH-500
SD-ZERO (He et al., 2026)	Self (revision-conditioned)	Qwen3-4B-Inst, OLMo-3-7B	On-policy self-distill	AIME, MATH, Codeforces
π -Play (Zhang et al., 2026f)	Self (QCP-conditioned)	Qwen3-4B/8B	Multi-agent self-distill	NQ, TriviaQA, HotpotQA
UnSD (Jin et al., 2025c)	Self (EMA + multi-teacher agree)	Qwen2.5-7B/0.5B-3B, Llama-3.1-8B, Gemma-3-4B	SciemerQA, EMA + contrastive + feat. match	SciemerQA, MBPP, CoS-E, ToolAlpaca
SDPO (Hubbott et al., 2026)	Self (textual feedback)	Qwen3-8B	Credit-assignment self-distill	Code, Math
SRPO (Li et al., 2026c)	Self (SDPO+GRPO)	Qwen3-8B	Sample-routed dual-objective	Chem, Phys, Bio, Mat, ToolUse
PRSD (Yu et al., 2026c)	Self (ctx-augmented ref.)	Qwen3-1.7B/4B/8B-Inst	Pre-biased reward-reg. KL	AIME24/25, HMMT25, Tool Use
TT-OPD (Jeong, 2026)	Self (EMA + outcome hints)	Qwen3-3.9B	Turn-level truncated KL	Healthcare AI GYM (18 benchmarks)
Kim & Lee (2026)	Self (GT-conditioned)	Qwen3-8B, R1-Distill-7B, AceReason-7B	Correct-only OPD	MATH-500, AIME24/25
ECRSD (Ke et al., 2026)	Self (entropy-gated)	Qwen3-4B/8B (thinking mode)	Entropy-guided confidence gate	Math reasoning (accuracy-length frontier)
RESO (Zhang et al., 2026n)	Self (reflection-enhanced)	LLM agents	Failure reflection + playbook	Continual agentic learning
SDAR (Liu et al., 2026)	Self (gated RL-OPSD)	Qwen2.5/Qwen3 families	Signoid-gated OPD auxiliary	ALFWorld, WebShop, SearchQA
GEAR (Li et al., 2026g)	Self (GT-conditioned)	LLM agents	Adaptive segment advantage	Agent tasks
<i>Black-Box</i>				
GAD (Ye et al., 2025)	GPT-5-Chat (API only)	Qwen2.5-14B-Inst	Adversarial minimax	LM5YS-Chat, Dolly
Lion (Jiang et al., 2025)	ChatGPT (API only)	LLaMA-7B/13B	NL critique curriculum	BIG-Bench Hard, AGIEval
OVD (Xiong et al., 2026)	QwQ-32B / Env. Agent	Qwen2-5.9B, Llama-3.2-3B	Verbal scores (0-9)	Web QA, MATH
ROPD (Fang et al., 2026)	GPT-5.2, Qwen3-30B-A3B	Qwen3-4B, Gemma3-4B	Rubric-based semantic scoring	AIME24/25, HMMT25, GPQA-D, IFFEval

3.3 Method Selection Considerations

Table 2 suggests that the most successful deployments tend to make deliberate choices at each of the three stages rather than optimizing a single axis in isolation. The constraints and interactions among these axes appear to reduce the space of viable configurations substantially, and we summarize the key selection factors below.

Teacher access constraints. The available signal source determines which objectives are feasible. White-box logit access enables exact token-level divergence computation, supporting methods such as DistiLLM (Ko et al., 2024), G-OPD (Yang et al., 2026e), and cross-tokenizer approaches like DSKD (Zhang et al., 2025b). API-only access restricts supervision to generated text or scalar scores, favoring methods designed for sparse feedback such as OVD (Xiong et al., 2026), GAD (Ye et al., 2025), and LUFFY (Yan et al., 2025). When no external teacher is available, self-distillation methods operate either through privileged information (OPSD (Zhao et al., 2026b), GATES (Stein et al., 2026)) or through verifier-guided self-improvement (SD-ZERO (He et al., 2026)).

Task characteristics. Reasoning and mathematical tasks benefit from mode-seeking objectives (reverse KL or RL-augmented variants) that concentrate probability mass on correct solutions (Gu et al., 2024; Yang et al., 2026e; Xu et al., 2025a). Open-ended generation tasks instead require mode-covering objectives (forward KL or adaptive entropy-aware divergences) that preserve output diversity. Instruction-following tasks occupy a middle ground where moderate on-policy mixing ($\lambda \geq 0.5$) with JSD provides strong quality-compute tradeoffs (Agarwal et al., 2024). Multi-step agentic tasks typically require sequence-level objectives with trajectory-aware credit assignment (Lyu et al., 2025; Zhang et al., 2026o).

Compute budget. Under constrained compute (<500 GPU-hours), prefix-truncated roll-outs (Zhang et al., 2026b) and offline caching strategies (Wu et al., 2026a) reduce on-policy overhead by a wide margin. At medium scale, full on-policy training with adaptive divergences (EOPD (Jin et al., 2026b), AOPD (Jia et al., 2026)) becomes feasible. Large-scale deployments (>5000 GPU-hours) typically adopt multi-stage pipelines that begin with off-

Table 3: OPD methods: Fixed divergence objectives.

Method	Year	Category	Divergence/Objective	Signal	Granularity	Key Innovation
GKD (Agarwal et al., 2024)	2023	Objective	FKL / RKL / JSD	White-box	Token	Unified OPD framework, DAgger for LLMs
MiniLLM (Gu et al., 2024)	2023	Objective	Reverse KL	White-box	Sequence	REINFORCE with teacher as reward
DistiLLM (Ko et al., 2024)	2024	Objective	Skew KL (SKL+SRKL)	White-box	Token	α -smoothing for KL stability
DistiLLM-2 (Ko et al., 2025)	2025	Objective	Contrastive SKL/SRKL	White-box	Token	Asymmetric loss per data source
KETCHUP (Fan et al., 2025)	2025	Objective	Reverse KL	White-box	Sequence	K -step Bellman return estimation
vOPD (Oh et al., 2026)	2026	Objective	Reverse KL + CV baseline	White-box	Token	Control variate variance reduction
AntiSD (Shen et al., 2026a)	2026	Objective	Divergence ascent	Self (PI)	Token	Deliberation-token boosting via PMI sign flip
OPD+ (Zhao et al., 2026a)	2026	Objective	f -div gradient-faithful	White-box	Token	Curvature-aware advantage w_f
Bridging (Jiang & Ferraro, 2026)	2026	Objective	Reverse KL + OT prefix signal	White-box	Token/Span	Near-future trajectory alignment
Decomposed-OPD (Yoon et al., 2026)	2026	Objective	Reverse KL + VGS	White-box	Token	Visual-language gradient steering
Surgical-PT (Lin & Han, 2026)	2026	Objective	Proximal KL-constrained	Oracle	Token/Sequence	Retention via KL-constrained reward
Distributional DAgger (Agrawal et al., 2026)	2026	Objective	Forward CE + REINFORCE	Self (EMA, feedback)	Token	DAgger with teacher-distribution feedback

policy warmup, transition to on-policy refinement, and optionally add RL-based exploration, as reported for Qwen3 (Yang et al., 2025a).

Stability requirements. Common failure modes motivate specific stabilization choices. Length inflation during on-policy training can be mitigated by reference divergence constraints and rollout mixing (Luo et al., 2026a). The flawed prefix trap, where teacher feedback on student-generated errors produces misleading gradients, is addressed by token-level reliability filtering (Xu et al., 2026c; Zheng et al., 2026a). Thinking-pattern mismatch between teacher and student is typically resolved through an off-policy cold-start phase (Li et al., 2026j) before transitioning to on-policy training.

4 Objective Functions and Optimization

The objective function determines what mathematical quantity the student optimizes at each training step, governing gradient geometry, mode coverage, and the resulting performance tradeoffs. Methods in this category have progressed from fixed divergences that apply a single metric uniformly, through adaptive divergences that modulate the metric per-token based on local distributional geometry, to RL-augmented objectives that incorporate reward signals beyond the teacher’s output distribution.

4.1 Fixed Divergence Objectives

GKD (Agarwal et al., 2024) established the canonical OPD framework by drawing on the on-policy imitation learning paradigm (Ross et al., 2011) for autoregressive language models. It samples trajectories from a mixture policy $\pi_{\text{mix}} = \lambda p_\theta + (1 - \lambda)p_{\text{data}}$ and applies standard f -divergences (Forward KL, Reverse KL, or JSD) at each token position. The mixture coefficient λ interpolates between fully off-policy ($\lambda = 0$) and fully on-policy ($\lambda = 1$), serving as a single control for exposure bias mitigation. On instruction-following and summarization benchmarks, moderate on-policy mixing ($\lambda \geq 0.5$) outperforms pure off-policy training, supporting the distributional alignment hypothesis.

Pure KL divergences, though elegant, exhibit numerical instability during on-policy exploration. When the student generates tokens to which the teacher assigns near-zero probability, Forward KL produces unbounded gradients. Reverse KL ignores these regions entirely, creating blind spots in the loss landscape. DistiLLM (Ko et al., 2024) tackles this through *Skewed KL* (SKL), which replaces q_θ in the standard KL with a skewed mixture $\tilde{p} = \alpha p_T + (1 - \alpha)p_\theta$ for $\alpha \in (0, 1]$. The token-level SKL loss is

$$\mathcal{L}_{\text{SKL}} = \mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{mix}}} \left[\sum_{t=1}^{|y|} D_{\text{KL}}(p_T(\cdot|y_{<t}) \parallel \alpha p_T(\cdot|y_{<t}) + (1 - \alpha)p_\theta(\cdot|y_{<t})) \right]$$

where $\mathcal{D}_{\text{mix}}$ mixes ground-truth sequences with cached student rollouts according to DistiLLM’s adaptive off-policy scheduler. Skewing ensures the target density is bounded below by $(1 - \alpha)p_\theta$, which avoids gradient explosion without requiring RL policy gradients. DistiLLM further proposes Skewed Reverse KL (SRKL) for mode-seeking applications, creating a symmetric pair of stabilized divergences. Its successor, DistiLLM-2 (Ko et al., 2025), recognizes that teacher-generated and student-generated data carry qualitatively different learning signals and applies *asymmetric* objectives, using Forward SKL on teacher-generated

data (where mode-covering teaches the student to explore the teacher’s full distribution) and Reverse SRKL on student-generated data (where mode-seeking teaches the student to sharpen its own best responses). This source-aware asymmetry produces consistent improvements over symmetric single-divergence baselines across instruction following and summarization tasks. The appropriate divergence is data-source-dependent.




Figure 2: Forward KL vs. Reverse KL divergence for fitting a student distribution P_S to a bimodal teacher P_T . **(a)** The teacher places equal probability mass on two valid modes (e.g., two correct answers or two stylistic variants). **(b)** Forward KL minimizes $D_{KL}(P_T \parallel P_S)$ and is mode-covering (zero-avoiding): the student spreads a single broad Gaussian over both modes and places substantial probability in the inter-mode hallucination zone (red dashed boundary), corresponding to outputs that are a blend of the two valid answers and therefore neither. **(c)** Reverse KL minimizes $D_{KL}(P_S \parallel P_T)$ and is mode-seeking (zero-forcing): the student concentrates sharply onto a single peak (Mode 2), dropping Mode 1 entirely but producing crisp, coherent generations within the chosen mode. This trade-off motivates the adaptive-divergence methods in §4 (EOPD, AOPD), which interpolate between the two regimes based on per-token teacher entropy or log-ratio statistics.

DistiLLM-2 applies different divergences to different data sources, but operates at a coarse granularity (entire sequences grouped by source). The fixed-divergence family faces a more fundamental limitation rooted in the heterogeneity of autoregressive token distributions *within* a single sequence. The choice between Forward KL (mode-covering) and Reverse KL (mode-seeking) is a global decision applied uniformly across all token positions (Figure 2). The teacher’s distributional structure, however, varies considerably within a single sequence. At a mathematical operator token, the teacher concentrates probability mass on one or two correct symbols, so mode-seeking is appropriate because the student must commit to the correct answer. At a conjunction or filler word, dozens of tokens share comparable probability, so mode-covering is appropriate because any synonym suffices. Applying a single divergence to both positions therefore wastes gradient signal in one regime or discards coverage information in the other. This per-position mismatch between the geometry of the loss and the geometry of the target distribution motivates the adaptive methods of Section 4.2.

Sequence-level objectives offer a sequence-level perspective. MiniLLM (Gu et al., 2024) elevates Reverse KL to the full sequence level, where the student p_θ appears inside the sampling expectation. The central challenge is computing the gradient of an expectation over p_θ when p_θ itself is parameterized by θ . Starting from $D_{KL}(p_\theta \parallel p_T) = \mathbb{E}_{y \sim p_\theta} [\log p_\theta(y|x) - \log p_T(y|x)]$ and applying the policy gradient theorem, MiniLLM derives (Eq. 2 of Gu et al., 2024):

$$\nabla_\theta \mathcal{L}_{\text{MiniLLM}} = -\mathbb{E}_{y \sim p_\theta} \left[\sum_{t=1}^{|y|} (R_t - 1) \nabla_\theta \log p_\theta(y_t | y_{<t}) \right], \quad R_t = \sum_{t'=t}^{|y|} \log \frac{p_T(y_{t'} | y_{<t'})}{p_\theta(y_{t'} | y_{<t'})}$$

where the -1 arises from the $-\log p_\theta(y|x)$ entropy term of the reverse-KL objective. Minimizing sequence-level Reverse KL is thus mathematically equivalent to policy gradient

RL with per-step reward $r_t = \log[p_T(y_t|y_{<t})/p_\theta(y_t|y_{<t})]$ and an entropy-induced constant. MiniLLM further decomposes R_t into a single-step quality term and a future-return term (Eq. 3 of Gu et al., 2024), computing $\mathbb{E}_{y_t \sim q_\theta}[r_t]$ in closed form over the vocabulary instead of by Monte Carlo sampling. This decomposition reduces variance by isolating the dominant source of gradient noise (single-step quality) from the accumulated future return, and supports credit assignment at the token level. Per-step returns R_t attribute the trajectory’s deviation from the teacher to specific token decisions rather than spreading uniform blame across the sequence, placing the optimization within the RL framework with the teacher’s log-probability as a dense reward. This connection to policy optimization aligns with MiniLLM’s empirical strength on reasoning tasks where sequence-level coherence matters more than token-level accuracy, at the cost of higher training compute because REINFORCE estimation in combinatorial output spaces requires more iterations and careful baseline subtraction to converge. KETCHUP (Fan et al., 2025) addresses this variance problem by replacing the full cumulative return G_t with a K -step Bellman return estimate derived from the optimality equation, reducing gradient variance at a power rate in K while preserving the sequence-level objective (Theorem 1 of Fan et al., 2025).

The token-level vs. sequence-level tradeoff. The choice between token-level and sequence-level objectives represents a core bias-variance tradeoff whose resolution foreshadows the RL-augmented methods of §4.3. Token-level methods (GKD, DistiLLM) produce $|V|$ -dimensional gradient signals at every position, resulting in low variance but potentially biased gradients when the teacher’s token-level distribution is poorly calibrated (the flawed prefix trap of Section 6.1). Sequence-level methods (MiniLLM, KETCHUP) yield unbiased gradients through the policy gradient theorem but suffer from high variance due to single-sample REINFORCE estimation over exponentially large sequence spaces. The practical resolution is task-dependent. For instruction following where token-level teacher quality is high, GKD-style methods tend to dominate. For reasoning tasks where sequence-level coherence matters more, MiniLLM-style methods are often preferred despite their variance, and KETCHUP’s K -step Bellman returns offer an intermediate bias-variance point on this spectrum. Recognizing that MiniLLM is already performing policy gradient RL (with the teacher’s log-probability as reward) suggests that sequence-level distillation and RL-augmented distillation are not distinct paradigms but endpoints of a continuum parameterized by the reward source, with pure teacher signal at one end, pure task reward at the other, and hybrid objectives (G-OPD, RLKD) interpolating between them.

Full-vocabulary vs. sampled-token KL. Orthogonal to the token-level vs. sequence-level choice, a practical implementation decision determines how much of the teacher’s distribution is used at each position. Many RL-based OPD implementations approximate the per-position KL by computing $\log p_\theta(y_t|y_{<t}) - \log p_T(y_t|y_{<t})$ only at the *sampled* token y_t , treating this scalar as a per-token advantage within a policy gradient framework (Lu & Thinking Machines Lab, 2025). While computationally lightweight (requiring only one teacher log-probability per position), this single-sample estimate incurs high gradient variance and discards the distributional information across non-sampled tokens. DeepSeek-V4 (DeepSeek-AI, 2026) reports that *full-vocabulary* logit distillation, computing the exact $D_{\text{KL}}(p_\theta(\cdot|y_{<t}) \parallel p_T(\cdot|y_{<t}))$ over all $|V|$ tokens at each position, yields more stable gradients and more faithful knowledge transfer at the cost of transmitting the full teacher logit vector. Their engineering solution (caching teacher hidden states and reconstructing logits on-the-fly through the prediction head) makes this feasible even with 10+ trillion-parameter teachers over vocabularies exceeding 100K tokens, which positions full-vocabulary KL as an attractive choice for industrial-scale OPD when infrastructure permits. Lightning-OPD (Wu et al., 2026a) further reports that the full teacher forward pass can be largely avoided by precomputing teacher log-probabilities over SFT rollouts, recording $4.0\times$ speedup over standard OPD under a *teacher consistency* condition (using the same teacher for both SFT data generation and distillation).

Divergence ascent for deliberation. AntiSD (Shen et al., 2026a) reverses the self-distillation gradient entirely. A pointwise mutual information analysis shows that privileged-context conditioning inflates the teacher’s confidence on tokens already implied by the solution (structural connectives, verifiable claims) while deflating it on deliberation tokens (*Wait, Let, Maybe*) that drive multi-step search. By *ascending* the student-teacher divergence instead of

descending it, AntiSD flips the per-token sign, yielding a naturally bounded advantage that boosts deliberation tokens in a single update step. An entropy-triggered gate disables the anti-distillation term once teacher entropy collapses, preventing degenerate behavior when the teacher distribution is already peaked. Across five models from 4B to 30B parameters, AntiSD reaches GRPO-level accuracy in $2\text{--}10\times$ fewer training steps and, given equal training budget, exceeds GRPO by up to 11.5 points on math reasoning, suggesting that the standard “descend toward the teacher” assumption can be relaxed when the privileged signal is systematically biased.

Control variate variance reduction. vOPD (Oh et al., 2026) offers a third resolution to this variance-cost tradeoff by importing the *control variate baseline* from policy-gradient RL into the OPD estimator. Casting the single-sample OPD gradient (Eq. 3 of Oh et al., 2026) as a REINFORCE update with per-token reward $r_t = \log p_T(y_t|c_t) - \log p_\theta(y_t|c_t)$, vOPD subtracts a baseline $b_t(c_t)$ from r_t before multiplying by $\nabla_\theta \log p_\theta$. The natural value function for this reward admits a closed form as the per-token negative reverse KL between student and teacher, $V^{p_\theta}(c_t) = -D_{\text{KL}}(p_\theta(\cdot|c_t) \parallel p_T(\cdot|c_t))$, computable directly from the forward pass that already produces the OPD objective without requiring an auxiliary critic or additional rollouts. Because the baseline depends only on the context c_t and not on the sampled token y_t , subtracting it preserves the unbiasedness of the gradient while reducing variance. A top- k approximation of this baseline (restricting the KL summation to the student’s k most probable tokens) further lowers compute, and because the approximation remains independent of the sampled action, unbiasedness still holds regardless of k . Across six reasoning benchmarks on Qwen3-1.7B and 4B, vOPD yields an average +3% absolute gain over vanilla single-sample OPD (up to +6.2% on MATH500), matches the accuracy of full-vocabulary OPD while reducing wall-clock time by up to 57.7%, and broadly outperforms top- k OPD, which adds compute but introduces bias. The result positions vOPD at a favorable point on the bias-variance-compute Pareto frontier for on-policy distillation of reasoning models.

Advantage and gradient redesigns within fixed-divergence training. The variance-control results above retain the conventional advantage and gradient definitions and tune only the estimator. A separate line of work asks whether the advantage itself is correctly specified once the on-policy rollout enters the picture. OPD+ (Zhao et al., 2026a) re-derives the per-token weight from an f -divergence gradient analysis and observes that the conventional stop-gradient reward $r_t = \log[p_T(y_t|y_{<t})/p_\theta(y_t|y_{<t})]$ corresponds to $f(u) = -\log u$ but ignores the second-order curvature, and the gradient-faithful weight $w_f(u) = -f(u) + u f'(u)$ aligns the OPD update with the true f -divergence gradient and reduces a systematic bias that grows when the student strays far from the teacher. Where OPD+ corrects the per-token advantage, Bridging Reasoning Trajectories (Jiang & Ferraro, 2026) addresses the short-horizon trajectory signal by detecting trajectory-level divergence through a short-window optimal-transport distance and injecting an OT-based teacher signal whenever the student’s near-future continuation drifts away from teacher trajectories. Low-OT but high-loss tokens are downweighted to suppress false alarms, and a two-stage schedule (off-policy warmup followed by on-policy distillation) confines the OT correction to the on-policy regime where it is most informative. The same per-position-mismatch concern reappears for vision-language students, where reverse-KL gradients can be dominated by the language prior and effectively bypass visual grounding. Decomposed-OPD (Yoon et al., 2026) responds with Visual Gradient Steering (VGS), a decomposition that splits the per-token gradient into a language-prior component and a visual-grounding component, normalizes the two by gradient norm, and amplifies the visual component by a Visual Dependency Score (VDS) read off from cross-attention statistics. A language-preservation regularizer prevents the rebalanced gradient from degrading text-only competence. Together, these three methods illustrate that the same fixed-divergence skeleton can be improved by acting on three distinct levels of the gradient signal, the per-token weight (OPD+), the short-horizon trajectory (Bridging), and the modality decomposition (Decomposed-OPD), without changing the divergence type itself.

Proximal and DAgger-style reformulations. Two further methods retain the fixed-divergence template but redefine *what is being constrained*. Surgical Post-Training (Lin & Han, 2026) treats reasoning injection as proximal on-policy distillation under a KL constraint to the pre-update policy, with the analytical and empirical claim that this KL-constrained

Table 4: OPD methods: Adaptive divergence objectives.

Method	Year	Category	Divergence/Objective	Signal	Granularity	Key Innovation
EOPD (Jin et al., 2026b)	2026	Objective	Continuous FKL+RKL blend	White-box	Token	Entropy-aware on-policy distillation
AOPD (Jia et al., 2026)	2026	Objective	Asymmetric PG/FKL switch	White-box	Token	Advantage-sign token-level switching
CREDIT (Shen et al., 2026b)	2026	Self (EF)	Contrastive PMI decomposition	Self (contrastive)	Token	Input-specific credit for OPSD reward
OGLS-SD (Yang et al., 2026f)	2026	Self (EF)	Outcome-guided logit steering	Self (outcome)	Token	Outcome-based teacher calibration
Position-Weighted OPSD (Liu et al., 2026d)	2026	Objective	Position-aware reliability	Self (PI)	Token	Per-position reliability score
Direction-Adaptive SD (Zhang et al., 2026c)	2026	Objective	Entropy-routed polarity	Self (PI)	Token	High-entropy teacher reversal
Token-Teachability (Wang et al., 2026j)	2026	Objective	Teachability-gated KL	White-box	Token	Fixed-context learnability diagnostic
Lookahead Group Reward (Liu et al., 2026f)	2026	Objective	RKL + look-ahead group reward	White-box	Token	One-step-ahead teacher confidence
RAFT (Li et al., 2026i)	2026	Objective	Top-K temperature distill	Self	Token	Data refinement + EMA loss balance
Trust-Region OPD (Xing et al., 2026)	2026	Objective	Adaptive trust-region KL	White-box	Token	Speculative-decoding-ratio gating
GNDPO (Hao et al., 2026)	2026	Objective	Global batch-normalized KL	White-box	Token	Batch-relative KL normalization for MLLM stability

reward formulation is what supports retention of existing capabilities during reasoning specialization. Distributional DAGger (Agrawal et al., 2026) returns to the original DAGger (Ross et al., 2011) blueprint that motivated GKD and generalizes scalar feedback into a teacher distribution. The student rolls out under its own policy, the feedback (e.g., a verifier or rich preference signal) conditions an EMA self-teacher, and the resulting forward cross-entropy is paired with a REINFORCE future-credit term so that token-level updates account for downstream consequences inside a PPO-style trust region. A top- K teacher truncation keeps the per-token cost bounded. These two reformulations frame fixed-divergence OPD as an instance of constrained policy improvement, with KL-constrained reward acting as the retention mechanism (Surgical Post-Training) and a future-aware DAGger update acting as the credit-assignment mechanism (Distributional DAGger), both of which set up the per-position adaptive divergences of Section 4.2.

4.2 Adaptive Divergence Objectives

Adaptive methods tackle this per-position mismatch by selecting or interpolating divergences based on local distributional geometry (Table 4). When the teacher is certain (low entropy), mode-seeking Reverse KL preserves precision, and when uncertain (high entropy), mode-covering Forward KL preserves diversity.

EOPD (Jin et al., 2026b) uses the teacher’s per-token entropy $H(p_T(\cdot|y_{<t}))$ as a gating signal. The base objective is Reverse KL at all positions, with Forward KL *additively* activated at high-entropy tokens where the teacher distributes probability across multiple valid continuations:

$$\mathcal{L}_{\text{EA}} = \mathbb{E}_{y \sim p_\theta} \left[\sum_{t=1}^{|y|} D_{\text{KL}}(p_\theta \parallel p_T) + \mathbb{I}[H_t > \tau_H] \cdot D_{\text{KL}}(p_T \parallel p_\theta) \right]$$

EOPD outperforms fixed divergence baselines across six competition-math reasoning benchmarks, confirming that the appropriate divergence is position-dependent. Conditioning on teacher uncertainty alone (a marginal property) gives EOPD a stable routing signal independent of student quality, making it suitable for early training when the student’s distribution is less reliable.

Self-distillation reward calibration. Two recent methods address pathologies specific to the self-distillation reward signal by injecting contrastive or outcome-based corrections. CREDIT (Shen et al., 2026b) decomposes the standard OPSD token reward into input-specific and input-generic components via a batch-contrastive baseline, suppressing generic shortcuts that inflate pointwise mutual information without contributing task-relevant reasoning (Section 5.2.3). OGLS-SD (Yang et al., 2026f) calibrates the self-reflected teacher distribution using verifiable outcome rewards to contrast successful and failed on-policy trajectories, addressing reflection-induced bias that shifts teacher logits away from optimal targets (Section 5.2.3). Both methods operate within the self-distillation paradigm and are detailed in Section 5.2.3.

AOPD (Jia et al., 2026) extends the per-token adaptive principle by diagnosing *why* standard OPD gradients fail at specific token positions. In non-positive advantage regions ($A_t \leq 0$), standard advantage-weighted policy gradients exhibit three pathologies that intensify under weak initialization. Heavy-tailed variance from unbounded negative advantages destabilizes updates, vanishing gradients at zero-advantage positions waste teacher signal, and “exploration black holes” trap learning because suppressing a wrong token only

redistributes mass proportional to the student’s own prior, unable to boost low-probability correct alternatives. AOPD addresses this by switching the objective at the token level. For positive-advantage tokens ($A_t > 0$), standard policy gradient exploitation is preserved. For non-positive tokens, the framework substitutes localized forward-KL guidance on the teacher’s top-K support, creating an asymmetric exploitation-imitation mechanism. A threshold parameter τ (set to zero by default, corresponding to the probability difference between teacher and student) determines the switching point, with $\tau = -1$ recovering standard OPD (intervention rarely triggers because $P_T - P_S \leq -1$ is almost never satisfied) and $\tau = 1$ recovering GKD (intervention everywhere because $P_T - P_S \leq 1$ generally holds). On competition-level mathematics (AIME 2024/2025, HMMT), AOPD attains average gains of +4.09 with strong initialization and +8.34 with weak initialization over standard on-policy baselines, the larger weak-initialization gap points to the exploration black hole as a dominant failure mode when the student starts far from the teacher. AOPD adapts the *objective class itself* (policy gradient vs. divergence minimization) based on the gradient’s structural health at each position, complementing entropy-based routing in EOPD that adapts the divergence *type*.

Token reliability and teachability. A parallel line of work refines per-token routing by replacing implicit signals (teacher entropy, advantage sign) with direct measurements of how trustworthy or learnable each position is. Position-Weighted OPSD (Liu et al., 2026d) observes that teacher entropy alone is insufficient to predict supervision quality at a given position, and that reliability degrades with depth into the student rollout because later positions condition on increasingly off-policy prefixes. The method modulates the per-token loss by a position-aware reliability score read off from teacher behavior on student-generated sequences, recovering accuracy where uniform weighting flattens the gradient. Token-Teachability OPD (Wang et al., 2026j) attacks the same problem from the student side. A fixed-context diagnostic measures whether the student can plausibly absorb the teacher’s preferred token at each position, separating genuinely learnable tokens from tokens that are merely high-disagreement. Selective OPD restricted to teachable tokens outperforms entropy-based and disagreement-based selection on reasoning benchmarks. Together, the two methods localize where the teacher’s signal carries useful information (Position-Weighted) and where the student is in a position to use it (Token-Teachability), tightening the routing logic that EOPD and AOPD apply at coarser per-position granularities. Direction-Adaptive Self-Distillation (Zhang et al., 2026c) pushes the routing further by reversing the teacher’s polarity at high-entropy tokens, similar in spirit to AntiSD’s divergence ascent (Section 4.1) but gated on entropy as in EOPD, so that the student is pulled away from the teacher at deliberation points where teacher uncertainty signals exploration value rather than guidance value.

Look-ahead and trust-region adaptation. Per-position routing can also be improved by extending the teacher signal across positions or by gating updates by stability. Lookahead Group Reward (Liu et al., 2026f) identifies *Supervision Fidelity Decay*, the observation that per-token teacher confidence on student rollouts degrades along the trajectory as the student drifts off-policy. The method augments the per-token signal with a one-step-ahead group reward, computing teacher confidence under a tree-attention mask after a single look-ahead step, group-normalizing the resulting score, and gating its activation on local entropy. Trust-Region OPD (Xing et al., 2026) adapts the per-step trust region directly. The acceptance ratio of speculative decoding is repurposed as a cheap proxy for student-teacher policy divergence, an outlier estimator clips top- k FKL excursions, and an off-policy teacher prefix is mixed in with cosine annealing during early training when the student rollout is least informative. The K1 reverse-KL estimator anchors the resulting trust region, recovering the EOPD and AOPD gating principle at the trajectory level. RAFT (Li et al., 2026l) couples the same adaptive principle to the data side by alternating self-conditioned rewriting (which refines noisy domain prompts via cosine-similarity semantic filtering and multi-expert fusion) with answer-conditioned on-policy distillation that uses top- K temperature distillation against the original instruction-tuned model. An EMA-based adaptive loss balancer keeps the distillation and SFT components from interfering, supporting the broader observation that on-policy distillation in domain specialization regimes benefits from adapting both the data distribution and the per-token weight.

Table 5: OPD methods: RL-augmented objectives.

Method	Year	Category	Divergence/Objective	Signal	Granularity	Key Innovation
G-OPD (Yang et al., 2026e)	2026	Objective	KL-Constrained RL	White-box	Token/Seq	OPD $\equiv$ dense KL-RL, reward extrapolation
RLKD (Xu et al., 2025b)	2025	Objective	KD + Reward Model	Reward Model	Sequence	Generative Structure Reward Model (GSRM)
KDRL (Xu et al., 2025a)	2025	Objective	Joint KD + RL	White-box	Token/Seq	On-policy KL regularizer during RL
RLAD (Zhang et al., 2026a)	2026	Objective	Trust Region Ratio	White-box	Token/Seq	PPO-style selective teacher following
AlignDistil (Zhang et al., 2025a)	2025	Objective	DPO + Reverse DPO	White-box	Token	Token-level RLHF via contrastive DPO reward
SuperCorrect (Yang et al., 2025b)	2024	Objective	Template + Cross-DPO	White-box	Sequence	Hierarchical thought scaffolding
SCoRe (Lyu et al., 2025)	2025	Objective	Earliest-error correction	White-box	Sequence	Root-cause error targeting
CMDP-KD (Zimmer et al., 2025)	2025	Objective	Constrained KL	White-box	Sequence	Unaugmented CMDP with hard KL budget
X-KD (Cai & Yuan, 2026)	2026	Objective	Inverse RL + KL	White-box	Token/Seq	AVRIL-based reward recovery distillation
TSO-KD (Kim & Baek, 2026)	2026	Objective	Indirect + Direct KD	White-box	Sequence	Indirect re-ranking + selective direct KL
REOPOLD (Ko et al., 2026)	2026	Objective	Log-ratio token reward	White-box	Token	Mixture reward clipping + entropy-dynamic sampling
CoDistill-GRPO (Kwon et al., 2026)	2026	Objective	GRPO + KD reward	White-box	Sequence	Bidirectional co-distillation, 18% speedup
dGRPO (Ramos et al., 2026)	2026	Objective	GRPO + dense OPD	White-box	Token/Seq	Long-context reasoning with LongBlocks data
Sparse-to-Dense (Xu et al., 2026d)	2026	Objective	Staged RL $\rightarrow$ OPD $\rightarrow$ RL	White-box	Token/Seq	Reward-density allocation principle
TGPO (Liu et al., 2026e)	2026	Objective	Directional teacher PG	White-box	Token	Dense directional teacher guidance on RKL
RWOPD (Zou et al., 2026)	2026	Objective	Verifier-weighted FKL	White-box	Sequence	Verifier-reward-weighted OPD for NL2SVA
AMR-SD (Wei et al., 2026)	2026	Objective	Asymmetric meta-reflective	Self	Token	Meta-reflection self-teacher in RLVR
OPPO (Li et al., 2026k)	2026	Objective	Bayesian value recursion	Oracle (PI)	Token	Likelihood-ratio per-token credit
StepOPSD (Zhang et al., 2026j)	2026	Objective	Step-aware preference + KL	Self (hindsight)	Step/Token	Dual-granularity for multi-turn agents
Self-Eval OPD (Zhang et al., 2026h)	2026	Objective	Cyclic GRPO + masked judge KD	External judge	Token	Self-evaluation channel calibration
RLCSD (Pan et al., 2026)	2026	Objective	Contrastive RKL in GRPO	White-box	Token	Correct-hint teacher pull, wrong-hint push in two-path distillation

Batch normalization for stable on-policy distillation. GNDPO (Hao et al., 2026) addresses an instability that arises when adaptive KL losses are applied to multimodal LLMs (MLLMs). Standard per-sample KL normalization can produce large gradient variance across mini-batches when the token-count distribution varies widely (e.g., vision inputs of different resolutions). GNDPO normalizes the KL objective across the entire batch relative to a global reference distribution, producing smoother gradient estimates and more stable training dynamics for MLLMs trained with on-policy distillation.

The field has moved noticeably from “which divergence?” to “which divergence *where*?” The evidence from EOPD and AOPD points to a consistent finding across evaluated settings, that no single fixed divergence dominates uniformly, and that even the simplest adaptive routing (a hard entropy threshold) often outperforms the best-tuned fixed alternative. Yet adaptive divergences, however sophisticated their per-token routing, share a structural limitation with their fixed counterparts. They optimize a distance to the teacher distribution. Perfect optimization under any f -divergence yields a student that matches the teacher but rarely exceeds it. When the teacher itself makes errors or lacks reasoning paths that a well-guided student could discover, divergence minimization can act more as a ceiling than as a floor.

4.3 RL-Augmented Objectives

The capability ceiling identified at the end of the previous section motivates a different class of objectives. RL-augmented methods inject an external performance signal that can steer the student toward solutions the teacher itself would be unlikely to produce.

OPD as KL-constrained RL. G-OPD (Yang et al., 2026e) formalizes an equivalence between standard OPD and dense KL-constrained reinforcement learning. It formulates distillation as maximizing a token-level reward defined as the log-probability ratio between teacher and reference policy, penalized by KL divergence from the reference:

$$\max_{\theta} \mathbb{E}_{y \sim p_{\theta}} \left[\sum_{t=1}^{|y|} \alpha \log \frac{p_T(y_t | y_{<t})}{p_{\text{ref}}(y_t | y_{<t})} - D_{\text{KL}}(p_{\theta}(\cdot | y_{<t}) \parallel p_{\text{ref}}(\cdot | y_{<t})) \right]$$

When $\alpha = 1$, this reduces to standard Reverse KL distillation. When $\alpha > 1$, the objective forces the student to *extrapolate* beyond the teacher’s probability mass, allowing a phenomenon documented as *Reward Extrapolation*, whereby the student discovers previously unseen, structurally sound reasoning paths in regions where the teacher’s generation probability $p_T(y|x)$ is low but the outcome reward remains high. In the multi-teacher setting where several domain-specific RL experts share the same base model, reward extrapolation (ExOPD) produces a unified student that surpasses all its same-size domain teachers, a result that positions OPD as a potential amplification mechanism rather than a lossy compression step.

This equivalence carries a broader consequence for the section as a whole. It connects the distillation and RL communities: the objectives discussed here (GKD’s token-level KL, MiniLLM’s sequence-level Reverse KL, G-OPD’s reward-augmented formulation) can

be viewed as special cases of KL-constrained policy optimization with different reward definitions and constraint strengths. Advances in KL-constrained RL (better trust regions, adaptive penalty coefficients, variance reduction) are therefore likely to transfer to OPD, and vice versa.

Gradient decomposition and the KD+RL complementarity. The unified KD+RL framework of Li et al. (2025) decomposes the hybrid gradient into two components, a dense KD gradient for token-level imitation and a Monte Carlo RL gradient for reward maximization. The KD component dampens the high variance inherent in policy gradient estimation, while the RL component prevents the student from collapsing onto teacher modes that are suboptimal for the downstream task. REOPOLD (Ko et al., 2026) instantiates this insight by interpreting on-policy distillation as policy optimization where the teacher-student log-likelihood ratio acts as a token reward. It stabilizes training through three mechanisms, mixture-based reward clipping to prevent over-trust of extreme teacher signals, entropy-based dynamic sampling that selectively weights teacher feedback at high-uncertainty tokens, and a unified exploration-to-refinement schedule. Empirically, REOPOLD attains $6.7\text{--}12\times$ greater sample efficiency than recent RL approaches and allows a 7B student to match a 32B teacher in visual reasoning with $\sim 3.3\times$ inference speedup.

Structure-aware reward as an alternative to dense KL. While G-OPD and REOPOLD retain token-level KL as the primary distillation signal (differing only in how they augment or stabilize it), RLKD (Xu et al., 2025b) departs from this template entirely by replacing token-level KL with a Generative Structure Reward Model (GSRM). The GSRM parses both teacher and student reasoning paths into sequences of meta-reasoning and solving steps, then scores structural alignment between the two sequences. This structural reward is combined with a task-level outcome reward (e.g., math accuracy) and optimized via GRPO:

$$R_{\text{total}}(x, y) = \alpha R_{\text{GSRM}}(y, y^{\text{teach}}) + (1 - \alpha) R_{\text{outcome}}(x, y)$$

The main idea is that reasoning distillation need not operate at the token level. By supervising at the *step structure* level, RLKD avoids the vocabulary-coupling assumption of logit-based methods while still transferring the teacher’s multi-branch reasoning strategy. Empirically, RLKD trained on only 0.1% of the data under a pure RL regime surpasses standard SFT-RL pipelines, showing that structural reward can substitute for dense KL when the teacher’s value lies in its reasoning organization rather than its token-level distribution. This trajectory from dense token-level KL (G-OPD) through variance-reduced token rewards (REOPOLD) to step-level structural matching (RLKD) traces an increasing abstraction of what “teacher knowledge” means, moving from distributional fidelity to behavioral organization.

Bidirectional co-distillation. CoDistill-GRPO (Kwon et al., 2026) addresses the common failure of GRPO on small models (where sparse rewards on difficult tasks yield negligible learning signal) through a co-distillation algorithm that simultaneously trains a large and a small model with carefully designed GRPO objectives. The small model uses an on-policy KD reward derived from the large model’s distribution as dense supervision, while the large model trains on the small model’s rollouts with importance reweighting, reducing the computational overhead of rollout generation. On Qwen2.5-Math-1.5B, CoDistill-GRPO reaches 32% on Minerva, an 11.6-point increase over the base model and 6.0 points over standard GRPO. The large model (Qwen2.5-Math-7B) nearly matches standard GRPO despite training on small-model rollouts, providing an approximate 18% speedup. The bidirectional design avoids the common assumption that a pre-trained oracle must exist, instead producing both the teacher signal and the student improvement simultaneously.

Directional teacher guidance on reverse KL. TGPO (Liu et al., 2026e) departs from the additive KD+RL template by embedding teacher guidance *inside* the policy gradient itself. Rather than treating the teacher’s distribution as a separate regularization term, TGPO uses the teacher’s per-token likelihood ratio as a directional signal that steers the reverse-KL gradient toward regions the teacher considers promising, concentrating updates on tokens where teacher and student disagree most. This dense directional coupling avoids the balancing hyperparameter that joint KD+RL objectives require (the α in RLKD, the KL penalty in KDRL) and instead lets the teacher’s confidence modulate the RL gradient at every position. On math reasoning benchmarks (e.g., MATH-500), TGPO outperforms both

standalone OPD and standalone GRPO, tighter coupling between teacher guidance and policy optimization can be more effective than additive combination.

Dense guidance for long-context reasoning. dGRPO (Ramos et al., 2026) targets the same sparse-reward failure mode as CoDistill-GRPO but through a different architecture. Rather than co-training two models, dGRPO augments a single student’s GRPO training with dense token-level guidance from a stronger teacher via an additive OPD term. The method specifically targets the long-context reasoning regime (up to 128K tokens), where sparse rewards are especially unstable because correct reasoning paths become increasingly rare in longer sequences and the effective learning signal degrades with context length. The accompanying LongBlocks dataset provides synthetic long-context tasks spanning multi-hop reasoning, contextual grounding, and long-form generation. The combined KD+RL objective provides more stable optimization than either component alone, addressing the regime where pure GRPO suffers from sample inefficiency (most long-context rollouts fail entirely) and pure OPD lacks direct reward optimization for discovering previously unseen long-range reasoning patterns. Where CoDistill-GRPO uses the KD signal bidirectionally between two models, dGRPO applies it unidirectionally from a fixed teacher, trading the mutual bootstrapping benefit for architectural simplicity.

Sparse-to-dense reward allocation. Xu et al. (2026d) articulate a reward-density principle that reframes GRPO and OPD not as competing recipes but as different reward-density regimes. Sparse sequence-level reward is most productive on the strongest model where exploration can discover new solutions (teacher-side RL), while dense token-level teacher reward is most effective for compressing learned behavior into a smaller model (student-side OPD). The operational lesson is a three-stage pipeline, GRPO on the teacher for capability discovery, a forward-KL warmup followed by OPD as a dense bridge, and optional student-side sparse RL after the bridge. At fixed Qwen3-1.7B deployment size, an RL-improved 8B teacher distilled through this dense bridge outperforms direct student-side GRPO (79.3% vs. 75.9% on MATH, 25.2 vs. 19.8 on AIME 2024), while transfer from the same teacher *before* RL underperforms. The bridge also makes subsequent student-side GRPO effective where it previously was not, lifting MATH from 75.4% to 78.5%. The teacher-quality ordering replicates on Llama-3.1-8B-Instruct with a Llama-3.3-70B-Instruct teacher.

Joint vs. sequential optimization. Given that KD and RL supply complementary gradient signals (dense imitation vs. sparse reward maximization), a persistent engineering challenge is *how* to combine them within a single training pipeline without mutual interference. Sequential approaches (distill-then-RL or RL-then-distill) tend to suffer from catastrophic interference, as the RL stage degrades distilled knowledge through policy drift, while the KD stage suppresses previously unseen solutions discovered through exploration. KDRL (Xu et al., 2025a) implements joint optimization by adding an on-policy KL regularizer between student and teacher *during* RL training, mitigating policy drift while still allowing reward-driven exploration. CMDP-KD (Zimmer et al., 2025) extends this regularization scheme by replacing the soft KL penalty with a *hard* KL budget via constrained MDP formulation. The state-augmented Saute formulation can be simplified for LLMs by removing state augmentation entirely, since the full history already determines the remaining budget, yielding constraint satisfaction guarantees without requiring teacher access at inference. RLAD (Zhang et al., 2026o) refines this further by recognizing that unconditional KL regularization is itself suboptimal. Its Trust Region Ratio Distillation (TRRD) replaces the standard KL penalty with a PPO-style likelihood-ratio objective anchored to a teacher-old-policy mixture. This selective strategy follows the teacher only when its signal improves the policy update, preventing blind KL minimization from dragging the student toward suboptimal trajectories on problems the teacher itself cannot solve. The evolution from KDRL’s uniform regularization to RLAD’s selective trust region mirrors the fixed-to-adaptive shift in divergence objectives (Section 4.1→4.2), echoing a recurring pattern across this survey that conditioning the learning signal on local geometry often works better than applying it uniformly.

Preference-based objectives as implicit RL. When logit access is unavailable or when sequence-level quality judgments are more natural than token-level distributions, several methods map the KD+RL problem onto Direct Preference Optimization (DPO). AlignDistil (Zhang et al., 2025a) derives token-level rewards from contrastive DPO and reverse

DPO models, converting sequence-level RLHF into per-token policy distillation that applies both standard and reverse preference objectives at the token granularity. OVD (Xiong et al., 2026) queries the teacher for pairwise preferences over student-generated responses, avoiding continuous probability extraction entirely. These preference-based objectives share a deep connection to token-level KL methods. Under DPO’s closed-form solution, the optimal policy satisfies $\pi^*(y|x) \propto \pi_{\text{ref}}(y|x) \exp(r(y, x)/\beta)$, and substituting the teacher’s log-probability as the implicit reward yields the geometric mixture distribution targeted by GKD with λ -mixing. The practical implication is that DPO-style methods are often most useful when constructing preference pairs is cheaper than computing full teacher logits. PBSO (Yu et al., 2026c) pushes this connection further in the self-distillation regime by deriving the reward-regularized optimum $\pi^* \propto \pi_{\text{teach}} \exp(r/\beta)$ against a context-augmented teacher rather than a static reference policy, a reward-tilted target that improves over the teacher itself under the same objective whenever the reward is non-degenerate (Section 5.2.1). The spectrum exposes a unifying gradient. Token-level KL (G-OPD, REOPOLD), step-level structure matching (RLKD), and sequence-level preference ranking (AlignDistil, OVD) all implement KL-constrained policy optimization at different abstraction levels, with the appropriate level determined by signal availability and the granularity at which teacher expertise manifests.

Verifier-reward-weighted distillation. When a domain-specific verifier can assign correctness scores to student rollouts, these scores offer a natural per-sample weighting for the distillation objective. RWOPD (Zou et al., 2026) instantiates this idea for NL-to-SystemVerilog generation, where an open property-equivalence verifier scores each student-generated hardware description against a formal specification. The verifier reward weights the forward-KL loss so that rollouts closer to functional correctness receive stronger teacher supervision, while clearly incorrect rollouts are attenuated rather than discarded entirely (preserving gradient signal from partial successes). The approach sets new benchmarks on NL2SVA benchmarks, a result that positions verifier-weighted OPD as well-suited to formal-verification domains where binary or graded correctness is cheaply computable and the reward landscape is sparse under standard RL.

Fine-grained credit assignment. All methods above, whether operating at the token, step, or sequence level, share one limitation. They apply supervision uniformly across the student’s trajectory. When the student makes a single error in step 3 of a 10-step proof, wholesale imitation retrains all 10 steps, diluting the corrective gradient on the actual failure point. SuperCorrect (Yang et al., 2025b) presents hierarchical thought templates with cross-model DPO for error localization, identifying which reasoning step diverged and concentrating feedback there. SCoRe (Lyu et al., 2025) takes a more aggressive approach, correcting *only the earliest error* in the student’s trajectory, concentrating all supervision at the single point of maximum corrective impact. This exploits the cascading nature of reasoning errors, where correcting the root cause can remove downstream consequences. The contrast with RLKD’s structural approach illustrates complementary design choices. RLKD rewards trajectories whose overall step organization matches the teacher’s, while SCoRe surgically corrects the first point of divergence. The two are complementary, and combining global structural reward with local error correction remains unexplored.

Token-level Bayesian and meta-reflective credit. The credit-assignment problem becomes more pressing when the teacher signal is itself unstable, as in self-distillation under RLVR or in multi-turn agent rollouts where rewards arrive only at the end of long action sequences. AMR-SD (Wei et al., 2026) stabilizes RLVR self-distillation by separating the per-token signal from the policy update through an asymmetric meta-reflective self-teacher. The student rolls out, a meta-reflection stage aggregates verifier outcomes, peer rollouts, and reference feedback into per-token supervision, and an asymmetric loss keeps the student-teacher loop from diverging without external regularization. OPPO (Li et al., 2026k) formalizes the same per-token credit problem inside a Bayesian framework, recursively estimating per-token value through oracle-conditioned likelihood ratios. The likelihood-ratio recursion produces a per-token advantage that is consistent with the trajectory-level outcome, giving fine-grained reward attribution without a separate critic network. Where AMR-SD and OPPO operate on single-turn reasoning trajectories, StepOPSD (Zhang et al., 2026j) extends step-aware preference distillation to multi-turn agent RL, segmenting agent trajectories into

causal interaction steps and applying step-level teacher scoring jointly with token-level distillation. The dual-granularity alignment matches sparse environment rewards (which arrive at step boundaries) with dense per-token supervision (which fills the gap inside each step), addressing the credit-assignment mismatch that limits OPD on long-horizon agent tasks.

Self-evaluation as judge distillation. Most of the methods above treat the judge or reward model as an external signal source. Zhang et al. (2026h) ask instead whether a base model already carries a latent judge that can be elicited through on-policy distillation. The training cycle alternates between a calibration-coupled GRPO phase (where the policy is updated on a stratified round-robin sample over attribute-score bins, with a non-linear calibration exponent shaping the reward) and a masked judge distillation phase that distills the external judge’s actual scores back onto the self-evaluation tokens of the rolled-out trajectories, with the answer tokens left untouched. The masking confines the OPD update to the self-evaluation channel, treating the judge as an on-policy teacher specialized to that channel. The result frames judge calibration as an OPD problem that can run inside the same RL loop as the policy update, in contrast to off-policy reward-model fine-tuning.

Contrastive distillation within RL. A related line of work applies contrastive reasoning to the teacher signal itself, rather than to judge calibration. RLCS (Pan et al., 2026) integrates contrastive reverse KL into the GRPO training loop via a two-path loss. When the teacher is conditioned on a correct hint, the student is pulled toward the teacher’s distribution (an imitation objective). When the teacher is conditioned on a wrong hint, the student is pushed away (a contrast objective). This correct/wrong decomposition creates a directional gradient that reinforces correct reasoning strategies while actively suppressing misleading ones, without requiring a separate discriminator or preference pair construction. The approach naturally extends GRPO’s group-relative advantage with a per-token distillation signal whose sign is determined by the hint’s correctness rather than the answer’s outcome.

From signal matching to environment reconstruction. A further extension asks whether the student can recover the teacher’s *underlying optimization landscape*. $\mathcal{X}$ -KD (Cai & Yuan, 2026) adopts the AVRIL framework to jointly model the teacher’s latent reward function and perform policy distillation, encouraging the student to optimize in a reconstructed version of the teacher’s training environment instead of mimicking its outputs. TSD-KD (Kim & Baek, 2026) combines indirect distillation (student generates candidates that the teacher re-ranks via preference feedback) with direct distillation (KL matching applied selectively to high-entropy tokens where teacher-student gaps are largest). These methods move beyond matching the teacher’s *behavior* toward recovering its *incentives*.

Across objective types, a rough progression is visible. Fixed divergences deliver stable, well-understood optimization but tend to be bounded by the teacher’s capability ceiling. Adaptive divergences can improve efficiency by matching the local geometry of the loss landscape. RL-augmented objectives offer a path around the teacher ceiling at the cost of added variance and instability. The appropriate choice depends on the specific application. For instruction following where teacher quality is high, fixed or adaptive divergences are often sufficient. For reasoning tasks where the student is expected to exceed the teacher, RL augmentation becomes attractive. For tasks where the teacher’s training process is accessible, environment reconstruction offers a more complete form of knowledge transfer.

5 Signal Source and Teacher Architecture

With the objective function fixed (Section 4), the next design axis concerns *where the teacher signal comes from*. We organize signal sources along two levels. The first level distinguishes external teacher distillation, where a separate stronger model provides supervision, from self-distillation, where the model generates its own training signal by exploiting internal asymmetries. Within external teacher distillation, the second level separates white-box methods that access the teacher’s full output distribution from black-box methods that operate with only generated text or scalar scores. These two levels track distinct properties of the teacher, its identity (external versus the model itself) and, when external, the access level at which its signal can be observed. The two are not fully independent, since access

level in turn constrains which objectives are applicable, but keeping them on separate levels avoids conflating a deployment constraint with a methodological choice. Each combination imposes distinct constraints on the applicable objectives and achievable performance.

Historically, signal sources have evolved along both levels. The access-level dimension moved first. Early OPD methods (GKD, DistiLLM) assumed full white-box access to model internals, a requirement that limits distillation to same-organization deployments, while the rise of proprietary API-only models (GPT-4, Claude) necessitated black-box methods that extract knowledge from text outputs alone. The teacher-identity dimension shifted later, as self-distillation methods removed the external teacher entirely, permitting continuous improvement without access to any stronger model. The movement along the access-level dimension, from full access through limited access, and along the identity dimension, from an external teacher to none, both represent increasing autonomy at the cost of signal density, and the two trade-offs are largely independent.

Complex reasoning is among the settings where OPD tends to show pronounced gains over off-policy alternatives. The underlying reason is that reasoning is *path-dependent* (Agarwal et al., 2024). A proof that begins “assume for contradiction...” requires entirely different subsequent tokens than one beginning “we proceed by induction...” If the student has never generated the inductive approach during training, it is unlikely to recover at inference time when that path would be optimal. Off-policy training over a fixed dataset of teacher reasoning traces only covers a finite set of solution paths, while the student needs to navigate its own terrain of partial solutions, dead ends, and recovery strategies. This observation motivates the on-policy formulation for reasoning transfer. The student generates its own intermediate reasoning steps $r^S \sim P_\theta(\cdot|x)$, and the teacher evaluates these student-generated chains instead of dictating its own:

$$\mathcal{L}_{\text{CoT-OPD}}(\theta) = \mathbb{E}_{x \sim \mathcal{D}, r^S \sim P_\theta(\cdot|x)} \left[\sum_{t=1}^{|r^S|} D_{\text{KL}} \left(p_T(\cdot|x, r_{<t}^S) \parallel p_\theta(\cdot|x, r_{<t}^S) \right) \right]$$

The Forward KL direction encourages mode-covering, so that the student allocates probability wherever the teacher does on the student’s own prefixes, receiving corrective feedback on its actual generation patterns instead of memorizing the teacher’s preferred trajectories. This self-consistent learning dynamic, where the student practices its own reasoning style under expert supervision, has made on-policy CoT distillation a common choice for transferring multi-step reasoning to smaller models. The principle that rationale transfer outperforms answer-only supervision was first established in the off-policy setting by Hsieh et al. (2023), and subsequent work showed that generating rationales on-policy (Agarwal et al., 2024) can yield further gains by reducing exposure bias, an effect that has also been observed at scale in DeepSeek-R1 (DeepSeek-AI et al., 2025) and its successors.

5.1 External Teacher Distillation

External teacher distillation assumes access to a separate model whose capabilities exceed the student’s on the target task distribution, and the teacher’s role is to supply corrective signals on the student’s on-policy rollouts. The two parts below separate methods by the density of that signal. Full logit access (Section 5.1.1) yields the richest per-token supervision but requires co-hosting the teacher, while API-constrained settings (Section 5.1.2) trade signal density for deployment flexibility.

5.1.1 White-Box Logit Supervision

White-box access yields the densest possible signal, the full probability distribution across $|V|$ tokens at every decoding step (Table 6). This exposes the teacher’s implicit decision boundaries, including the relative probabilities among incorrect tokens (Hinton’s “dark knowledge”), which carry information about similarity structure that binary labels cannot capture.

Scope of this section. White-box logit access is the implicit assumption of the majority of methods cataloged elsewhere in this survey, including most fixed-divergence (Section 4.1),

Table 6: OPD methods: White-box logit supervision innovations.

Method	Year	Category	Divergence/Objective	Signal	Granularity	Key Innovation
DSKD (Zhang et al., 2025b)	2025	Signal	Dual-Space KL	White-box	Token	Cross-vocabulary projection
PromptKD (Kim et al., 2024)	2024	Signal	Prompt-based elicitation	White-box	Token	Input-side distillation
MAD-OPD (Wang et al., 2026c)	2026	Signal	Multi-agent debate ensemble	White-box	Token	Confidence-weighted multi-teacher
Veto (Jang et al., 2026)	2026	Signal	Geometric bridge KL	White-box	Token	Adaptive target reformulation
SimCT (Sun et al., 2026)	2026	Signal	Multi-token continuation KL	White-box	Token	Cross-tokenizer supervision recovery
CSD (Kim et al., 2026c)	2025	Signal	Continuous score matching	White-box	Token	Bypasses discrete vocabulary mismatch
MPD (Yang et al., 2026c)	2026	Signal	Mixed-policy KL	White-box	Token	Teacher rewrites student traces concisely
BRTS (Zhang et al., 2026e)	2026	Signal	Best-of-N teacher selection	White-box	Sequence	Priority rollout selection for teacher reliability
Pair-In, Pair-Out (Tan et al., 2026)	2026	Signal	MTP confidence-head KL	White-box (full backbone)	Token	OPD-supervised latent-MTP head
DuDi (Payoungkhamdee et al., 2026)	2026	Signal	Dual-signal cross-lingual KD	White-box	Token/Sequence	SPIN + token KD via cross-lingual verbalizer
OPRD (Yang et al., 2026d)	2026	Signal	Representation-space OPD	White-box	Layer	Hidden-state alignment extends logit-space distillation
BTB (Niu et al., 2026)	2026	Signal	Cross-tokenizer OPD	White-box	Token	Token-mapping bridges heterogeneous model families

adaptive-divergence (Section 4.2), and RL-augmented (Section 4.3) objectives, the privileged-information self-distillation methods in Section 5.2.1 that condition on full teacher logits, and most weighting and curriculum methods in Section 6. To avoid duplicating those methods under multiple categories, Section 5.1.1 (and the figure-1 leaf) covers the smaller subset whose *primary contribution* is an innovation on the white-box signal interface itself, specifically cross-tokenizer alignment, full-vocabulary compression, teacher-rollout selection, and confidence-head construction. Readers interested in the full set of white-box methods should consult the Signal column (“White-box”) across Tables 3–9.

Same-Family Distillation. When teacher and student share the same tokenizer (as within a single model family), white-box logit transfer is straightforward and methods can directly compute per-token KL divergence without alignment overhead.

Standard logit distillation. The majority of methods surveyed in Section 4 assume same-family white-box access and compute $p_T(\cdot|x, y_{<t})$ at each student-generated token position. The computational cost is dominated by the teacher forward pass, which can be amortized by caching teacher logits for an entire rollout before computing the student’s loss. When teacher and student share the same architecture, this caching reduces wallclock overhead to roughly $1.5\text{--}2\times$ per training step compared to student-only RL (Agarwal et al., 2024), a cost that has driven the systems-level optimizations discussed in Section 6.3.

Mixed-policy reasoning compression. MPD (Yang et al., 2026c) transfers concise reasoning behavior from a larger teacher to a smaller student through a mixed-policy mechanism. Given a student-sampled trajectory, the teacher rewrites it into a more concise reasoning trace, and the student trains via KL-based alignment on the compressed trajectory. This preserves student-policy exploration (unlike off-policy distillation) while injecting teacher-guided compression (unlike standard on-policy distillation, which aligns the student with teacher distributions over the student’s verbose trajectories). On Qwen3-1.7B, MPD reduces token usage by up to 27.1% while improving performance across multiple reasoning benchmarks, exploiting the observation that larger reasoning models produce more concise traces for the same problems.

Best-of-N teacher rollout selection. BRTS (Zhang et al., 2026e) addresses a variance problem in standard OPD. Teacher supervision computed under noisy student-generated contexts often relies on a single stochastic teacher rollout that can be incorrect or poorly matched to the student’s reasoning behavior. BRTS samples a small pool of teacher trajectories and selects the best using a priority rule (correctness first, student alignment second), with a ground-truth-conditioned recovery step for harder prompts where unconditioned teacher samples fail. The selected trajectory provides reliable teacher-context supervision inside the OPD loop. On AIME 2024, AIME 2025, and AMC 2023, BRTS improves over standard OPD with pronounced gains on harder datasets. Teacher signal quality matters more when problems are near the boundary of the student’s competence.

Token-level adaptive supervision. Within the white-box setting, several methods refine *which* tokens receive supervision, and they can be organized along the axis of what drives the selection decision. SelectKD (Huang et al., 2025) routes based on student-teacher agreement, applying a propose-and-verify mechanism where the student proposes tokens that the teacher verifies, and only positions where agreement holds receive full loss while rejected tokens are masked or down-weighted. AdaSwitch (Peng et al., 2025) routes based on student prefix reliability, maintaining a running error estimate and switching between

exploration (student generates freely) and guidance (teacher supplies dense supervision) so that guidance kicks in only when the trajectory has drifted far enough to warrant teacher intervention. These selection criteria, agreement and prefix quality, are additive rather than competing, each capturing a different signal about where supervision is most informative. They also compose naturally with entropy-based divergence routing (EOPD), which selects *which loss function* applies at each position while token-level selection determines *how much gradient signal* propagates. The combination yields an orthogonal two-dimensional design space that current methods have only begun to explore jointly. These token-selection mechanisms complement the sample-level and token-level weighting schemes discussed in Section 6.

Latent compression with OPD-supervised confidence heads. A separate use of white-box supervision combines OPD with multi-token prediction (MTP) for inference-time efficiency. Pair-In, Pair-Out (Tan et al., 2026) compresses the student’s input sequence into latent codes and replaces the autoregressive output head with a parallel MTP head, training a confidence head on top of the MTP outputs to decide which positions can commit speculatively. The full-backbone version of the same model serves as the OPD teacher, scoring compressed student rollouts and supplying the per-token signal that calibrates this confidence head, so that latent compression is constrained by the teacher’s own per-token preferences rather than by an unsupervised bottleneck. The construction frames OPD as a calibration mechanism for architectural compressions, complementing the per-token logit transfer methods above.

Cross-Family Distillation. The methods above assume that teacher and student share a common tokenizer, which holds when both belong to the same model family (e.g., Qwen3-30B→Qwen3-4B). A key challenge arises when teacher and student use different tokenizers or vocabulary sizes, as is common in cross-family transfer (e.g., Llama teacher→Qwen student). Direct KL divergence over mismatched vocabularies is undefined. DSKD (Zhang et al., 2025b) addresses this through *dual-space* knowledge distillation, employing projectors that map between teacher and student representation spaces:

$$\mathcal{L}_{\text{DSKD}} = D_{\text{KL}}(P_{T \rightarrow S} \parallel P_S) + D_{\text{KL}}(P_T \parallel P_{S \rightarrow T})$$

where $P_{T \rightarrow S}$ denotes the teacher’s hidden states projected into the student’s space and decoded through the student’s prediction head. An exact token alignment algorithm handles vocabulary mismatches by identifying identical tokens across tokenizers. Vocabulary tokens with similar semantic embeddings should receive similar probability mass even when the tokenization schemes differ, allowing distillation across different architectures (e.g., Llama teacher to Qwen student) that share no vocabulary overlap. CSD (Kim et al., 2026c) offers another solution, distilling via a continuous score function that bypasses the discrete vocabulary mismatch.

SimCT (Sun et al., 2026) takes a different approach that operates entirely within the OPD loss form itself. Rather than projecting representations or transporting probability mass, SimCT enlarges the supervision space by comparing teacher and student over short multi-token continuation units that both tokenizers can realize. These units represent the finest jointly tokenizable supervision interface, recovering teacher signal that exact shared-token matching silently discards at positions where vocabularies disagree. Across three heterogeneous teacher-student pairs (Qwen2.5-7B→Phi-4-mini, Qwen2.5-7B→Gemma-2-2B, Phi-4-mini→Gemma-2-2B) on math reasoning and code generation, SimCT outperforms both shared-vocabulary OPD and coarser cross-tokenizer baselines, with pronounced gains at positions of maximal vocabulary disagreement. The result suggests that a substantial fraction of the cross-tokenizer supervision loss documented under shared-token matching stems not from inherent incompatibility but from an unnecessarily restrictive matching granularity.

Together, these methods move OPD toward being more architecture-agnostic, allowing organizations to distill from teachers in one model family (e.g., Llama-based) into students in another (e.g., Qwen-based) without the historically restrictive requirement of shared vocabulary, and supporting heterogeneous teacher-student pairs that were previously difficult to handle.

Table 7: OPD methods: Black-box and API-constrained distillation.

Method	Year	Category	Divergence/Objective	Signal	Granularity	Key Innovation
GAD (Ye et al., 2025)	2025	Signal	Adversarial matching	Black-box	Sequence	Discriminator as proxy reward
Lion (Jiang et al., 2023)	2023	Signal	Verbal feedback	Black-box	Sequence	NL critique curriculum
OVD (Xiong et al., 2026)	2026	Signal	Verbal scores (0-9)	Black-box	Sequence	Score-based trajectory matching
LUFFY (Yan et al., 2025)	2025	Signal	Mixed-policy GRPO	Off-policy + RL	Sequence	Importance-weighted off-policy
PRISM (Wang et al., 2026g)	2026	Signal	Black-box pre-alignment	Black-box (API)	Sequence	Multimodal OPD before RLVR
ROPD (Fang et al., 2026)	2026	Signal	Rubric-based scoring	Black-box (API)	Sequence	Structured semantic rubrics, 10 $\times$ sample efficiency
DASD (Yan et al., 2026)	2026	Signal	Sequence-level distribution match	Black-box	Sequence	Multi-trace distribution alignment
DOT (Zhang et al., 2026f)	2026	Signal	Theoretical RL+KD framework	Black-box	Sequence	On-policy generalization theory
ORPO-Distill (Singh et al., 2025)	2025	Signal	ORPO contrastive preference	Black-box (API)	Sequence	Mixed-policy negative sampling
OmnioPD (Zhou et al., 2026b)	2026	Signal	Chunk MC verification + KL anchor	Black-box (API)	Chunk/Token	Logit-free Bayesian smoothing

A separate challenge in cross-family transfer is teacher *overconfidence*. When the teacher assigns extreme probability to tokens outside the student’s reachable space, the resulting optimization targets are simultaneously precise and useless. Two solutions tackle this from opposite directions. Veto (Jang et al., 2026) (Adaptive Target Reformulation) refines the teacher’s output *post-hoc*, constructing a geometric bridge in logit space with a parameter β that serves as both an Adaptive Gradient Veto (suppressing pathological gradients on low-confidence tokens) and a Decisiveness Knob (balancing performance with diversity), dynamically reformulating the teacher’s target distribution based on the student’s current capability. PromptKD (Kim et al., 2024) takes the opposite approach, adapting the teacher *pre-hoc* by prepending learnable prompt tokens (adding only 0.0007% of the teacher’s parameters) that cause the teacher to produce “student-friendly” distributions more accessible to the student’s limited capacity. Post-hoc methods (Veto) operate on the supervisory signal after it is produced, while pre-hoc methods (PromptKD) modify the teacher’s generation process itself. Both are effective, and they compose without interference.

Cross-lingual logit transfer via shared verbalizers. The cross-tokenizer methods above operate on monolingual reasoning data. DuDi (Payoungkhamdee et al., 2026) adapts the same dual-space principle to the multilingual setting for Southeast Asian small language models. A cross-lingual verbalizer maps semantically aligned content across teacher-student vocabularies, allowing token-level distillation to operate over a shared verbalizer space even when the surface tokenizations differ across languages. The training signal is dual-component, an online sequence-level SPIN-style objective that improves the student’s preferred-language responses and a token-level KD term that mixes off-policy teacher data with on-policy student rollouts. The on-policy stream supplies corrective per-token feedback on the student’s self-generated mistakes in lower-resource languages, complementing off-policy demonstrations on languages with abundant teacher data. The construction shows that the SimCT idea of widening the matching granularity also applies across languages once a shared verbalizer is available.

Representation-space and cross-tokenizer extensions. Two further methods push white-box OPD beyond the logit level. OPRD (Yang et al., 2026d) extends OPD into the representation space by adding a hidden-state alignment term alongside standard logit distillation. On-policy student rollouts trigger both the usual per-token KL loss and a layer-wise representation matching objective, so that the teacher’s internal activations, which encode intermediate reasoning structure, are transferred in addition to its output distribution. BTB (Niu et al., 2026) addresses heterogeneous model families by constructing an explicit token-to-token mapping bridge across incompatible tokenizers. Rather than relying on shared vocabulary subsets or projection-based alignment (as in DSKD and SimCT), BTB learns a lightweight mapping from each teacher token to the closest student-side token and applies standard on-policy KD over this bridged vocabulary. The design extends OPD to cross-family pairs such as Llama-to-Qwen that share no tokenizer overlap, complementing SimCT’s multi-token continuation approach with a direct token-level mapping.

5.1.2 Black-Box and API-Constrained Distillation

When only API access is available, the student observes the teacher’s top-1 text output (or at best, top- k log-probabilities) rather than the full distribution (Table 7). This information bottleneck rules out token-level divergence matching and forces methods to operate at the sequence level. Despite this limitation, black-box methods have been reported to work well

empirically. In the settings studied, a substantial portion of the teacher’s knowledge can be recovered from text-only observations.

Distilling Step-by-Step. Hsieh et al. (2023) introduced the foundational black-box approach, extracting not just the teacher’s final answers but its intermediate reasoning steps (rationales) as additional supervision. Although technically off-policy (the student fine-tunes on static teacher-generated rationales rather than its own), this work is the conceptual ancestor of on-policy CoT distillation and established a key principle. The student is trained with a multi-task objective combining answer prediction and rationale generation, allowing a 770M-parameter model to outperform a 540B-parameter teacher with over $700\times$ fewer parameters and substantially less training data (50% fewer examples on average, up to 85% reduction on some tasks). This result suggests that *how* the teacher reasons can be more valuable than *what* it concludes, a principle that recurs throughout modern OPD and motivates the on-policy extensions below.

If extracting textual rationales recovers partial teacher reasoning, a further question is whether a richer proxy for the teacher’s full distribution can be reconstructed from black-box outputs alone. The following methods attempt increasingly aggressive reconstructions, from adversarial discrimination through verbal scoring to preference ranking.

Adversarial distribution matching. GAD (Ye et al., 2025) reformulates the distillation problem as a generative adversarial game:

$$\max_G \min_D V(G, D) = \mathbb{E}_{(x, y^T) \sim \mathcal{T}} \left[-\log \sigma \left(D(y^T) - D(G(x)) \right) \right]$$

A discriminator D (initialized from the student with a scalar prediction head) distinguishes student rollouts from teacher API outputs via a Bradley-Terry preference model, while the generator G maximizes V via policy gradient (implemented as GRPO) using the discriminator score as an evolving on-policy reward signal. Unlike white-box methods that match full distributions, GAD extracts knowledge through a scalar quality signal, trading distributional fidelity for API compatibility. It inherits the training instability of adversarial objectives, as the discriminator can overfit to stylistic cues rather than semantic quality, but produces reliable gains over SFT baselines.

Verbal feedback and curriculum. Lion (Jiang et al., 2023) exploits the teacher’s capacity as a natural language critic through a three-stage adversarial loop comprising (i) an *imitation* stage where the student is fine-tuned on teacher-generated responses, (ii) a *discrimination* stage where the teacher identifies instructions on which the student still falls short, and (iii) a *generation* stage where the teacher creates harder instructions targeting the student’s identified weaknesses. This closed-loop curriculum steadily narrows the teacher-student gap without requiring any gradient information from the teacher, making it applicable to any instruction-following API. Lion-13B reaches competitive performance on BIG-Bench Hard and AGIEval using only 70K training examples.

Distribution-Aligned Sequence Distillation (DASD) (Yan et al., 2026) reexamines the dominant practice of SFT on teacher-generated reasoning traces, identifying three limitations, namely (i) inadequate representation of the teacher’s sequence-level distribution from single best-of- n samples, (ii) misalignment between teacher output distribution and student capacity, and (iii) exposure bias from teacher-forced training. Its distribution-aligned pipeline addresses all three by generating multiple diverse teacher traces per prompt and aligning the student’s generation distribution with the teacher’s through sequence-level matching, reaching strong performance using only 448K training samples. DDT (Zhang et al., 2026f) complements DASD’s empirical findings with a theoretical framework explaining why RL generalizes better than SFT for reasoning. RL’s on-policy data generation keeps training aligned with the model’s evolving distribution instead of a static snapshot.

Verbal score distillation. OVD (Xiong et al., 2026) replaces token-level probability matching with trajectory matching using discrete verbal scores (0–9) from teacher models. This removes the requirement for token-level alignment between vocabularies, reduces memory consumption by avoiding logit storage, and reports up to +12.9% absolute EM improvement on web QA and +25.7% on math benchmarks in a single-sample training regime. The effectiveness of such coarse supervision suggests that on-policy exploration itself supplies a

substantial fraction of the learning signal, with the teacher’s role being closer to *selecting among student trajectories* than to *correcting individual tokens*.

These methods form a spectrum of signal richness versus access requirements. Adversarial methods recover a dense proxy for the teacher’s distribution but require complex training dynamics (discriminator co-training, generator-discriminator instability). Score-based methods are simpler to train but lose distributional information, recovering only a scalar quality signal. The signal type therefore maps naturally to API access tier, with full logit access enabling KL methods (Section 4) and unavailable token-level logits motivating adversarial or score-based approaches.

MoE discriminator for pipeline-staged OPD. PRISM (Wang et al., 2026g) contributes a distinct architectural innovation to black-box distillation by inserting an OPD-based *pre-alignment* stage between SFT and RLVR, tackling the observation that direct RLVR from SFT checkpoints suffers from cold-start instability due to the large distributional gap between SFT outputs and reward-aligned behavior. Rather than using a monolithic discriminator (as in GAD), PRISM employs a Mixture-of-Experts discriminator with specialized perception and reasoning expert heads, delivering *disentangled* corrective signals that distinguish visual grounding errors from logical reasoning failures. The adversarial game operates entirely at the response level without requiring teacher logits. The student generates on-policy responses, a black-box teacher (Gemini 3 Flash) generates reference demonstrations, and the MoE discriminator learns to distinguish them while delivering differentiated feedback through its expert routing. This disentangled feedback matches the structure of multimodal models, where failure modes are heterogeneous, since a model may correctly perceive an image but reason incorrectly about it, or vice versa. On Qwen3-VL at 4B and 8B scales, the pre-alignment stage alone yields +4.4 and +6.0 points over direct SFT-to-RLVR pipelines, and the subsequent RLVR training converges faster and more stably. Black-box OPD therefore serves as a structured warm-up that shapes the policy landscape for downstream RL optimization.

Rubric-based semantic distillation. ROPD (Fang et al., 2026) challenges the assumption that effective OPD requires token-level logit access by proposing structured semantic rubrics as a scalable alternative to teacher logits. A Rubricator induces prompt-specific evaluation rubrics by contrasting teacher and student rollouts, and a Verifier scores student rollouts against these rubrics to drive on-policy optimization. The teacher model assumes both roles, and its independence from the training loop enables offline rubric generation that lowers GPU requirements considerably. Across diverse benchmarks (AIME 24/25, HMMT 25, GPQA-Diamond, HealthBench, IFEval) and configurations (Qwen3-4B and Gemma3-4B students with GPT-5.2 and Qwen3-30B teachers), ROPD matches or surpasses logit-based OPD methods while achieving approximately $10\times$ sample efficiency. On AIME 2025 with thinking mode, the Qwen3-4B student (68.75%) surpasses its GPT-5.2 teacher (67.08%), rubric-augmented optimization can, in favorable configurations, support student-exceeds-teacher performance even without gradient-level access. ROPD occupies a distinctive position in the black-box spectrum. Where GAD recovers a scalar quality signal through adversarial discrimination and OVD through verbal scores, ROPD recovers *structured multi-dimensional* evaluation criteria that decompose quality into interpretable rubric dimensions, offering richer feedback than scores while remaining fully black-box compatible.

Preference distillation across architectures. ORPO-Distill (Singh et al., 2025) recasts cross-architecture black-box distillation as a preference-optimization problem with mixed-policy negatives. Diverse teacher CoT traces (sampled with $K=8$, temperature 0.8, deduplicated by ROUGE-L) serve as positives, and student-generated incorrect traces serve as negatives, with the ORPO contrastive loss ($\lambda=1$) replacing token-level KL. A mixed-policy sampler with $\phi=0.5$ alternates between negatives drawn from the base checkpoint θ_0 and from the latest checkpoint θ_{t-1} , balancing exposure-bias mitigation against the stability of a fixed reference policy. Across five multi-choice QA benchmarks the construction outperforms both pure off-policy and pure on-policy baselines, and the distillation transfers across architectures (e.g., Llama-family teacher into Qwen-family student) without logit access on either side.

Chunk-level speculative verification. OmniOPD (Zhou et al., 2026b) narrows the gap between black-box and white-box OPD without requiring teacher logits. A peak-entropy scheduler picks chunk boundaries inside the student’s rollout, the black-box teacher is queried for N Monte Carlo continuations conditioned on the same prefix, and a Dirichlet-Multinomial Bayesian smoother converts chunk-level semantic similarity into a per-chunk reward distribution. Tokens inside an audited chunk receive supervision from the verified teacher distribution, while a trust-region KL anchor on the unaudited tokens keeps the off-chunk gradient from drifting. The method recovers a substantial fraction of white-box OPD’s performance using only black-box queries, chunk-level Monte Carlo verification can substitute for the dense logit signal when teacher access is API-only.

Off-policy knowledge integration. Not all training signal must originate on-policy. Incorporating knowledge from off-policy sources (e.g., pre-existing teacher-generated datasets) into an on-policy framework presents its own difficulties. LUFFY (Yan et al., 2025) handles this by extending GRPO to a mixed-policy objective that incorporates off-policy reasoning traces alongside the student’s on-policy rollouts. The key design challenge is mitigating *superficial imitation*, because without correction the student tends to copy high-probability tokens from off-policy traces while ignoring low-probability but important reasoning steps. LUFFY’s policy shaping via regularized importance sampling reweights gradients to emphasize learning from unfamiliar yet effective actions, posting +6.4 average points over standard RLVR.

Across the two access levels, external teacher distillation offers the densest supervision in OPD at the cost of access-dependent constraints. White-box methods reach the tightest distributional alignment through full-vocabulary KL but require co-hosting the teacher, while black-box methods trade signal density for deployment flexibility. The unresolved tension is that the strongest teachers, the proprietary frontier models, are typically reachable only through APIs, which is the setting where the dense per-token signal that OPD relies on is hardest to obtain. The self-distillation methods in the next section remove the external teacher entirely, trading this dependence for a different limitation that we examine below.

5.2 Self-Distillation

Self-distillation removes the need for an external teacher entirely (Table 8). Instead, the model exploits asymmetries *within itself* to construct a training signal. Three families of methods have emerged, distinguished by the source of this internal asymmetry, namely privileged information that is available at training but not inference time, pure self-distillation through rollout diversity and architectural self-training, and external feedback from verifiers or environments that anchors the self-generated signal.

Scope boundary. This section covers methods whose training signal arises from distributional alignment between different views of the same model (conditioned vs. unconditioned, frozen vs. updated, temperature-sampled vs. greedy). We exclude game-theoretic self-play methods such as SPIN (Chen et al., 2024) and IRIS (Liao et al., 2026), whose core mechanism is adversarial competition between successive checkpoints aimed at convergence to a Nash equilibrium over a reference distribution. That framing is conceptually distinct from the distributional alignment that defines on-policy distillation. These game-theoretic approaches are briefly discussed as an adjacent paradigm in Section 9.

5.2.1 Privileged Information

The first family creates asymmetry through *privileged information* (PI), meaning training-time access to context unavailable at inference. Conditioning the same model on additional context produces a stronger policy that can supervise the unconditioned version, turning a single network into both teacher and student. The methods below differ mostly in what plays the role of that context, ranging from a ground-truth answer through retrieved documents, spatial masks, structured judge feedback, and future observations to a weak external critique. This surface variety is largely orthogonal to the distillation mechanism itself, which remains token-level matching of a PI-conditioned self-teacher over student rollouts. What does vary in consequence is the *type* of privileged signal. A shared latent

Table 8: OPD methods: Self-distillation.

Method	Year	Category	Divergence/Objective	Signal	Granularity	Key Innovation
OPSD (Zhao et al., 2026b)	2026	Self (PI)	KL on GT-conditioned	Self	Token	Ground-truth as privileged info
CRSP (Sang et al., 2026)	2026	Self (PI)	Consensus PI	Self	Token	37% token reduction, +9% accuracy
GATES (Stein et al., 2026)	2026	Self (PI)	Gated consensus KL	Self	Token	Document as PI, gated supervision
OPSDL (Zhang et al., 2026g)	2026	Self (PI)	Short-ctx $\rightarrow$ long-ctx	Self	Token	Long-context self-distillation
SD-ZERO (He et al., 2026)	2026	Self (EP)	Self-revision preference	Self + Verifier	Sequence	Binary reward $\rightarrow$ dense self-supervision
π -Play (Zhang et al., 2026d)	2026	Self (EP)	Multi-agent co-evolution	Token	Token	Internally generated PI from co-evolution
UnSD (Jin et al., 2026c)	2026	Self (SD)	Multi-axis unified SD	Self (EMA + agreement)	Token	Systematic component analysis
RLRT (Kim et al., 2026a)	2026	Self (EP)	Reversed teacher signal	Self (reversed)	Token	Reinforcing self-driven reasoning tokens
SDPO (Hibotter et al., 2026)	2026	Self (EP)	Textual feedback + DPO	Self + Verifier	Sequence	Credit assignment from outcomes
RLTF (Song et al., 2026)	2026	Self (EP)	Critique-conditioned	Self + Verifier	Sequence	Text feedback for self-distillation
SRPO (Li et al., 2026c)	2026	Self (EP)	CRPO + SDPO routing	Self + Verifier	Sequence	Sample-routed multi-objective
CoPD (Gu et al., 2026)	2026	Self (EP)	Co-evolving bidirectional OPD	Mutual (parallel)	Sequence	Multi-capability consolidation
VPD (Li et al., 2026a)	2026	Self (EP)	Variational EM co-evolution	Self (co-evolved)	Token	Trust-region teacher adaptation from language feedback
PAINT (Tan & Hong, 2026a)	2026	Self (EP)	Partial-solution masking	Self (re-score)	Token	Overlap-adaptive self-distillation
RESL (Zhang et al., 2026f)	2026	Self (EP)	Reflection-enhanced self-distill	Self + Env. Feedback	Token	Failure reflection + persistent playbook
MSD (Qin et al., 2026)	2026	Self (PI)	Cross-lingual safety distill	Self (English PI)	Token	Multilingual safety alignment
GUI-SD (Zhang et al., 2026i)	2026	Self (PI)	Visual spatial PI for GUI	Self (Gaussian mask)	Token	Multimodal GUI grounding
PBSD (Yu et al., 2026c)	2026	Self (PI)	Pref.-based reward-reg. KL	Self (ctx-augmented)	Sequence	DPO-style loss, preference margin
TT-OPD (Jeong, 2026)	2026	Self (PI)	Turn-level truncated KL	Self (EMA + outcome hints)	Turn	Multi-turn agentic OPD
VSD (Lin et al., 2026a)	2026	Self (PI)	Direction-magnitude decouple	Self (structured judge)	Token	Structured PI for video reasoning
π -Distill (Peralta et al., 2026)	2026	Self (PI)	Unified PI framework	Self	Token	Theoretical foundations for PI self-distill
OPCD (Ye et al., 2026b)	2026	Self (PI)	Context-conditioned RKL	Self (experience PI)	Token	Experiential knowledge internalization
OEL (Ye et al., 2026a)	2026	Self (PI)	Online experiential learning	Self (trajectory PI)	Token	Continuous post-deployment distillation loop
HDPO (Ding, 2026)	2026	Self (PI)	GRPO + JSD self-distill	Self (GT-conditioned)	Sequence	Cliff-prompt non-zero gradient patch
MTP-SD (Kirichenbauer et al., 2026)	2026	Self (SD)	Multi-token span scoring	Self (frozen copy)	Span	Architectural self-distillation for > 3x decoding
SDFT (Shenfeld et al., 2026)	2026	Self (SD)	In-trust-region demo teacher RKL	Self (ICL-conditioned)	Token	Forgetting resistance via trust-region implicit KL
RLSD (Yang et al., 2026a)	2026	Self (EP)	RLVR + self-distill decomp.	Self + Verifier	Token	Magnitude (self-distill) + direction (RLVR), 2x efficiency
SDAR (Liu et al., 2026)	2026	Self (EP)	Self-distill OPD + RL	Self (GT-conditioned)	Token	Signpost-gated auxiliary self-distill for agents
OPHSD (Zhao et al., 2026c)	2026	Self (PI)	Harness-augmented self-KL	Self (harness PI)	Token	Internalizes inference scaffolds as PI
COPSD (Liu et al., 2026h)	2026	Self (PI)	Cross-lingual token KL	Self (English PI)	Token	17 African languages via cross-lingual PI
ATESD (Han et al., 2026b)	2026	Self (PI)	Beta-policy exposure control	Self (adaptive PI)	Token	Learnable teacher exposure ratio
TRACE (Wang et al., 2026e)	2026	Self (PI)	Token-routed FKL/RKL	Self (annotator PI)	Span	Critical-span OPD + GRPO on rest
LisOPD (Li et al., 2026h)	2026	Theory	Reward-extrapolation clip safety	White-box	Token	Closed-form λ^* for format collapse
Rock Tokens (Jiang et al., 2026b)	2026	Theory	Persistent high-loss analysis	White-box	Token	18% structurally uninformative tokens
Many Faces (Zhu et al., 2026)	2026	Theory	Comprehensive failure catalog	White-box/Self	Token	Biased top-K KL + PI type limitation
Unmasking OPD (Armandpour et al., 2026)	2026	Theory	Per-token gradient alignment	White-box/Self	Token	Training-free diagnostic framework
Local Collapse (Liu et al., 2026b)	2026	Theory	BIC-style release rule	White-box	Segment	Local teachability collapse in suffix OPD
Safactory (Chen et al., 2026d)	2026	Application	Safety evolution pipeline	White-box	Trajectory	Closed-loop OPD for agent safety
HyperEyes (Li et al., 2026c)	2026	Application	Duplication-efficiency RKL	White-box	Token/Seq	Multi-modal search agent with OPD
ProteinOPD (Zhang et al., 2026m)	2026	Application	Multi-teacher consensus OPD	White-box (multi)	Token	Protein PLM preference alignment
TTS (Kumar et al., 2026)	2026	Application	Online test-time OPD	White-box	Token	Draft-model adaptation for speculative decoding
Dagger-LM (Li et al., 2026a)	2026	Application	Turn-level student-teacher mix	White-box	Turn	Dagger for multi-turn LLM agents
AVSD (Nguyen et al., 2026)	2026	Self (PI)	Multi-view consensus + residual	Self (multi-view)	Token	Decompose consensus vs view-specific
Skill-Cond. Gated SD (Huang et al., 2026)	2026	Self (PI)	Outcome-validated polarity	Self (multi-teacher)	Token	Outcome-gated bounded SD
Critique Distill (Jin et al., 2026a)	2026	Self (PI)	Weak-critic conditioned KL	Self (weak critic)	Token	Scalable oversight inversion
Constitutional Cross-SFT (Wen et al., 2026)	2026	Self (PI)	Cross-SFT cold-start + RKL	Self (constitution)	Token	Safety-expressiveness preservation
World-Model PI (Zhou et al., 2026a)	2026	Self (PI)	Advantage-weighted decision-node KL	Self (future-video PI)	Step	Verify/invoke/rely on world-model rollouts
SD-PC (Liu et al., 2026g)	2026	Self (PI)	Full-vocab RKL + REINFORCE	Self	Token	Full-vocab SD with positive-advantage gating
VPG (Chen et al., 2026e)	2026	Self (SD)	Length-aware variational posterior	Self	Sequence	Efficiency-aware reasoning self-distill
Self-Sup OPD (Tan & Hong, 2026b)	2026	Self (SD)	Intra-group KL	Self (GRPO group)	Token	Correct/incorrect-rollout contrast
HINT-SD (Yeo et al., 2026)	2026	Self (SD)	Targeted hindsight relabel	Self	Turn	Failure-relevant turn selection
SD-Search (Ma et al., 2026)	2026	Self (SD)	Step hindsight relabel	Self	Step	Search-query-level supervision
Vision-OPD (Yuan et al., 2026)	2026	Self (SD)	Regional-to-global RKL	Self (regional PI)	Token	Fine-grained MLLM grounding
It Takes Two (Park et al., 2026)	2026	Self (SD)	Complementary parallel SD	Self (mutual splits)	Token	Contextual-integrity privacy alignment
TOD Proactivity (Zhang et al., 2026a)	2026	Self (SD)	Asymmetric user-concern PI	Self (concerns)	Token	Proactive task-oriented dialogue
Search-EI (Liang et al., 2026b)	2026	Self (SD)	GRPO + offline SD evolution	Self (snapshot)	Sequence	Self-evolution loop for search reasoning
MAIGO (Zheng et al., 2026b)	2026	Self (SD)	History-cleaned RKL	Self (single-turn PI)	Token	Multi-turn lost-in-conversation gap
RQSD (Zhao et al., 2026d)	2026	Self (SD)	Reflection + quote-localized	Self (reflector)	Token/Span	Error-source targeted SD
Canonical-Ctx OPD (Lin et al., 2026b)	2026	Self (SD)	Answer-masked RKL	Self (full-context PI)	Token	Multi-turn canonical-context alignment
COMAP (Liu et al., 2026i)	2026	Self (SD)	EMA SD + reflection gate	Self (world model)	Trajectory	Co-evolve agent and world model
OPCT (Han et al., 2026a)	2026	Self (EP)	Frozen-copy contrastive RKL	Self (frozen-copy)	Token	Sycophancy/jailbreak consistency
MCSD (Luo et al., 2026b)	2026	Self (PI)	Modality-gap-aware two-stage SD	Self (modality PI)	Token	Spatial planning with modality gap bridging
PTD-PO (Xiang et al., 2026)	2026	Self (PI)	Top-K JSD privileged tutoring	Self (spatial-reasoning PI)	Token	Spatial-reasoning hints for multimodal policy optimization
TWI (Cai et al., 2026a)	2026	Self (PI)	Cropped-image zoom-in PI	Self (cropped PI)	Token	Privileged zoom-in reasoning internalizes visual grounding
PBSD (Tan et al., 2026)	2026	Self (PI)	Bayesian turn-level SD	Self (trajectory reward)	Turn	Sparsified trajectory rewards - turn-level credit
Visual-SDPO (Dong, 2026)	2026	Self (PI)	Statement-weighted KL + GRPO	Self (visual artifact PI)	Token	Rendered visual artifact as privileged feedback
AR-OPD (Zhang, 2026)	2026	Self (PI)	Anchor+oracle residual OPD	Self (oracle PI)	Token	Prevents hindsight leakage; +2.3 vs. full OPD
HERO (Liu et al., 2026a)	2026	Self (PI)	Turn-level KL with future-obs PI	Self (future-env PI)	Turn	Future environment observations as privileged context
When Ctx Returns (Wang et al., 2026h)	2026	Self (PI)	PKL no-context anchoring	Self (no-ctx PI)	Token	Anchoring regularizer prevents context-induced degradation
Rubric-Guided SD (Rezaei et al., 2026)	2026	Self (PI)	JSD with rubric PI	Self (rubric PI)	Token	Removes need for external verifier via rubric-conditioned SD
AR-to-Diffusion (Su et al., 2026)	2026	Self (SD)	AR $\rightarrow$ diffusion on-policy SD	Self (frozen AR)	Token	Converts AR LM to diffusion LM; 15x $\sim$ 7000x fewer training tokens
SC-OPD (Xu et al., 2026a)	2026	Self (EP)	Sign-consistency gated SD	Self + Verifier	Token	Phased teacher sampling; +1.98 avg@32 / +7.50 pass@632
OPCoD (Byeon et al., 2026)	2026	Self (EP)	Coupled co-distillation + peer feedback	Mutual (peer)	Token	Cross-domain mutual Pareto improvement

rule that applies across instances (a system prompt, a constitution, a grading rubric) can be internalized into a PI-free policy, whereas an instance-specific answer cannot be aggregated the same way, a distinction that bounds when PI distillation succeeds and that we return to in the failure analysis of Section 7.2. We therefore present the methods in order of increasing distance from the original ground-truth-as-PI construction, noting for each what privileged signal it introduces and which limitation it addresses.

Ground-truth as PI. OPSD (Zhao et al., 2026b) constructs the teacher policy by conditioning on both the problem x and the ground-truth answer y^* , while the student conditions only on x . The student generates on-policy rollouts, and the teacher delivers dense token-level supervision.

$$\mathcal{L}_{\text{OPSD}}(\theta) = \mathbb{E}_{(x, y^*) \sim \mathcal{S}} \mathbb{E}_{\hat{y} \sim p_{\theta}(\cdot | x)} \left[\sum_{n=1}^{|y|} D(p_{\theta}(\cdot | x, y^*, \hat{y}_{<n}) \parallel p_{\theta}(\cdot | x, \hat{y}_{<n})) \right]$$

On competition-level math benchmarks (AIME 2024/2025, HMMT), OPSD matches or exceeds GRPO at the 4B and 8B scales while using only a single rollout per problem versus GRPO’s eight, yielding a significant computational advantage. The gains are largest at 8B (+0.9 average points over GRPO) and 4B (+0.8 over GRPO), but it underperforms GRPO at 1.7B (-0.5), likely because OPSD’s reliance on self-rationalization requires sufficient model capacity to produce a meaningful dense token-level signal. The authors note a theoretical limitation, that when problems exceed a model’s comprehension threshold, even the ground-truth-conditioned teacher cannot supply meaningful supervision, evidence of a problem-difficulty ceiling in the PI mechanism that scales with model capacity.

A separate limitation arises when all rollouts for a given prompt fail (the “cliff prompt” problem), causing the policy gradient to vanish entirely. This failure mode directly par-

allels PACED’s finding (Section 6.2) that gradient SNR vanishes when student pass rate approaches zero, but with an important difference. PACED tackles this through curriculum selection (avoiding problematic prompts), while HDPO (Ding, 2026) tackles it through signal augmentation. HDPO augments GRPO with a JSD-based self-distillation term on ground-truth-conditioned rollouts, providing non-zero gradients even at the boundary of the student’s competence, directly patching the cliff-prompt vulnerability without restricting the problem distribution.

Relaxing the ground-truth requirement. Ground-truth labels are not always available. GATES (Stein et al., 2026) shows that diverse forms of training-time context can serve as PI by having a single model act as tutor (conditioned on a source document) and student (answering from the question alone). Because source documents do not guarantee correct answers, GATES presents a consensus-based gating mechanism that suppresses the gradient when tutor uncertainty is high, preventing the model from reinforcing unreliable self-supervision. π -Distill (Penaloza et al., 2026), an earlier formalization, casts the OPSD-style privileged self-distillation paradigm as a unified theoretical framework, identifying the conditions under which privileged self-distillation provably improves over standard training. RESD (Zhang et al., 2026n) extends the PI paradigm to settings where successful rollouts are extremely rare (below 5% pass rate). Rather than discarding failed trajectories, RESD treats them as a form of negative privileged information, where a reflection module distills structured diagnostic feedback from failures into a persistent playbook that conditions subsequent on-policy generation, providing dense supervision that outperforms GRPO by approximately $8\times$ in sample efficiency on low-pass-rate tasks.

Behavioral PI: context distillation. The PI paradigm extends beyond factual ground truth to behavioral instructions. OPCD (Ye et al., 2026b) treats accumulated experiential knowledge (historical solution traces) and optimized system prompts as privileged context, training the student on its own rollouts while minimizing reverse KL against a context-conditioned teacher, allowing a model to internalize behaviors that would otherwise require expensive in-context demonstrations at inference time. OPSDL (Zhang et al., 2026g) extends this to long-context settings, where the model’s own short-context capability serves as the self-teacher to supervise its on-policy long-context generation via token-level reverse KL, enhancing long-context capabilities without relying on high-quality supervision data or sparse sequence-level rewards. OEL (Ye et al., 2026a) extends to continuous post-deployment learning, extracting transferable experiential knowledge from interaction trajectories as dynamic PI and iterating knowledge extraction with on-policy consolidation to form an online learning loop, closing a loop that OPCD opened.

Reasoning compression via PI. CRISP (Sang et al., 2026) finds that a model prompted to “be concise” can serve as a teacher for its own verbose version, reducing chain-of-thought token count by 57–59% on MATH-500 while *improving* accuracy by 9–16 percentage points. Much of the verbosity in distilled reasoning models is trainable inefficiency rather than a necessary feature of competent reasoning.

Visual PI for GUI grounding. GUI-SD (Zhang et al., 2026i) extends the privileged information paradigm to multimodal GUI grounding, where the task is to predict click coordinates on a screen given a natural language instruction. The privileged context is a Gaussian soft mask centered on the target element’s bounding box, giving the teacher version explicit spatial localization that the student must learn to infer from raw screenshots alone. This represents a qualitatively different form of PI from textual ground truth (OPSD) or document context (GATES), one that is *visual* and *spatial*, evidence that the PI mechanism extends across modalities. Beyond the new PI type, GUI-SD proposes entropy-guided token weighting via a dual-criterion, where tokens receive higher loss weight based on both their digit significance (coordinate digits carry more semantic weight than filler tokens) and the teacher’s confidence at that position. With only a single on-policy rollout per instance, GUI-SD outperforms GRPO (which requires multiple rollouts for reward estimation) on six GUI grounding benchmarks including ScreenSpot-v2, ScreenSpot-Pro, and OSWorld. The success of single-rollout PI distillation in the visual domain mirrors OPSD’s finding in mathematical reasoning, offering further evidence that privileged information can deliver sufficiently dense gradient signal to reduce reliance on the sample-intensive reward estimation required by RL-based alternatives.

Internalizing inference harnesses as PI. OPHSD (Zhao et al., 2026c) extends the PI paradigm from factual ground truth to *inference-time scaffolds*. When LLMs are paired with external harnesses (draft-verify loops, plan-solve workflows), the gain comes from the procedure rather than the model. OPHSD uses the harness-augmented model as teacher for self-distillation, internalizing task-specific harness capabilities into the student. Evaluated on draft-verify for text classification and plan-solve for mathematical reasoning, OPHSD outperforms baselines (+10.83% over OPSD on HMMT 2025), and reattaching the harness at inference yields no additional benefit, evidence that complex harnesses can serve as temporary training scaffolds whose benefits are permanently absorbed.

Cross-lingual reasoning via PI. COPSD (Liu et al., 2026h) transfers a model’s own high-resource reasoning behavior to low-resource languages. The teacher receives privileged cross-lingual context (the problem translation and reference solution in English) while the student sees only the low-resource problem. Training minimizes full-distribution token-level divergence on the student’s own rollouts, providing dense supervision while avoiding the sparsity and instability of outcome-only RL. Across 17 low-resource African languages, COPSD improves mathematical reasoning and outperforms GRPO by a wide margin, with the largest gains on the lowest-resource languages where the supervision gap is widest. MSD (Qin et al., 2026) (below) is a contemporaneous work that independently applies cross-lingual privileged-information-based self-distillation to safety alignment rather than reasoning, evidence that English-as-PI generalizes across capability dimensions.

Adaptive teacher exposure. ATESD (Han et al., 2026b) questions an assumption shared by nearly all PI methods, that the teacher should always see the *full* reference reasoning. A controlled exposure sweep shows that full exposure is not reliably optimal, and that student-teacher mismatch grows monotonically as the teacher sees more privileged context. ATESD models the reveal ratio with a lightweight Beta-policy controller conditioned on compact training-state statistics, optimized with a discounted learning-progress reward that scores each exposure decision by its effect on the student’s future improvement rather than immediate loss. Across Qwen3-{1.7B, 4B, 8B}, ATESD outperforms OPSD (+0.95 to +2.33 Average@12 points), positioning adaptive teacher exposure as a new axis for reasoning self-distillation.

Token-routed sparse self-OPD. TRACE (Wang et al., 2026e) argues that all-token KL in self-OPD spends gradients on mostly redundant positions and amplifies privileged-information leakage, producing entropy rise, shortened reasoning, and out-of-distribution degradation. TRACE uses a privileged annotator to mark critical spans in each rollout, applying forward KL to key spans of correct rollouts and optionally reverse KL to localized error spans, while leaving all other tokens to GRPO. A KL annealing schedule keeps cumulative privileged-gradient exposure finite. On four math benchmarks plus GPQA-Diamond, TRACE improves over GRPO by 2.76 points on average and is the only trained method in their evaluation that preserves OOD performance where GRPO and all-token self-OPD baselines degrade. Gains persist under online self-annotation (+1.90 points using the training student as annotator), reducing dependence on external supervisors.

Cross-lingual safety via PI. MSD (Qin et al., 2026) and COPSD (Liu et al., 2026h) independently explore cross-lingual privileged-information-based self-distillation in different domains, with MSD focusing on multilingual safety alignment and COPSD on cross-lingual mathematical reasoning (above). In MSD, the same model’s English-language reasoning serves as privileged information for its low-resource-language student, transferring safety boundaries without target-language safety data. The two parallel works indicate that cross-lingual PI is capability-agnostic. Its Dual-Perspective Safety Weighting (DPSW) concentrates distillation gradients on safety-critical tokens by combining teacher confidence with student disagreement. We discuss MSD’s broader implications for safety-aware OPD in Section 9.

Beyond KL matching via reward regularization. A recurring concern across the PI literature above is that directly matching the privileged teacher’s distribution can destabilize training and, in the worst case, degrade reasoning capability as training proceeds. PBSO (Yu et al., 2026c) diagnoses this pathology through a reward-regularized lens, replacing pure KL matching with the objective $\max_{\pi} \mathbb{E}[r(x, y)] - \beta D_{\text{KL}}(\pi \parallel \pi_{\text{teach}})$, whose analytic optimum is a reward-reweighted teacher $\pi^* \propto \pi_{\text{teach}}(y|x) \exp(r(x, y)/\beta)$ that provably dominates the

teacher itself whenever the reward is informative. Practically, PBSO realizes this optimum through a DPO-style pairwise loss where the preferred sample y^+ is drawn from the context-augmented teacher and the dispreferred sample y^- from the current student, preserving on-policy sampling while avoiding the forced imitation of every teacher token. On Qwen3-1.7B/4B/8B, PBSO matches OPSO’s token efficiency while surpassing its peak accuracy and avoiding the post-peak decline documented in Figure 2 of Yu et al. (2026c), preference-based targets can therefore serve as a more stable substitute for direct KL matching than curriculum or signal-augmentation patches alone. This post-peak decline is the same pathology that CaOPD (Zhang et al., 2026d) attributes to the training-deployment context mismatch and that Kim et al. (2026b) diagnose as epistemic suppression (Section 7.2), so PBSO can also be read as a self-distillation mechanism that converts these diagnostic insights into a corrective optimization target rather than a post-hoc recalibration step.

Outcome-conditioned PI at the turn level for agentic reasoning. The PI methods above all target single-turn generation. Extending PI to multi-turn agentic settings opens a separate question of granularity, at which temporal level should privileged information intervene? Existing agentic OPD work has explored the *trajectory* granularity (TCOD’s temporal curriculum over turn prefixes, Section 6.2), the *step* granularity (MAD-OPD’s step-level decomposition, Section 8), and the *skill* granularity (Skill-SD’s skill-conditioned updates, Section 8), but none of them supply PI at the *turn* level itself. TT-OPD (Jeong, 2026) fills this gap. A gradient-free EMA teacher receives correctness-dependent cues (reinforcing phrases for trajectories scored correct, corrective phrases for those scored incorrect) that are inserted at the prompt-response boundary and then stripped from the teacher’s output log-probabilities, so the student never observes the hints directly but still receives an outcome-aware KL signal at every conversation turn. Two aspects distinguish this instance from prior PI mechanisms. *First, the privileged context encodes trajectory-level correctness rather than problem-level ground truth*, which is one PI realization we are aware of where the teacher is conditioned on a derived property of the student’s own rollout (its outcome correctness) rather than on the problem-level ground truth directly. This generalizes the PI paradigm from single-response answer-conditioned supervision (GATES, OPSO, GUI-SD) to multi-turn behavioral regularization. *Second, the ablation isolates why turn-level PI matters here*: EMA teacher updates alone reduce the KL-collapse pathology but still allow the multi-turn structure to erode, and outcome hints without length control trigger response explosion, so neither EMA nor hints in isolation suffice. These pathologies are among the failure modes documented in Section 7.2, and the interaction between teacher dynamics and outcome-conditioned PI is what appears to keep agentic OPD stable in this setup.

Viewed together with TCO, Skill-SD, and MAD-OPD, TT-OPD completes a granularity spectrum for agentic distillation in which the choice of temporal unit (trajectory, turn, step, skill) is matched to the level at which errors compound in the target task.

Structured PI for video reasoning. VISO (Lin et al., 2026a) extends the privileged information paradigm from generic binary signals (OPSO’s correctness conditioning) to *structured multi-dimensional* feedback for VideoLLM reasoning. A video-aware judge model evaluates rollouts along answer correctness, logical consistency, and spatio-temporal grounding quality, producing structured feedback that decomposes reasoning errors into diagnostically meaningful subtypes. The central mechanism is direction-magnitude decoupling. Rollout-level advantages from environmental rewards determine the *update direction* (which trajectories to reinforce or penalize), while the structured teacher-student discrepancy conditioned on judge feedback modulates *token-level update magnitudes*. This ensures RL stability while allowing fine-grained credit assignment aligned with specific reasoning error types rather than undifferentiated token-level matching. Combined with curriculum scheduling (distillation→RL transition) and EMA teacher stabilization for long video sequences, VISO attains approximately $2\times$ faster convergence than RLVR baselines with consistent accuracy improvement on video reasoning benchmarks. The direction-magnitude decoupling parallels Self-Distilled RLVR’s decomposition (Section 5.2.3), where self-distillation sets magnitudes and RL sets directions, but VISO’s structured judge provides richer magnitude signals by distinguishing *why* a token matters (grounding error vs. logical inconsistency vs. factual mistake). VISO thus extends the PI spectrum from binary (OPSO) through spatial

(GUI-SD) to multi-dimensional structured feedback, with domain-specific error taxonomies determining what the PI encodes.

Multi-view, critique, and future-observation PI. As the privileged signal moves further from a literal answer, it begins to test the boundary of what can still serve as a usable self-teacher. AVSD (Nguyen et al., 2026) takes a modest step by treating multiple training-time views of the same input as parallel PI sources, decomposing the per-token signal into a consensus component shared across views and a view-specific residual so that the student does not collapse onto a single teacher view when views disagree. A larger departure comes from On-Policy Critique Distillation (Jin et al., 2026a), which draws PI not from the input but from a *weak* external critic. The strong student rolls out, a weak model critiques those rollouts, an outcome-and-rubric filter keeps only the useful critiques, and a critique-conditioned self-teacher then supervises the strong student via token-level KL. This inverts the conventional teacher-student capability ordering and instantiates scalable oversight, where weak supervision suffices once it is concentrated on the positions a rubric flags as informative. World-Model PI (Zhou et al., 2026a) pushes furthest, into multimodal foresight. The privileged context is the ground-truth future video and final answer, available to a training-time evaluator but absent from the student’s perception. The student rolls out under the unprivileged context while the evaluator builds a candidate set at intermediate decision nodes, whose advantage-weighted distillation teaches the student when to invoke a world-model rollout, when to verify it, and when to trust its result. The gradient from input perturbations through weak critiques to future observations shows that the self-teacher tolerates progressively more abstract privileged signals, with the binding constraint being not the signal’s form but whether it remains absent at inference.

Composing PI with outcome verification and calibration. A recurring concern is that a PI-conditioned self-teacher, however rich its context, can still point the student toward confident mistakes, so a second group of methods pairs the privileged signal with a corrective mechanism that vetoes or reweights it. Skill-Conditioned Gated SD (Huang et al., 2026) makes the correction a polarity switch. It maintains a skill-conditioned multi-teacher pool and uses a binary correctness signal $r \in \{-1, +1\}$ to distill toward the teacher when a rollout is correct and away from it when incorrect, with masking, clipping, and thresholding keeping the bounded gated objective stable. Constitutional Cross-SFT (Wen et al., 2026) applies the correction earlier, at teacher construction. To avoid the safety-expressiveness collapse documented for naive safety SFT, it inserts a Cross-SFT cold-start that calibrates the constitution-conditioned teacher before on-policy distillation begins, then optimizes a token-level reverse-KL reward under the calibrated teacher, with a geometric leakage analysis bounding how much constitutional content a reader could recover from the student. Self-Distilled Policy Gradient (Liu et al., 2026g) places the correction alongside the distillation term itself, composing full-vocabulary reverse-KL self-distillation with a REINFORCE objective on group-relative advantages, gating the SD term on positive advantage so that it tracks successful rollouts while the policy-gradient term drives discovery of solutions outside the teacher’s mode. Across the three, the privileged signal supplies density and the verification signal supplies direction, a division of labor that recurs whenever PI alone risks reinforcing plausible-but-wrong continuations.

Multimodal spatial and visual PI. The visual modality supplies its own catalog of privileged signals, and a cluster of methods shows that the same OPSD construction carries over once a vision-specific asymmetry is identified, each at a different level of visual granularity. At the representational level, MGSD (Luo et al., 2026b) confronts the modality gap directly, bridging the discrepancy between visual and language representations through a two-stage procedure in which a gap-bridging stage aligns feature distributions before the PI-based stage trains the student on its own rollouts under a gap-aware self-teacher. At the level of reasoning content, PTD-PO (Xiang et al., 2026) gives the teacher both spatial context and full reasoning traces withheld from the student, then concentrates a top- K JSD objective on the most informative positions in each rollout. At the level of raw perception, Thinking Without Images (Cai et al., 2026a) (TWI) lets the teacher see a full-resolution crop of the target region while the student sees only the original image, so on-policy KD internalizes the fine-grained grounding the crop reveals. Visual-SDPO (Dong, 2026) closes the loop back to symbolic structure, rendering a diagram or table constructed from the problem as a training-time PI

channel and combining statement-weighted KL with GRPO so that each token’s loss tracks the relevance of the rendered artifact. Read together, these four locate visual PI along a perception-to-abstraction axis, the same organizing move that the textual methods above make with answers and rules.

Future-observation and context-anchoring PI. HERO (Liu et al., 2026a) extends agentic PI from ground-truth conditioning (TT-OPD) to future environment observations. At training time, the teacher receives not only the current state but also the next environment observations, producing a policy that anticipates future context. The student distills this foresight via turn-level KL without access to future observations at inference. Where HERO injects future context, When Context Returns (Wang et al., 2026h) addresses a related problem, that adding long context at inference degrades a model previously trained without it. The method constructs a “no-context” self-teacher from the same model run without the additional context, applying forward KL as an anchoring regularizer that prevents the student from abandoning its context-free capabilities under context-induced distribution shift.

Bayesian credit redistribution and oracle residual PI. PBSO* (Tian et al., 2026) (Yang Tian et al., not to be confused with PBSO (Yu et al., 2026c)) tackles sparse trajectory-level reward in multi-turn agentic OPD through Bayesian self-distillation. Turn-level credit is derived from a Bayesian update rule that redistributes trajectory reward to individual turns using the self-teacher’s log-probability ratios, enabling fine-grained per-turn supervision without a separate reward model for each turn. AR-OPD (Zhang, 2026) identifies hindsight leakage as a failure mode of full oracle conditioning, where the privileged teacher sees the complete oracle answer and can “leak” backward, distorting the student’s earlier reasoning tokens. AR-OPD mitigates this by decomposing the distillation objective into an anchor term (standard on-policy KL without oracle) and a residual oracle correction applied only to tokens where the anchor’s uncertainty exceeds a threshold, achieving +2.3 points over full oracle OPD and +7.9 over SFT. Rubric-Guided SD (Rezaei et al., 2026) provides the rubric evaluation criteria as PI, conditioning the teacher on a structured grading rubric that the student must learn to satisfy without access to the rubric at inference. The JSD-based objective with rubric conditioning removes the external verifier call at inference while matching verifier-supervised performance, framing rubric internalization as an OPD task.

5.2.2 Pure Self-Distillation

Where privileged information creates asymmetry through additional context, the methods in this section extract a training signal from the model’s own output distribution without any auxiliary input. The common mechanism is *rollout diversity*: by sampling at non-zero temperature, the model produces a population of completions whose quality varies, and the variance itself supplies a selection or distillation signal. These methods sit on a staleness spectrum that ranges from fully on-policy (the training distribution is regenerated every gradient step) to mildly stale (rollouts are generated once per epoch and reused), with the degree of staleness trading off compute cost against distributional alignment.

Forgetting resistance through self-distillation. SDFT (Shenfeld et al., 2026) positions on-policy self-distillation as a mechanism for forgetting resistance. It uses the model’s own in-context learning ability to construct a demonstration-conditioned teacher from itself, then trains the student via reverse KL on its own rollouts, which is mathematically equivalent to trust-region RL with implicit KL regularization. This reframing converts continual learning from demonstrations into an on-policy distillation problem, where the trust-region constraint prevents catastrophic forgetting of previously acquired capabilities while accommodating new ones. The approach is fully on-policy (fresh rollouts at each step), placing it at the freshest end of the staleness spectrum.

Architectural self-distillation. MTP-SD (Kirchenbauer et al., 2026) targets *architectural* rather than behavioral improvement, converting a pre-trained autoregressive model into a standalone multi-token predictor by scoring student-generated token spans against a frozen copy of itself as teacher. This delivers $> 3\times$ faster decoding at typically 3–7% accuracy drop (model-dependent) without requiring any auxiliary module at inference time, expanding the scope of what self-distillation can optimize from output distributions

to generation architectures. AR-to-Diffusion (Su et al., 2026) extends the same principle across generation paradigms, using an autoregressive LM as a frozen self-teacher to train a diffusion LM through on-policy token-level distillation. The student generates diffusion-denoising trajectories, and the frozen AR teacher assigns token-level probabilities that guide the denoising distribution. By converting the AR LM into a diffusion LM, the method reaches its reported quality with $15\times-7000\times$ greater training-token efficiency than standard diffusion-LM pretraining, establishing that on-policy self-distillation can transfer knowledge across different generation architectures.

Systematic unification of self-distillation components. The methods above each investigate individual mechanisms (in-context demonstration teachers, multi-token span scoring), but the interaction effects across these design choices remain unclear. UniSD (Jin et al., 2026c) addresses this gap by proposing a unified framework that systematically decomposes LLM self-distillation along three complementary axes: *supervision reliability* (multi-teacher agreement to identify trustworthy self-derived signals, token-level contrastive learning to separate correct from plausible-but-wrong alternatives), *representation alignment* (feature matching to transfer internal structure beyond output distributions), and *training stability* (EMA teacher for temporal smoothing, divergence clipping to prevent rare high-divergence tokens from dominating updates). Evaluated across six benchmarks and six models spanning three families (Qwen2.5, Llama-3.1, Gemma-3), UniSD finds that on-policy trajectories combined with multi-teacher agreement and EMA stabilization form the core ingredients of effective self-distillation. The integrated variant UniSD* improves over the base model by +5.4 points and over the strongest baseline (GKD) by +2.8 on Qwen2.5-7B, without relying on any external teacher. UniSD’s three-axis decomposition provides a design vocabulary for future self-distillation methods, surfacing which design choices matter most.

Self-supervised reasoning signals. Three reasoning-focused methods extract a learning signal from contrasts that the model itself can score. Self-Supervised OPD (Tan & Hong, 2026b) repurposes the per-prompt group of GRPO rollouts as a self-supervised signal, contrasting correct and incorrect rollouts inside the same group to identify the per-token positions that separate them and applying token-level KL only on these contrastive positions, reusing GRPO’s existing rollouts at no additional inference cost. ROSD (Zhao et al., 2026d) replaces OPD’s solution-conditioned self-teacher with a reflector that, given a wrong rollout, extracts a corrective idea explaining the error and locates the first erroneous span via an error quote. The corrective idea conditions the self-teacher’s distribution toward error-specific supervision, and the error quote restricts distillation to the suffix starting from the error so that valid reasoning prefixes are preserved. Variational Posterior Guidance (Chen et al., 2026e) adds an efficiency-aware variational posterior that scores each candidate rollout by a length-aware utility, biasing the self-distillation target toward shorter correct trajectories rather than verbose ones. The three methods illustrate that the asymmetry needed for self-distillation can be sourced from intra-group contrast (Self-Supervised OPD), reflective error localization (ROSD), or posterior-weighted utility (VPG), without auxiliary models.

Multi-turn and agentic self-distillation. A second cluster extends pure self-distillation to settings where the student rolls out a long trajectory and the supervision target is constructed retrospectively. HINT-SD (Yeo et al., 2026) and SD-Search (Ma et al., 2026) both apply hindsight relabeling but at different granularities. HINT-SD selects only the failure-relevant turns from a long-horizon agent rollout and runs token-level self-distillation against a hindsight-relabeled successful continuation, while SD-Search localizes the supervision to step-level search queries inside a search-augmented reasoning trajectory, combining the resulting per-step signal with outer-loop GRPO. Search-E1 (Liang et al., 2026b) alternates outer-loop GRPO with offline self-distillation over the GRPO-improved checkpoints, treating each round of distillation as a self-evolution step that consolidates the policy improvements before the next RL round. MAIGO (Zheng et al., 2026b) addresses the lost-in-conversation gap that emerges in pure multi-turn settings by constructing a history-cleaned privileged teacher (the same model conditioned on a single-turn reformulation of the task) and supervising student rollouts on the original noisy multi-turn conversations, raising the sharded-to-full accuracy ratio from 66.5% to 84.1% on Qwen2.5-7B-Instruct while keeping full-context accuracy within 2.3 points. Canonical-Context OPD (Lin et al., 2026b) pursues the same multi-turn alignment from a different angle. The student rolls out from a raw-sharded

conversation history, while a frozen same-backbone teacher answers the same prompt under a canonical full-context history, and an answer-masked reverse KL aligns the student’s behavior with the canonical-context teacher across same-prefix relabeling. COMAP (Liu et al., 2026i) couples the agent and a textual world model in a co-evolution loop where on-policy self-distillation with an EMA teacher is applied to the world model while the agent uses future-aware reflection with an action gate, with a world-state gate curriculum and a suffix-based refinement initialization keeping the loop stable. Across these six methods, the common pattern is that the supervision target is reconstructed from the rollout itself (hind-sight relabeling, history-cleaned reformulation, canonical context, world-model future), turning the multi-turn structure into the source of asymmetry rather than a complication to be tolerated.

Domain-specific self-distillation. Three further methods specialize pure self-distillation to domains where the asymmetry comes from a domain-native split. Vision-OPD (Yuan et al., 2026) introduces a regional-to-global split for fine-grained MLLM understanding, where the regional-conditioned variant of the model serves as the self-teacher for the global-conditioned student that lacks region annotations at inference, applying token-level reverse KL on student rollouts. TOD Proactivity (Zhang et al., 2026a) treats privileged user-concern annotations as the asymmetric signal in task-oriented dialogue, training an asymmetric self-teacher conditioned on annotated user concerns and distilling onto a student that observes only raw utterances, so that the proactive information-gathering behavior is internalized without inference-time access to the annotation. It Takes Two (Park et al., 2026) pursues complementary self-distillation for contextual-integrity privacy alignment, where two parallel rollouts under complementary contextual splits supply mutual supervision, anchoring privacy decisions across context perturbations rather than against any single static teacher. The three methods adapt the rollout-diversity principle of pure self-distillation to multimodal grounding (Vision-OPD), dialogue policy (TOD Proactivity), and privacy alignment (It Takes Two), evidence that the architectural pattern transfers across modalities once a domain-appropriate asymmetry is identified.

Cross-domain co-distillation. OPCoD (Byeon et al., 2026) extends the mutual-supervision idea from a single model’s complementary rollouts to two specialist models that tutor each other across domains. Each model runs on-policy self-distillation whose self-teacher is conditioned on both its own verified-correct rollout and a natural-language feedback message from its peer, so the two self-distillation loops are coupled and co-evolve during training. A cognizance-based gate exchanges feedback only when the peer is more competent on the current domain, mitigating the negative transfer that ordinarily arises when mixing data from different fields, and a feedback-anchoring term grounds the message in the problem to prevent it from corrupting an already-correct rollout. On science question-answering across domain pairs such as physics-chemistry and chemistry-materials, OPCoD reaches mutual Pareto improvement on Qwen3-8B, with each model gaining out-of-domain capability without losing its original specialty. Where CoPD (Gu et al., 2026) couples two models within a single capability and It Takes Two splits one model’s context, OPCoD targets the multi-domain regime where the asymmetry is the peer’s distinct specialization.

These methods collectively show that pure self-distillation can deliver meaningful gains across coding, reasoning, and instruction-following tasks without external supervision. The gains are bounded, however, by the information already present in the model’s pre-training distribution. When temperature diversity alone is insufficient to surface high-quality completions, bridging this boundary requires either connecting the model to external verification signals or internally generating privileged information, both of which motivate the external feedback mechanisms below.

5.2.3 External Feedback

The pure self-distillation methods above generate training signal from the model’s own distribution, but this signal is bounded by the information already encoded in the model’s parameters. Without access to ground truth or environmental verification, the model cannot distinguish confident-but-wrong outputs from genuinely correct ones, a limitation that leads to the saturation problem analyzed in Section 7.2. The methods in this section retain

the self-distillation architecture (no separate teacher model) while injecting external signals, whether from verifiers, environments, multi-agent co-evolution, or partial-solution oracles, that anchor the training loop and prevent it from collapsing onto self-reinforcing errors.

From binary rewards to dense supervision. A line of work converts the sparse outcome rewards of RLVR into dense token-level self-supervision through different forms of feedback. SDPO (Hübotter et al., 2026) is an early instance, attributing success or failure to specific reasoning steps using *structured textual* feedback (runtime errors, failing unit tests, and LLM judge evaluations), with the feedback-conditioned model serving as a self-teacher for token-level KL distillation. RL from Text Feedback (RLTF) (Song et al., 2026) concurrently uses free-form natural language critiques from an automated judge, training the model’s single-turn policy to match its own feedback-conditioned second-turn generations. SD-ZERO (He et al., 2026), a more recent instance, takes a different angle by converting the same binary verifier signal into dense self-supervision through a dual-role Generator-Reviser architecture, where the Reviser is conditioned on the Generator’s output and its binary correctness signal and produces an improved version, and the Generator then distills from the Reviser’s token-level distribution. On Qwen3-4B-Instruct, SD-ZERO records 68.3% average accuracy across eight benchmarks (avg@8), outperforming GRPO (64.2%), showing that self-revision can be more effective than standard RL for converting sparse rewards into learning signal.

From SD-ZERO’s binary signal through SDPO’s structured errors to RLTF’s free-form critiques, a clear trajectory emerges. As external feedback becomes richer, the self-distillation objective becomes more precisely targeted, allowing the model to correct specific failure modes rather than globally adjusting its policy. RESD (Zhang et al., 2026n) pushes this trajectory further by transforming raw failure feedback into an active source of corrective supervision rather than treating it as a passive conditioning signal. When the student’s on-policy rollouts fail, RESD generates retrospective reflections that diagnose local errors, and curates a persistent global playbook of reusable lessons across training steps. The enriched context enables the self-teacher to provide actionable token-level supervision even without any successful demonstrations, directly addressing the rare-success regime where GRPO and standard OPSD lose signal because reward variance collapses to zero. On continual agentic learning tasks, RESD outperforms GRPO by a factor of $8\times$ in sample efficiency (Amazon/UCSD), evidence that structured reflection on failures can substitute for the successful demonstrations that most self-distillation methods implicitly require. Sample-Routed Policy Optimization (SRPO) (Li et al., 2026c) shows that these approaches can be composed by routing. Correct samples follow GRPO’s reinforcement path (reward-weighted likelihood maximization) while failed samples follow SDPO’s targeted logit-level correction. Correct and incorrect samples carry distinct types of learning signal, and treating both identically wastes the information each carries. SRPO raises the five-benchmark average on Qwen3-8B by 3.4% over GRPO and 6.3% over SDPO alone across science and tool-use tasks. A useful external-feedback strategy routes each sample to its most informative objective.

Decomposing magnitude and direction. Self-Distilled RLVR (RLSD) (Yang et al., 2026a) decomposes the gradient update into two functionally distinct components. Self-distillation determines the *magnitude* of per-token policy updates (how much to change each token’s probability), while environmental RLVR feedback from verifiers anchors the *direction* (which tokens should increase vs. decrease in probability). At just 200 training steps, RLSD surpasses GRPO trained for 400 steps on Qwen3-VL-8B-Instruct, exhibiting $2\times$ sample efficiency from the dual signal structure. This decomposition clarifies why external grounding and self-distillation are additive rather than competing. The self-distillation component calibrates update magnitudes using the model’s internal distributional geometry, while the external verifier ensures these updates point in the right direction.

RLRT (Kim et al., 2026a) inverts this self-distillation paradigm entirely. While standard self-distillation guides the student toward the teacher’s predictions, RLRT observes that on *successful* rollouts (as judged by an external verifier), tokens where the teacher would not have predicted the student’s choice reflect the student’s self-driven reasoning rather than teacher imitation. By reading the self-distillation signal in reverse and reinforcing these “rebellious” tokens via GRPO augmentation, RLRT frames exploration not as uniform diversity but as valuable deviation grounded in the student’s own success. On base,

instruction-tuned, and thinking-tuned Qwen3 checkpoints, RLRT outperforms standard self-distillation and exploration-based baselines, positioning information asymmetry as a design axis for externally grounded self-improvement.

Gated self-distillation for multi-turn agents. Where RLSD decomposes gradients into magnitude (self-distillation) and direction (RL), SDAR (Lu et al., 2026) takes this decomposition into multi-turn agentic settings and identifies a failure mode specific to that regime, compounding multi-turn instability destabilizes the self-teacher’s supervision, while skill-conditioned privileged guidance produces negative teacher signals that arise from imperfect skill retrieval rather than genuinely poor student behavior. SDAR treats on-policy self-distillation as a gated auxiliary objective while keeping RL as the primary optimization backbone, mapping detached token-level teacher-student log-ratio signals through a sigmoid gate that strengthens distillation on teacher-endorsed positive-gap tokens and softly attenuates negative teacher rejections. Across the Qwen2.5 and Qwen3 families on ALF-World, WebShop, and SearchQA, SDAR improves over GRPO by +9.4%, +10.2%, and +7.0% respectively. The asymmetric treatment of positive and negative teacher signals appears to be a key factor when transferring self-distillation from single-turn to multi-turn settings. More broadly, the external feedback methods in this section all share one property that distinguishes them from pure self-distillation: compiler outputs, unit test results, and symbolic verifiers supply unambiguous correction signals that break self-reinforcing echo chambers, because even a model that is highly confident in a flawed derivation receives $R = 0$ from the verifier.

Input-specific credit via pointwise MI decomposition. CREDIT (Shen et al., 2026b) addresses a subtler pathology in the self-distillation reward. Under a posterior-compatibility interpretation of feedback conditioning, the standard OPSD token reward reduces to a Bayesian filtering increment whose trajectory sum equals the pointwise mutual information (pMI) between the response and the feedback given the input. This pMI can be inflated by input-generic shortcuts (responses that are generically likely under any positive feedback) rather than input-specific reasoning. CREDIT isolates the input-specific component through a batch-contrastive baseline that penalizes responses remaining likely under unrelated inputs. At the sequence level, this functions as a teacher-side surrogate for a contrastive pMI objective. Across coding, scientific reasoning, and tool-use benchmarks on two model families, CREDIT delivers strong aggregate performance at negligible additional compute. Generic-shortcut suppression is therefore a broadly useful correction for self-distillation rewards.

Outcome-guided logit steering. OGLS-SD (Yang et al., 2026f) identifies a mismatch between teacher and student in OPSD that prior adaptive methods overlook, reflection-induced bias in teacher responses. Self-reflected teacher distributions can be shifted by response templates and solution-conditioned artifacts, producing miscalibrated token-level supervision. OGLS-SD addresses this by using verifiable outcome rewards to contrast successful and failed on-policy trajectories, using the outcome signal to steer teacher logits toward calibrated targets. By combining outcome-level correctness with dense token-level guidance through logit steering, OGLS-SD stabilizes self-distillation and improves reasoning performance over standard OPSD across diverse benchmarks. Outcome information can therefore serve as an independent calibration axis for self-distillation objectives.

Multi-agent co-evolution and partial-solution oracles. The external signals above all come from environment-side verifiers. An independent source of external grounding emerges when the model generates its own verification signal through multi-agent interaction or partial-solution disclosure. π -Play (Zhang et al., 2026l) unifies privileged information and self-play within a single multi-agent self-evolution framework. Three co-evolving agents (examiner π_ϕ^E , teacher π_ψ^T conditioned on a question construction path, student π_θ^S without it) are jointly optimized through alternating updates. Self-play produces a byproduct during task generation, as the examiner agent constructs a question construction path (QCP) that provides high-quality PI at no additional cost. As the student strengthens, the examiner generates harder tasks, and the teacher’s behavior is softly constrained to track the student via a KL penalty, creating a self-adjusting curriculum that requires no external data. On Qwen3 (4B, 8B), data-free π -Play surpasses fully supervised search agents such as Search-R1

and delivers markedly higher evolutionary efficiency than conventional self-play (Dr.Zero), with pronounced gains on multi-hop benchmarks where the teacher’s token-level credit assignment has considerable effect.

PAINT (Tan & Hong, 2026a) occupies a distinctive position between pure self-distillation (no external information) and full PI (complete answer revealed). Rather than conditioning on the entire ground-truth solution (as in OPSD), PAINT exposes only a *partial* solution whose extent is adaptively controlled by the overlap between the student’s rollout and the reference. It computes a recall-style overlap score based on the fraction of reference anchors (boxed answers, formulas, key numbers) that appear in the student’s rollout, hiding more of the reference when the student is close (encouraging broader generalization since the rollout already captures the essential structure) and revealing more when far (providing stronger corrective guidance when the student has drifted). This overlap-adaptive partial-solution masking creates a natural curriculum within the self-distillation mechanism itself. PAINT further employs energy-space interpolation that applies the distillation loss only at positions where the teacher-student entropy mismatch exceeds a threshold, producing a sparse loss that concentrates gradient signal on the most informative positions. On competition-level mathematics (AIME 2024/2025 and HMMT 2025 average), PAINT records +2.1 over the OPSD baseline and +2.9 over GRPO, partial revelation of privileged information can outperform both full revelation and no revelation, an information disclosure curve whose theoretical characterization remains open.

Co-evolutionary distillation. A separate thread abandons the fixed-teacher assumption entirely, training teacher and student jointly so that both improve during distillation. CoPD (Gu et al., 2026) shows that the sequential pipeline (train experts first, then distill) can be replaced by co-evolutionary training where RLVR and bidirectional OPD are interleaved, maintaining behavioral proximity that enables continuous knowledge absorption instead of post-hoc transfer. Variational Policy Distillation (VPD) (Li et al., 2026i) formalizes this co-evolution as a Variational Expectation-Maximization problem, where an E-step refines the teacher on trajectory outcomes via an adaptive trust-region update that translates textual feedback into a dynamically improved target distribution, and an M-step has the student internalize this distributional guidance on its own rollouts. By continuously improving the teacher’s ability to extract actionable signals from diagnostic critique, VPD mitigates the plateau that passive (fixed-teacher) self-distillation encounters as the student improves. On scientific reasoning and code generation tasks, VPD outperforms both standard RLVR and existing self-distillation baselines, while stress-tests on mathematical reasoning and cold-start regimes illuminate the practical bounds of feedback-driven self-distillation compared to pure environment-driven RL.

Saturation analysis. Kim et al. (2026b) trace self-distillation degradation in mathematical reasoning to the suppression of *epistemic verbalization*, the model’s expression of uncertainty during reasoning. When self-distillation shortens reasoning traces, it disproportionately removes hedging phrases and uncertainty markers that serve as implicit “attention flags” for difficult sub-problems. The result is shorter but less calibrated reasoning, where the model confidently commits to flawed steps. This connects to the teacher uncertainty challenge in Section 7, where an overconfident student produces unreliable self-distillation targets, and the saturation problem tends to require either external grounding or the internally generated PI mechanism of π -Play to overcome.

Frozen-copy consistency for safety alignment. On-Policy Consistency Training (Han et al., 2026a) adapts the external-feedback template to safety alignment without verifiers or reward models. A frozen copy of the pre-update policy serves as the self-teacher, the student rolls out under contrastive prompt pairs (one canonical and one perturbed by sycophancy or jailbreak triggers), and per-token reverse KL between the frozen-copy teacher on the canonical prompt and the student on the perturbed prompt aligns the student’s response on the perturbed prompt with the teacher’s response on the canonical one. Because the teacher and student share a backbone and only the conditioning prompt differs, the construction supplies dense per-token consistency supervision while keeping the safety update local to the perturbed branch. Across sycophancy, jailbreak-defense, and safety-awareness benchmarks the method improves safety scores with minimal capability degradation. Consistency-against-

Table 9: OPD methods: Training efficiency and stabilization.

Method	Year	Category	Divergence/Objective	Signal	Granularity	Key Innovation
TIP (Xu et al., 2026c)	2026	Dynamics	Soft-OR token weighting	White-box	Token	Entropy x divergence quadrant scoring
SCOPE (Zheng et al., 2026a)	2026	Dynamics	Rollout routing	White-box	Token	Redirect on flawed prefix
SelecTKD (Huang et al., 2025)	2025	Dynamics	Confidence filtering	White-box	Token	Top-1 probability threshold
AdaSwitch (Peng et al., 2025)	2025	Dynamics	Exploration/guidance switch	White-box	Token	Binary mode switching
PACED (Xu et al., 2026b)	2026	Dynamics	Beta-kernel curriculum	White-box	Sample	Frontier difficulty sampling
FOPD (Zhang et al., 2026b)	2026	Dynamics	Prefix truncation	White-box	Hybrid	Early-token focus
Lightning-OPD (Wu et al., 2026a)	2026	Dynamics	Offline caching	White-box	Token	4.0x cost reduction
SKD (Xu et al., 2025c)	2025	Dynamics	Draft-verify	White-box	Block	Amortized teacher cost
NFD (Rang et al., 2026)	2026	Dynamics	Async gen-train decoupling	White-box	Token	8.1x speedup via Δ -IFD filtering
Prune-OPD (Yang et al., 2026g)	2026	Dynamics	Top-k overlap drift detection	White-box	Token	Dynamic trajectory pruning, 37-68% time reduction
R-OPD (Fu et al., 2026b)	2026	Dynamics	Top-p rollout + masking	White-box	Token	Empirical failure mode fixes
SOD (Zhong et al., 2026a)	2026	Dynamics	Step-level divergence reweighting	White-box	Step	Tool-call cascade attenuation
Stable-OPD (Luo et al., 2026a)	2026	Dynamics	RKL + reference div.	White-box	Token	Rollout mixing for stability
TCOD (Wang et al., 2026d)	2026	Dynamics	Temporal curriculum	White-box	Trajectory	Multi-turn KL stability
Uni-OPD (Hou et al., 2026)	2026	Dynamics	Dual-perspective data balancing	White-box (multi-teacher)	Token	Unified LLM+MLLM OPD
CaOPD (Zhang et al., 2026d)	2026	Dynamics	Calibration-aware OPD	White-box	Token	Diagnoses scaling law of miscalibration
EGRSD (Ke et al., 2026)	2026	Dynamics	Entropy-gated self-distillation	Self (PI)	Token	Teacher-entropy confidence gate + causal lookahead
MOPD (Yu et al., 2026b)	2026	Dynamics	Peer-conditioned multi-rollout	White-box	Sample	Success/failure rollout group weighting
GEAR (Li et al., 2026g)	2026	Dynamics	Adaptive segment reweighting	Self (PI)	Segment	Granularity-adaptive advantage boundaries
EHOPE (Cai et al., 2026b)	2026	Dynamics	Parameter-dynamics foresight	White-box	Token	3x acceleration via update extrapolation
Visual-Adv OPD (Liu et al., 2026c)	2026	Dynamics	Visual-grounded reweighting	White-box	Token	Per-token visual-advantage gradient
Filter-Then-Reweight (Li et al., 2026m)	2026	Dynamics	Trajectory filter + token reweight	White-box	Token/Trajectory	Dual-granularity OPD
Less-is-More (Ziheng et al., 2026)	2026	Dynamics	Early-stop rollout	White-box	Token	Off-policy teacher-decay truncation
DeltaPrompts (Jung et al., 2026)	2026	Dynamics	Prompt-divergence curriculum	White-box (VLM)	Prompt	Zero-delta-trap synthesis
f-OPD (Chen et al., 2026c)	2026	Dynamics	Freshness-aware control	White-box	Trajectory	Async-OPD lag bounding
TR-Behav. Blend (Plyusov et al., 2026)	2026	Dynamics	Trust-region warmup	White-box	Token	Closed-form geometric blend
Adaptive Refresh (Guo et al., 2026)	2026	Dynamics	Three-gate teacher refresh	Self (PI)	Trajectory	Isolation/reward/length-tail gating
CEI (Chen et al., 2026a)	2026	Dynamics	Iterative experience internalization	Self	Step	Multi-iter agent self-evolution
POPD/TOFD (Zhang et al., 2026k)	2026	Dynamics	Horizon-controlled rollout	White-box	Token	Progressive/truncated horizon
SafeSteer (Li et al., 2026e)	2026	Dynamics	Safety-token localized RKL	Self (steered)	Token	Activation-steered safety teacher
TRD (Jiang et al., 2026a)	2026	Dynamics	Trajectory-refined distillation	White-box	Sequence	Teacher correction at prefix-failure points
KAT (Xin et al., 2026)	2026	Dynamics	KL-agreement trap truncation	White-box	Token	Online rollout truncation at KL-agreement trap regions

perturbation can therefore substitute for an external safety reward when the asymmetry can be encoded in the prompt alone.

Sign-consistency gating with phased teacher sampling. SG-OPD (Xu et al., 2026a) introduces a verifier-guided gating mechanism that filters the token-level distillation signal by the sign consistency between the student’s rollout gradient and a separately computed verifier signal. A phased teacher-sampling schedule warms up on high-quality teacher responses in early training before transitioning to on-policy rollouts, smoothing the distribution shift that arises when a weak student is forced to match a strong teacher from the first iteration. On competition-level mathematics benchmarks (AIME 2024/2025 and HMMT 2025), SG-OPD improves over the base GRPO setting by +1.98 at the per-sample level (avg@32) and +7.50 at the per-question level (pass@32), attributing the gains to discarding spurious token-level gradients that would otherwise reinforce confident-but-incorrect reasoning steps.

Across signal sources, from white-box logits through black-box APIs to self-distillation, a fundamental trade-off between signal density and autonomy governs the design space. White-box methods supply the densest supervision ($|V|$ -dimensional distributions at every token) but require same-organization deployment and co-location of teacher and student. Black-box methods sacrifice distributional information but operate across organizational boundaries, with recent methods (GAD, OVD) recovering effective proxy signals. Self-distillation removes external dependencies but is bounded by the model’s pre-existing capabilities unless augmented by external verification or internally generated PI. This continuum also reflects the field’s maturity gradient. The white-box extreme is relatively mature (GKD, DistiLLM, DSKD), while the self-distillation extreme still lacks principled convergence guarantees for complex reasoning. Recent activity concentrates at the boundaries between categories, with methods like LUFFY blurring the line between off-policy demonstrations and on-policy RL, and π -Play bridging self-distillation with externally grounded co-evolution.

6 Training Efficiency and Stabilization

Given an objective function (Section 4) and a signal source (Section 5), the final engineering decision is *how to stabilize and accelerate the on-policy training loop*. On-policy generation introduces optimization challenges largely absent from standard supervised fine-tuning, including non-stationary data distributions (the student’s policy shifts during training, making older rollouts stale), gradient signal-to-noise ratio (SNR) collapse on hard prompts (where all rollouts fail and no useful gradient emerges), and the computational overhead of autoregressive student rollouts followed by teacher scoring (several times more expensive than off-policy training).

These challenges are not independent. SNR collapse on hard prompts wastes compute, motivating curriculum strategies that avoid such prompts. Non-stationarity invalidates cached teacher signals, motivating efficient online scoring. And the compute overhead of autoregressive generation creates a systems-level bottleneck that motivates architectural solutions. The methods in this section tackle these challenges through three complementary mechanisms.

6.1 Token and Sample Weighting

Not all teacher supervision is equally reliable in the on-policy setting. The core problem is the *flawed prefix trap*, where a student generating an erroneous token early in a sequence causes all subsequent teacher predictions to be conditioned on this out-of-distribution prefix. Because the teacher has never been trained on such prefixes, its conditional distribution $p_T(\cdot|x, \hat{y}_{<t}^{\text{bad}})$ becomes noisy and unreliable. Naively matching this corrupted signal propagates errors instead of correcting them.

Teacher reliability filtering. TIP (Xu et al., 2026c) (Token Importance Profiling) offers a principled framework for token weighting by organizing importance along two axes, student entropy h_t and teacher-student divergence δ_t . Crossing these axes yields four quadrants with qualitatively distinct learning roles. High-entropy, high-divergence tokens (Q1) carry the dominant corrective signal, high-entropy, low-divergence tokens (Q2) indicate agreement on uncertainty, low-entropy, high-divergence tokens (Q3) represent *overconfident errors* where the student is certain but wrong, and low-entropy, low-divergence tokens (Q4) are genuinely mastered. A theoretical result establishes that Q3 tokens are structurally invisible to any entropy-only weighting scheme, because no non-decreasing function of entropy with $f(0) = 0$ can distinguish them from Q4. TIP’s parameter-free token selection recovers Q3 coverage while preserving Q1/Q2 sensitivity. On long-horizon planning benchmarks like DeepPlanning, training on exclusively Q3 tokens (less than 20% of all tokens) surpasses full-token OPD.

SCOPE (Zheng et al., 2026a) elevates this analysis from tokens to *rollouts*. Its dual-path architecture routes rollouts by correctness. Incorrect trajectories receive teacher-perplexity-weighted KL distillation (down-weighting unreliable guidance on flawed prefixes), while correct trajectories receive student-perplexity-weighted MLE (concentrating reinforcement at the capability boundary). This counters diversity collapse, the counterintuitive phenomenon where Pass@1 improves at the expense of Pass@k, yielding a 7.3% relative Pass@32 gain over competitive baselines across six reasoning benchmarks.

EGRSD (Ke et al., 2026) unifies token-level weighting with the self-distillation setting by introducing an entropy-guided confidence gate that down-weights high-entropy teacher positions while maintaining a nonzero lower bound on every token weight. The central observation is that teacher entropy varies widely across a chain-of-thought sequence, and uniformly weighting all positions forces the student to match noisy teacher predictions at genuinely uncertain reasoning transitions. EGRSD further proposes a causal-lookahead variant (CL-EGRSD) that distinguishes sustained high-entropy spans (where the teacher is genuinely uncertain and supervision is unreliable) from transient high-entropy positions whose following context rapidly resolves to low entropy (where the momentary uncertainty reflects a branching point rather than teacher ignorance). This temporal distinction complements TIP’s static quadrant analysis with a dynamic view of how entropy evolves across the reasoning trace. On Qwen3-4B and Qwen3-8B in thinking mode, EGRSD and CL-EGRSD advance the accuracy-length frontier among compared trainable methods, entropy-aware gating improves self-distillation efficiency without sacrificing coverage.

SelectKD (Huang et al., 2025) applies a propose-and-verify mechanism at the token level, where the student proposes tokens that the teacher verifies through greedy Top-k or non-greedy Spec-k matching. Accepted tokens receive full loss while rejected tokens are masked or down-weighted, inducing an implicit curriculum quantified by Token Acceptance Rate (TAR). Positions where teacher-student agreement fails indicate either genuine ambiguity or distributional divergence, and in neither case does the unfiltered gradient carry reliable information.

AdaSwitch (Peng et al., 2025) presents a dynamic switching mechanism between pure exploration and guided distillation. Unlike static filtering that simply attenuates unreliable signals, AdaSwitch maintains a sliding window of recent token-level divergences (KL or JSD between student and teacher logits) and computes a running average $\bar{d}_{i-1} = \frac{1}{L} \sum_{j=i-L}^{i-1} d_j$ over a window of length L . When the current token’s divergence exceeds a context-adaptive threshold $\tau_i = K \cdot \bar{d}_{i-1}$ (where K is a sensitivity multiplier), the method switches from on-policy exploration to off-policy teacher guidance. This single-switch design preserves semantic coherence (avoiding the jarring token-level alternation of prior mixing methods) while providing high-quality supervision precisely when the student’s generation drifts into high-divergence regions. The threshold’s adaptive nature matters here, as a fixed threshold tends to be either too aggressive (switching on easy prefixes where minor divergence is harmless) or too permissive (failing to intervene when the student enters genuinely problematic states). By conditioning on recent divergence history, AdaSwitch calibrates its intervention threshold to the local sequence difficulty.

These four methods span a design spectrum from fine-grained to coarse-grained intervention. TIP applies continuous weighting over every token, SCOPE routes entire rollouts into distinct loss pathways, SelecTKD makes binary accept/reject decisions per token, and AdaSwitch commits to a single trajectory-level mode switch. The ordering exposes an underlying tradeoff between signal fidelity (finer granularity preserves more information) and optimization stability (coarser decisions reduce gradient variance). This tradeoff is not merely theoretical. Fine-grained methods like TIP require reliable token-level teacher signals, which degrade precisely when the flawed prefix problem is severe. Coarse-grained methods sacrifice per-token optimality but are robust to local signal corruption because a single bad token cannot contaminate the entire trajectory’s gradient. The choice along this spectrum correlates with the dominant failure mode. Token-level methods tend to excel when teacher reliability varies sharply within a sequence (e.g., long reasoning chains where some steps are easy and others require critical insight), while trajectory-level methods are preferred when the primary concern is prefix-quality drift across the generation (e.g., multi-turn dialogue where early errors shift the entire conversation topic).

Sample-level weighting. Beyond token-level filtering, several methods weight entire samples based on their informativeness. The flawed prefix trap manifests most severely on *hard* prompts where the student’s pass rate approaches zero. On these prompts, nearly all rollouts contain early catastrophic errors, and the resulting teacher signals are dominated by noise. Methods that uniformly sample prompts waste most of their gradient budget on uninformative hard prompts or trivially correct easy prompts, with the productive learning concentrated on prompts at the boundary of the student’s competence.

Failure-mode analysis. Fu et al. (2026b) present a systematic empirical analysis of why sampled-token OPD is fragile in long-horizon settings. They identify three failure modes: (1) an imbalanced token-level supervision, (2) unreliable teacher guidance on student-generated prefixes, and (3) tokenizer or special-token mismatch. Their proposed fix, teacher top- K local support matching, a truncated reverse-KL objective with top- p rollout sampling and special-token masking, produces more stable optimization across reasoning and agentic settings and yields a 19.8% performance gain over standard sampled-token OPD baselines.

Step-level reweighting for tool-integrated reasoning. The weighting mechanisms above all operate at token or trajectory granularity. SOD (Zhong et al., 2026a) identifies a failure mode specific to tool-integrated reasoning (TIR) agents that motivates an intermediate *step-level* granularity. In TIR settings, erroneous tool calls inject corrupted observations into the reasoning context, causing student-teacher divergence to accelerate sharply at subsequent steps rather than drift gradually as in text-only reasoning. Standard token-level OPD cannot detect this discontinuous state transition because it lacks visibility into inter-step boundaries. SOD adaptively reweights distillation strength at each reasoning step based on the accumulated step-level divergence, attenuating potentially misleading teacher signals in high-divergence regions while preserving dense guidance where teacher-student alignment remains intact. Using Qwen3-4B as teacher and Qwen3-0.6B/1.7B as students, SOD outperforms the second-best baseline by up to 20.86% across math, science, and code benchmarks, with the 0.6B student reaching 26.13% on AIME 2025 (average@32). The

step-level granularity fills a gap between TIP’s per-token weighting (too fine for detecting tool-call cascades) and SCOPE’s trajectory-level routing (too coarse for preserving useful supervision at well-aligned steps within a partially corrupted trajectory).

Peer-conditioned multi-rollout weighting. MOPD (Yu et al., 2026b) identifies a limitation shared by the weighting methods above. They condition the distillation signal on the relationship between teacher and student at each position, without exploiting information from *other student rollouts* for the same prompt. By partitioning a batch of student rollouts into success and failure groups, MOPD constructs peer-conditioned teacher signals that contrast what the teacher would predict given a successful student trajectory versus a failed one. This peer conditioning provides a form of contrastive credit assignment that is unavailable from any single rollout, because the teacher’s signal on a failed rollout can be decomposed into a component shared with successful rollouts (uninformative) and a residual that isolates the failure-specific divergence. The approach draws on a different source of information (cross-rollout structure) than the token-level entropy (TIP, EGRSD) or step-level divergence (SOD) methods, and the two are in principle complementary.

Granularity-adaptive advantage reweighting. GEAR (Li et al., 2026g) adaptively segments each trajectory into variable-length spans based on the divergence between the on-policy student and a ground-truth-conditioned teacher. Rather than applying uniform token-level or predetermined step-level boundaries, GEAR detects segment boundaries where the student-teacher advantage profile shifts, then computes per-segment advantages that concentrate the learning signal at the granularity most informative for each region of the trajectory. This adaptive segmentation addresses a gap between token-level methods (which can be noisy on individual positions) and trajectory-level methods (which average over heterogeneous regions). On LLM agent tasks, GEAR improves over fixed-granularity baselines. The optimal weighting granularity therefore varies not only across tasks (as SOD’s step-level analysis implies) but also within a single trajectory.

Multimodal and horizon-aware reweighting. Three further methods extend the weighting principle to multimodal and horizon-aware settings while keeping the underlying mechanism intact. Visual-Advantage OPD (Liu et al., 2026c) adapts token-level reweighting to vision-language students by computing a per-token visual advantage that quantifies how much each token of the student rollout depends on visual evidence rather than language priors, and then concentrating the OPD gradient on visually grounded tokens. The construction parallels Decomposed-OPD’s gradient steering (Section 4.1) but lives in the weighting layer rather than the loss form. Filter-Then-Reweight (Li et al., 2026m) composes coarse and fine granularities. A trajectory-level filter on teacher log-probability discards rollouts that the teacher itself cannot anchor, and a soft per-token reweight then blends teacher confidence with student confusion to retain partial information from the surviving trajectories. A PPO-style clipped objective on the weighted advantages keeps the update bounded even when the per-token weights vary widely. Less-is-More (Ziheng et al., 2026) attacks the same teacher-reliability question from a horizon angle, identifying *Off-policy Teacher Decay*, the observation that later positions in a student rollout become increasingly off-policy for the teacher and the supervision quality drops. Truncating the rollout at the position where teacher reliability begins to deteriorate is shown to recover most of the OPD benefit at a fraction of the rollout length. The three methods occupy distinct points on the granularity spectrum, namely modality-specific token weighting (Visual-Advantage), dual-granularity trajectory-plus-token reweighting (Filter-Then-Reweight), and horizon-level rollout truncation (Less-is-More).

The token and sample weighting methods above operate *reactively*, attenuating the loss signal *after* the student has already generated a potentially uninformative rollout. A more proactive approach asks whether we can select prompts *before* generation to maximize the expected information gain per rollout, motivating the curriculum methods below.

6.2 Curriculum and Difficulty Adaptation

The sample weighting problem motivates a parallel solution. *Curriculum design* actively selects prompts matching the student’s current competence level. Where weighting operates

on the loss *after* generation, curriculum design intervenes *before* generation by choosing which prompts to present. The design space spans multiple axes, including prompt difficulty, temporal depth within multi-turn interactions, semantic granularity of supervision, and initialization strategy, with each axis targeting a distinct failure mode of naive uniform sampling.

Competence-boundary sampling. PACED (Xu et al., 2026b) models prompt difficulty dynamically using a Beta-kernel distribution centered on the student’s current pass rate. The curriculum shifts toward harder prompts as the student improves, while avoiding prompts where the pass rate is near zero (where gradient SNR vanishes) or near one (where the student has already mastered the content). This “frontier sampling” strategy concentrates the gradient budget on the narrow band of prompts where learning tends to be most efficient.

The mathematical justification connects to the gradient SNR analysis. For a prompt with pass rate p , the expected gradient magnitude under binary reward vanishes at both extremes ($p \rightarrow 0$ and $p \rightarrow 1$) and peaks at intermediate difficulty. PACED formalizes this through a general Beta-kernel weight family $w(p) = p^\alpha(1-p)^\beta$, where the symmetric case $\alpha = \beta = 1$ yields the familiar $p(1-p)$ with peak at $p^* = 0.5$, and asymmetric choices ($\alpha \neq \beta$) shift the peak to $p^* = \alpha/(\alpha + \beta)$, allowing the curriculum to be tilted toward harder problems ($\alpha < \beta$, peak left of 0.5) or easier problems ($\alpha > \beta$, peak right of 0.5) when prior knowledge of the pass-rate distribution motivates such tilting. PACED proves this Beta family is the leading-order minimax-optimal weighting under bounded misspecification of the SNR model, and shows empirically that the symmetric default $w(p) = p(1-p)$ already captures most of the achievable gain. Concretely, PACED estimates per-prompt difficulty $\hat{p}_i$ via a one-shot evaluation phase (typically $K=8$ rollouts per prompt) and weights each prompt by $\hat{p}_i(1-\hat{p}_i)$, concentrating training on the productive learning range while avoiding already-mastered or intractable problems. The single-pass estimation suffices because the Beta kernel is minimax-robust to stale pass rates, though iterative recomputation can yield modest additional gains.

This SNR analysis has a deeper implication that extends beyond curriculum design, as it explains *why* on-policy training tends to outperform off-policy approaches on reasoning tasks without explicit curriculum engineering. In off-policy settings, the problem difficulty distribution is fixed by the dataset. As the student improves, an increasing fraction of training data falls into the “too easy” zone where gradients are negligible, leading to diminishing returns. On-policy training, by contrast, automatically adjusts the effective difficulty because the student’s rollouts reflect its current capabilities, naturally generating sequences at the boundary of its competence. PACED makes this automatic adjustment explicit and tunable, but even naive on-policy training captures much of the benefit simply by sampling from p_θ rather than p_{data} .

Self-adjusting curricula. Several methods from earlier sections supply implicit curriculum mechanisms without explicit difficulty estimation. π -Play’s multi-agent framework (Section 5.2.3) automatically generates harder tasks as the student strengthens, PAINT (Tan & Hong, 2026a) (Section 5.2.3) adapts its partial-solution disclosure based on student-teacher overlap, and PRISM (Wang et al., 2026g) (Section 5.1.2) shapes the policy landscape for subsequent RL through a structured multimodal warm-up. These complement PACED’s explicit difficulty estimation with emergent difficulty adaptation.

Dual-perspective curriculum. Where PACED treats the student as a black-box function whose pass rate fully characterizes prompt difficulty, Uni-OPD (Hou et al., 2026) argues that pass rate alone conflates two distinct failure modes, problems the student cannot solve (low pass rate due to capability gaps) and problems the student solves incorrectly in subtle ways (moderate pass rate but low-quality solutions). Separating these requires examining not just *whether* the student succeeds but *how* it fails. From the *student* perspective, Uni-OPD implements difficulty-aware and correctness-aware data balancing that promotes exploration of informative states during rollout sampling, extending PACED’s frontier-sampling insight from a single pass-rate metric to a dual-criterion that jointly considers problem difficulty and rollout correctness. From the *teacher* perspective, it tackles a subtler failure mode. Even when the student generates informative rollouts, aggregated token-level

supervision can become inconsistent with outcome-level reward signals. A trajectory that receives higher token-level teacher scores may actually produce a worse final answer. Uni-OPD’s outcome-guided margin calibration restores order consistency between correct and incorrect trajectories, so that the token-level gradient direction aligns with the sequence-level quality ordering. This dual-perspective framework generalizes across LLMs and MLLMs (5 domains, 16 benchmarks) and accommodates both single-teacher and multi-teacher settings, implying that the curriculum and supervision reliability problems are largely orthogonal and can be solved independently.

Temporal curriculum for multi-turn distillation. PACED, Uni-OPD, and the self-adjusting methods all operate along the *difficulty* axis, selecting *which prompts* to train on. TCO (Wang et al., 2026d) opens an orthogonal axis by asking not *which prompts* but *how deep* within a single prompt’s trajectory to apply supervision, tackling a failure mode unique to multi-turn agent distillation. In standard OPD applied to multi-turn tasks, the student generates entire interaction trajectories and receives teacher supervision at every turn. TCO identifies “Trajectory-Level KL Instability,” where as conversation depth increases, KL divergence between student and teacher distributions escalates due to compounding errors in early turns that shift the state distribution, making subsequent teacher supervision increasingly misaligned with the student’s actual trajectory context. The proposed temporal curriculum progressively expands the horizon of teacher supervision rather than applying it uniformly across all turns from the start. Two scheduling variants target complementary failure modes. Forward-to-Backward (F2B) begins supervision at early turns and progressively extends to later turns, allowing the student to first establish reliable early-turn behavior before tackling the compounding-error regime. Backward-to-Forward (B2F) uses the teacher to generate early turns (supplying a clean state distribution) and applies supervision only at later turns, progressively requiring the student to handle earlier turns independently as training progresses. On ALFWorld, WebShop, and ScienceWorld, TCO delivers gains of up to +18 points over vanilla multi-turn OPD, with the F2B variant showing particular strength on tasks where early-turn errors are most consequential. The temporal curriculum insight connects directly to the gradient SNR framework developed above. Just as PACED identifies that prompts with near-zero pass rates yield negligible gradients, TCO identifies that later turns in long trajectories suffer analogous gradient degradation because the student’s state has drifted too far from where the teacher’s supervision is meaningful.

Off-policy cold start. Li et al. (2026j) identify that OPD can fail catastrophically when the student’s initial policy is too far from the teacher’s distribution (the “thinking-pattern mismatch” problem). Their proposed solution is an off-policy cold-start phase, a brief period of standard SFT on teacher-generated data that reduces the pattern gap before switching to on-policy training. This two-phase approach, off-policy warmup followed by on-policy refinement, has become a common industrial pattern (Section 8).

The hybrid SFT+OPD pipeline. The off-policy cold start (Li et al. (2026j)) combined with on-policy refinement has emerged as a common industrial recipe. Empirically, a useful transition point from off-policy to on-policy training appears to be when the student’s generation distribution reaches sufficient overlap with the teacher’s. Transitioning too early can waste on-policy compute on uninformative rollouts (the student is too weak for meaningful self-generation). Transitioning too late tends to leave the student stuck near the off-policy ceiling. Qwen3 (Yang et al., 2025a) follows this two-phase approach, using off-policy distillation to establish basic reasoning and mode-switching capabilities before transitioning to on-policy distillation where the student generates its own sequences for logit-level alignment. DeepSeek-R1 (DeepSeek-AI et al., 2025) uses 100% off-policy (no transition), which works well at extreme teacher capability but may leave performance on the table for reasoning tasks where the student could exceed the teacher through on-policy exploration.

Data freshness and teacher-refresh curricula. Curriculum design generalizes beyond prompt difficulty to the data, the rollout staleness, and the teacher-refresh schedule. DeltaPrompts (Jung et al., 2026) identifies the *zero-delta trap* in VLM distillation, where a substantial fraction of training prompts elicit indistinguishable teacher and student outputs and therefore produce zero OPD signal. Synthesizing high-divergence prompts that maximize teacher-student disagreement under a controlled difficulty band converts the dataset into an

active OPD curriculum that allocates compute to informative examples. f-OPD (Chen et al., 2026c) addresses the freshness side of the curriculum. In long-horizon asynchronous OPD pipelines (where rollout generation, teacher scoring, and student updates run on separate workers), staleness builds up and destabilizes the loss. A freshness-aware control framework bounds the lag between the rolling-out policy and the updating policy, scheduling rollouts according to a freshness budget rather than a fixed cadence and restoring the stability that purely synchronous pipelines provide at lower throughput. Trust-Region Behavior Blending (Plyusov et al., 2026) brings the same insight into the warmup phase. A teacher-guided behavior policy under a student-centered KL constraint blends the teacher and student token distributions in closed form during early training, and an annealing schedule retreats the blend toward pure on-policy as the student’s policy converges. The construction recovers the conventional OPD optimum at full annealing while supplying the implicit curriculum that practitioners would otherwise build manually. Adaptive Teacher-Refresh (Guo et al., 2026) then asks when in self-OPD the teacher should move at all. Three gates (an isolation gate enforcing a minimum freeze period, a reward-ratchet gate requiring measurable improvement, and a length-tail gate guarding against degenerate sequence-length explosions) condition each hard refresh of the self-teacher, supplying a principled answer to a hyperparameter that is usually tuned by trial and error. Continual Experience Internalization (Chen et al., 2026a) closes the loop for self-evolving agent settings. Principle-level experience extraction, step-wise injection, and off-policy context-distillation with rejection sampling are interleaved into an iterative self-evolution loop that keeps the agent stable across many internalization rounds. Together, the five methods extend curriculum design from *which prompt* (PACED, Uni-OPD) and *how deep within a trajectory* (TCOD) to *which data, how fresh, how blended, when to refresh, and how to internalize*, surfacing the temporal and structural axes of OPD curricula that earlier work treated as fixed.

Trajectory-refined distillation via prefix-failure correction. TRD (Jiang et al., 2026a) adds a further axis by asking not just which prompts or how deep, but *at which point within a corrupted rollout* the teacher should intervene. During on-policy generation, TRD monitors the student’s trajectory for prefix-failure events, positions where the student’s accumulated token sequence has drifted so far from the teacher’s manifold that continued teacher supervision would be applied over an off-distribution prefix. Rather than discarding the whole rollout or applying uniform supervision, TRD inserts teacher-generated corrections at the detected failure points, concatenates the corrected suffix back into the rollout, and applies distillation on the corrected trajectory. This trajectory-refinement step converts low-quality rollouts that would otherwise contribute noisy or zero-gradient signal into usable training examples, improving both sample efficiency and training stability without requiring additional teacher queries beyond the correction inference calls.

6.3 Compute Optimization

Token weighting (Section 6.1) reduces the effective gradient budget wasted on uninformative tokens, and curriculum design (Section 6.2) avoids generating uninformative rollouts altogether. Together they improve *sample efficiency*, extracting more learning per rollout. Yet neither removes the underlying *systems-level* bottleneck. OPD requires generating student rollouts *and* scoring them with the teacher at every training step, creating substantial compute overhead relative to off-policy SFT regardless of how efficiently each rollout is utilized. The methods below attack this orthogonal cost dimension.

The cost breakdown for a typical white-box OPD step consists of

1. **Student rollout:** autoregressive generation of $B \times K$ completions (batch size B , K rollouts per prompt), memory-bound due to KV cache growth. In practice this component typically dominates wall-clock time, since generation is sequential and cannot be trivially parallelized across tokens.
2. **Teacher scoring:** forward pass through the teacher to obtain full-vocabulary logits on student-generated sequences, compute-bound. When the teacher is considerably larger than the student, this pass can rival or exceed the student’s backward pass in cost.

3. **Student update:** standard backward pass computing gradients of the distillation loss, dominated by optimizer state memory. This component is usually the cheapest of the three but sets the lower bound on memory usage.

Each component admits independent optimization, and the methods below attack different bottlenecks.

Prefix truncation. FOPD (Zhang et al., 2026b) observes that the useful distillation signal (measured by per-token reverse-KL loss) is concentrated in the prefix of student-generated sequences, conjecturing that the student is weakest at high-level planning decisions captured in early tokens. Later tokens contribute diminishing marginal signal as the generation becomes increasingly constrained by the prefix context. Exploiting this observation, FOPD truncates on-policy rollouts to a prefix of length k and reverts to off-policy KD for the remainder. On-policy prefix distillation (FOPD) matches the performance of full OPD while reducing training FLOP by $2\times\text{--}47\times$. The optimal truncation point k depends on the information density profile of the task. Using Qwen3-8B as teacher and Qwen3-8B-Base as student, a prefix scheduling strategy that progressively increases the training prefix from 1 token by a fixed increment ($\Delta L = 256$) per gradient step closely matches full OPD quality while lowering compute.

Offline teacher caching. Lightning-OPD (Wu et al., 2026a) decouples student optimization from teacher inference entirely by precomputing teacher log-probabilities once over fixed SFT rollouts, eliminating the need for a live teacher server during training. The enabling condition is *teacher consistency*, which requires that the same teacher model be used for both the SFT stage (generating training trajectories) and the OPD stage (providing the reference distribution). Under teacher consistency, Lightning-OPD provably shares the same optimum as standard online OPD, with bounded gradient discrepancy and an implicit regularization effect that prevents policy drift without explicit KL penalties. The framework achieves $4.0\times$ higher training efficiency than standard OPD, substantially lowering the barrier for academic research.

These two methods expose a deeper structural point. The “on-policy” requirement is not binary but admits degrees of approximation. FOPD shows that partial rollouts preserve most of the distributional alignment benefit. Lightning-OPD shows that fixed (stale) teacher signals remain effective when teacher consistency holds and the implicit regularization bounds policy drift. Together, they define a *fidelity-efficiency frontier* along which the degree of on-policy approximation trades off against compute cost, with full on-policy training at one extreme (maximum alignment, maximum cost) and fully offline SFT at the other (minimum cost, exposure bias ceiling).

Speculative knowledge distillation. DistillSpec (Zhou et al., 2024) first applied on-policy KD to improve the *draft* model used by speculative decoding at *inference* time, training the draft to better match the target model’s distribution and thereby raising the acceptance rate of speculatively proposed tokens. SKD (Xu et al., 2025c) subsequently adapts the speculative-decoding mechanism to the *training* side, where the student generates candidate tokens that the teacher verifies and selectively replaces based on its distribution, creating a data generation strategy that dynamically interpolates between supervised and on-policy KD. When student samples are consistently rejected, SKD degenerates to supervised KD. When they are consistently accepted, it degenerates to on-policy KD. The expected wall-clock speedup of speculative decoding is

$$E[S] = \frac{\tau(x)}{c\gamma + 1} \quad \text{where } c = \frac{\text{cost(draft forward)}}{\text{cost(target forward)}}$$

with $\tau(x)$ denoting the block efficiency (expected tokens generated per verification step) and γ the block size. DistillSpec improves $\tau(x)$ by aligning the draft model’s distribution with the target model’s, yielding 10%–45% speedup improvements over standard speculative decoding. Combined with lossy speculative decoding under both greedy and non-greedy sampling, DistillSpec also enables fine-grained control over the latency-performance trade-off. In practical multi-scale deployments, first distilling a stronger target model and then applying DistillSpec to train a well-aligned draft model can reduce decoding latency by $6\text{--}10\times$ with minimal performance drop relative to standard decoding without distillation.

GPU memory management. Beyond FLOPs, the most severe constraint for white-box OPD is GPU memory. Holding a 70B teacher alongside a 7B student requires the teacher’s weights (~ 140 GB in BF16), the student’s weights and optimizer states (~ 84 GB), and the full-vocabulary logits tensor $[B, T, |V|]$ for distillation. Peak memory can easily exceed an 8×80 GB H100 node. Practitioners alleviate this via (1) *teacher quantization*, serving the teacher in FP8 or INT4 for a $2\text{--}4\times$ memory reduction at negligible precision loss, (2) *logit offloading*, immediately streaming logits to CPU RAM or retaining only the top- k entries (since 99% of the probability mass is typically concentrated in fewer than 100 tokens), and (3) *aggressive gradient checkpointing* to minimize student activation memory during the backward pass. Combining these techniques is typically required for viable white-box OPD on standard clusters.

Concrete cost example. To ground these formulas, consider distilling a 70B teacher into a 7B student on $8 \times$ H100 GPUs (representative estimates based on scaling from reported benchmarks). Off-policy over 1B tokens involves three stages. The teacher generates the dataset offline (~ 200 GPU-hours), the student trains for ~ 100 GPU-hours, totaling ~ 300 GPU-hours. On-policy over 1B tokens, each step requires (1) student generation ($\sim 3\times$ forward-pass cost due to autoregressive decoding), (2) teacher scoring (one 70B forward pass), and (3) student backward pass, yielding $\sim 1,200\text{--}1,500$ GPU-hours, a $4\text{--}5\times$ overhead consistent with empirical measurements at smaller scales (Wu et al., 2026a). For reasoning tasks, on-policy methods tend to outperform off-policy SFT, and the off-policy performance ceiling is lower because static datasets cannot cover the combinatorial space of valid reasoning paths. This “ceiling gap” widens with task complexity, providing a strong cost-benefit justification for on-policy methods. Section 7.4 formalizes this cost-quality tradeoff into a decision framework.

Asynchronous generation-training decoupling. NPD (Rang et al., 2026) attacks a major compute bottleneck of OPD, the synchronous coupling between student generation and gradient updates that forces the GPU to alternate between memory-bound decoding and compute-bound backpropagation. NPD reformulates on-policy distillation as a three-stage asynchronous pipeline. Student rollouts are batch-generated via vLLM, teacher top- k logits are computed through parallel prefill with sequence packing, and the student trains with a composite CE+KD loss on packed sequences. The principal challenge in async OPD is policy lag, where the generation policy drifts from the training policy across async iterations, potentially degrading into off-policy learning. NPD counters this with Δ -IFD filtering, which estimates teacher-student discrepancy via an Instruction-Following Difficulty metric and rejects extreme out-of-distribution samples that would push updates outside a “safe proximal learning zone.” The resulting framework records $8.1\times$ throughput speedup over synchronous on-policy baselines while outperforming SFT by $+8.09\%$ averaged across 11 evaluation benchmarks spanning general, math, reasoning, and code tasks. NPD also shows that on-policy distillation produces superior initializations for subsequent RL, narrowing the exploration space for GRPO. The NPD $\rightarrow$ GRPO pipeline enables openPangu-Embedded-1B to reach 68.73% average accuracy, surpassing the larger Qwen3-1.7B (63.69%). The relationship between NPD and the other compute methods is one of orthogonal scope. FOPD reduces the *length* of each rollout, Lightning-OPD removes the *freshness* requirement for teacher signals, and NPD removes the *synchronization* requirement between generation and training. All three can be composed within a single system.

Parameter-dynamics foresight. EffOPD (Cai et al., 2026b) takes a different angle on efficiency, asking *why* OPD converges faster than RL at the parameter level. The analysis shows that OPD’s efficiency stems from a form of “foresight”: it establishes a stable update trajectory toward the final model early in training. At the module-allocation level, OPD identifies regions with low marginal utility and concentrates updates on modules most relevant to reasoning. At the update-direction level, OPD exhibits stronger low-rank concentration, with dominant subspaces aligning closely with the final update subspace early in training. Building on these findings, EffOPD proposes a plug-and-play acceleration method that adaptively selects an extrapolation step size and moves along the current update direction, achieving $3\times$ training acceleration while maintaining comparable final performance without additional trainable modules or complex hyperparameter tuning.

Drift-aware dynamic truncation. Prune-OPD (Yang et al., 2026g) attacks a subtler source of compute waste that arises specifically in long-horizon reasoning. As the student generates extended sequences, its prefix progressively diverges from the teacher’s reasoning path, making subsequent teacher rewards increasingly unreliable. Rather than imposing a fixed truncation length (as in FOPD), Prune-OPD continuously monitors local student-teacher compatibility via per-position top- k overlap ratios. When this metric falls below a threshold, a *prefix-drift event* is recorded and subsequent OPD rewards are monotonically down-weighted. Once cumulative drift exceeds a budget, the rollout is truncated entirely, reallocating compute strictly to regions of reliable teacher supervision. Across diverse teacher-student pairs (DeepSeek-R1-Distill-Qwen-7B→R1-Distill-Qwen-1.5B, JustRL-DeepSeek-1.5B→R1-Distill-Qwen-1.5B, Qwen3-4B→Qwen3-1.7B/4B-Base), Prune-OPD reduces training time by 37.6–68.0% while preserving or improving performance on AMC, AIME, and HMMT. The adaptive behavior is key. When student-teacher compatibility remains high, Prune-OPD automatically preserves long-context supervision by expanding the training window, avoiding the information loss of fixed truncation. FOPD and Prune-OPD thus represent complementary truncation philosophies. FOPD uses a *scheduled* prefix that grows uniformly regardless of compatibility, while Prune-OPD uses a *reactive* budget that contracts or expands based on real-time drift detection, better suited for long-horizon tasks where drift timing varies across instances.

Rollout truncation and localized distillation. Two further compute-side innovations target the rollout cost itself rather than the systems-level memory wall. POPD/TOPD (Zhang et al., 2026k) examines whether full rollouts are necessary at all. Progressive Horizon Expansion (POPD) starts with a short rollout horizon and progressively grows it, allowing early training to focus on the prefix where signal density is highest (consistent with FOPD’s prefix-truncation finding) and later training to extend coverage as the student stabilizes. Truncated Rollout Distillation (TOPD) instead trains on a single short log-ratio horizon throughout, trading sequence coverage for per-step cost reduction. SafeSteer (Li et al., 2026e) attacks the cost from a different direction by confining the OPD update to a small set of safety-relevant tokens. An activation-steered safety teacher supplies the per-token reverse-KL signal only on tokens flagged by contrastive log-probability voting as safety-critical, leaving the rest of the rollout untouched. The result is a localized distillation procedure whose compute cost scales with the safety-token fraction rather than the full rollout length, illustrating how task-specific structure can recover the efficiency that prefix-truncation methods achieve through generic horizon control. KAT (Xin et al., 2026) addresses a different failure mode by identifying *KL-agreement trap* regions, contiguous token spans where the student’s and teacher’s top-1 predictions agree but the full distributions differ substantially, creating a deceptive gradient landscape that wastes compute on low-value updates. An online rollout monitor detects these trap regions in real time and truncates the current rollout at the trap entry point, reallocating the generation budget to fresh rollouts where student-teacher divergence is genuine and informative. The combination of real-time detection and online truncation requires no additional forward passes and integrates directly into the existing OPD training loop, complementing Prune-OPD’s drift-based budget with an agreement-trap-based criterion that targets a structurally different source of compute waste.

The training efficiency stack. The three subsections above form a layered efficiency stack. Token weighting (Section 6.1) reduces noise at the gradient level, so that each rollout contributes more effectively to learning. Curriculum design (Section 6.2) reduces waste at the sampling level, so that each rollout targets the student’s competence frontier. Compute optimization (this section) reduces overhead at the systems level, lowering the cost of producing and scoring each rollout. These layers are complementary rather than substitutive. PACED’s frontier sampling generates maximally informative prompts, TIP’s token weighting extracts maximal signal from the resulting rollouts, and FOPD’s prefix truncation minimizes the generation cost of those rollouts. Deploying all three layers simultaneously can compound their individual 2–4× efficiency gains, potentially reducing the effective cost gap between on-policy and off-policy training from 4–5× to near parity while retaining the quality ceiling advantage.

7 Understanding OPD: Theory, Failure Modes, and Cost

The preceding sections addressed the mechanics of on-policy distillation, covering what to optimize (Section 4), where the signal comes from (Section 5), and how to stabilize training (Section 6). This section turns to the equally important questions of *why* and *when*: under what conditions does OPD succeed, how does it fail, what theoretical frameworks unify the methods surveyed above, and when is OPD preferable to simpler off-policy pipelines?

These four dimensions of understanding are not independent but form a cumulative argument. Identifying *when* OPD works (Section 7.1) establishes the preconditions that the theoretical framework must explain. Cataloging *how* it fails (Section 7.2) motivates the specific design choices, adaptive divergences, curriculum weighting, and hybrid RL objectives, whose theoretical grounding Section 7.3 supplies. Both the success conditions and the failure modes then inform the practical cost-benefit analysis in Section 7.4, which distills the theoretical insights into engineering guidelines. Resolving these questions moves OPD from a largely empirical practice toward a more systematic engineering discipline (Zhang et al., 2026f; Li et al., 2026j).

7.1 Success Conditions

Li et al. (2026j) present a systematic investigation of OPD training dynamics, identifying two necessary conditions for effective distillation. First, the student and teacher must share compatible reasoning patterns (operationalized as high overlap in their top- k token distributions). Even a superior teacher fails to transfer knowledge if its thinking style is structurally mismatched with the student’s (e.g., a non-thinking teacher distilling into a thinking student), because low initial distributional overlap cannot be recovered through training alone. Second, the teacher must supply new capabilities beyond what the student has already acquired. When both models are trained on the same data and recipe, they converge to similar distributions at their respective scales, leaving the teacher with little transferable signal, as reverse distillation experiments show that same-family 1.5B and 7B teachers become distributionally indistinguishable from the student’s perspective.

At the token level, successful OPD is characterized by progressive alignment on *high-probability overlap tokens*, a small shared token set that concentrates the bulk of the probability mass. When these conditions fail, recovery strategies include off-policy cold start (SFT initialization that reduces the pattern gap) and teacher-aligned prompt selection. A persistent challenge is that the quality of OPD’s dense supervision degrades systematically with trajectory depth, with instability originating at later tokens and propagating backward, raising open questions about OPD’s applicability to long-horizon reasoning.

These findings point toward a general principle. OPD’s benefit scales with the *exploitable gap* between teacher and student. When this gap is too small (same-recipe models), the teacher offers little new information. When too large (thinking-pattern mismatch), the student cannot absorb the supervision. The productive regime lies in between, and identifying it remains a key diagnostic challenge.

Training-free per-token diagnostics. Armandpour et al. (2026) introduce a complementary diagnostic framework that operates at the highest resolution, per token, per question, per teacher. They derive an ideal per-node gradient defined as the parameter update that maximally increases the student’s probability of success, and develop a scalable targeted-rollout algorithm to estimate it efficiently even for long chains of intermediate thoughts. The gradient alignment score (cosine similarity between this ideal gradient and a given distillation gradient) quantifies how well a particular teacher configuration approximates the ideal signal. Across self-distillation settings and external teacher models, distillation guidance exhibits considerably higher alignment with the ideal on incorrect rollouts than on correct ones, where the student already performs well and the teacher’s signal tends to become noisy. The optimal distillation context depends jointly on the student model’s capacity and the target task, with no single universally effective configuration emerging. These findings provide a principled, training-free basis for the per-task diagnostic analyses that the success conditions above motivate.

Diagnostic checklist. Based on the surveyed literature, four pre-training diagnostics emerge as predictors of OPD benefit:

1. *Token overlap ratio*: Compute top- k overlap between teacher and student on a held-out set. If too low, apply off-policy cold start first.
2. *Pass rate on target prompts*: If the student cannot solve any problems in the target domain, OPD gradients will vanish. Use curriculum pacing (Xu et al., 2026b) or easier warm-up tasks.
3. *Teacher calibration*: Verify that the teacher’s confidence correlates with correctness on student-generated prefixes (Zhang et al., 2026d). Poorly calibrated teachers inject noise.
4. *Length stability*: Monitor output length during early training. Abrupt length inflation indicates the self-reinforcing repetition cycle identified by Luo et al. (2026a). Apply Stable-OPD corrections preemptively.

OPSD as compression, not correction. Kim & Lee (2026) sharpen the characterization of *what* on-policy self-distillation actually accomplishes in thinking-enabled mathematical reasoning. By applying OPSD separately to correct-only versus incorrect-only rollout groups, they isolate two potential mechanisms. Correct-only OPSD preserves accuracy while substantially shortening reasoning traces (compression), whereas incorrect-only OPSD degrades accuracy (correction fails). Three alternative explanations (richer teacher context, mid-trace feedback reinjection, and extended training) are tested and ruled out, supporting the interpretation that the hindsight-conditioned self-teacher can identify redundancy in long thinking traces but does not reliably supply better alternative reasoning steps. This finding refines the success conditions above. OPSD’s productive regime appears to be “making the model express known solutions more efficiently” rather than “making the model solve harder problems”. The practical implication is a refined pipeline ordering. SFT establishes format compliance, RLVR expands the set of reachable correct trajectories, and OPSD compacts them for cheaper inference. Reversing the last two stages (compressing before exploration) tends to remove the raw material that exploration would produce, suggesting why post-RL placement is preferred. Across Qwen3-8B and AceReason-Nemotron-7B (with DeepSeek-R1-Distill-Qwen-7B as a supplementary pre-RL reference), the post-RLVR + correct-only OPSD configuration reaches the most favorable position in the accuracy-length plane in their experiments, consistent with OPSD’s mechanistic role rather than implementation details.

7.2 Failure Modes

Complementing the success conditions, a growing body of work catalogs the specific ways OPD can fail even when structural preconditions hold, as artifacts of the on-policy training dynamics themselves. Understanding these failure modes helps debug OPD pipelines, and each mode has a corresponding mitigation strategy developed in the methods sections above.

We organize failure modes by their *root cause* instead of their symptom, as multiple failure modes can produce similar surface-level degradation (e.g., accuracy collapse) through different mechanisms.

The flawed prefix trap. As detailed in Section 6.1, when the student generates an erroneous prefix, the teacher’s conditional distribution over subsequent tokens is poorly calibrated because the teacher was never trained on such out-of-distribution inputs. Fu et al. (2026b) formally characterize three specific failure modes, including (1) an imbalanced one-token signal that reduces distribution matching to a single sample, (2) unreliable teacher guidance on out-of-distribution student prefixes, and (3) tokenizer mismatch distortions. The theoretical analysis identifies a bias-variance tradeoff, where token-level OPD is biased relative to sequence-level Reverse KL but enjoys tighter worst-case variance bounds.

Extrapolation cliff in structured outputs. ListOPD (Li et al., 2026h) identifies a sharp failure boundary specific to reward-extrapolation OPD ($\lambda > 1$) on structured-output tasks. In a single-position Bernoulli reduction, the authors derive a closed-form base-relative clip-safety threshold $\lambda^*(p, b, c)$ determined by three measurable quantities, namely the teacher

modal probability, the warm-start mass, and the importance-sampling clip strength. Above λ^* , the extrapolated fixed point exits the clip-safe region, switching training from format-preserving to format-collapsing. On calibrated listwise JSON tasks (Amazon Fashion), three pre-registered tests corroborate the predicted cliff interval within locked prediction windows. Operating just below λ^* , ListOPD brings a 1.7B Qwen3 student to in-domain parity with an 8B-SFT baseline at one-fifth the parameters, with the gain driven primarily by format adherence rather than ranking quality. The extrapolation cliff connects directly to G-OPD’s reward extrapolation mechanism (Section 4.3), providing a formal characterization of *when* extrapolation beyond the teacher becomes destructive rather than productive.

Persistent high-loss tokens as structural residuals. Jiang et al. (2026b) investigate a related failure phenomenon. Even after OPD training reaches apparent saturation, a substantial subset of tokens (up to 18%) continues to exhibit persistently high loss. These “Rock Tokens” resist teacher-driven corrections despite providing a disproportionately large share of total gradient norms due to their high occurrence frequency. Causal intervention shows that Rock Tokens provide negligible functional contribution to actual reasoning performance, suggesting they represent structural and discourse residuals (connectives, formatting, hedging) that the student model cannot or need not internalize from the teacher. The finding implies that a substantial fraction of OPD’s optimization bandwidth is spent on tokens that neither improve nor inform the student’s reasoning capability, complementing TIP’s token-weighting framework (Section 6.1) with a characterization of *which* tokens are structurally uninformative rather than merely currently uninformative.

Comprehensive failure analysis across settings. Zhu et al. (2026) present a systematic empirical study across both OPD and OPSD, cataloging when each works and when each fails. Three failure mechanisms emerge: (1) distribution mismatch between teacher and student caused by conditioning on student-generated prefixes (the flawed prefix trap above, now measured across diverse settings), (2) optimization instability from biased top-K reverse-KL gradients that standard implementations silently introduce, and (3) an OPSD-specific limitation where the student learns a PI-free policy that aggregates PI-conditioned teachers, which is insufficient when PI is instance-specific (such as a ground-truth solution) but effective when PI represents a shared latent rule (such as a system prompt). Stop-gradient top-K objectives, RLVR-adapted teachers, and SFT-stabilized students mitigate these failures. The finding that OPSD’s effectiveness depends on the *type* of privileged information (shared rule vs. instance-specific answer) refines the PI taxonomy in Section 5.2.1 by exposing a structural limitation that curriculum or signal-augmentation patches may not fully address.

Two subtler failure modes compound the prefix trap. *Gradient SNR collapse* occurs when the student’s pass rate on a prompt approaches zero. In this state, all rollouts contain early errors, teacher signals are dominated by noise, and the gradient SNR vanishes. This helps explain why hard prompts yield no learning signal in standard OPD and motivates PACED’s competence-boundary curriculum (Section 6.2). *Epistemic suppression* (Kim et al., 2026b) manifests in self-distillation. When the process shortens reasoning traces, it disproportionately removes hedging phrases and uncertainty markers that serve as implicit attention flags for difficult sub-problems. This produces shorter but less calibrated reasoning where the model confidently commits to flawed steps. This calibration degradation connects directly to the teacher uncertainty challenge, where an overconfident student produces unreliable self-distillation targets, creating a self-reinforcing cycle that compounds with each iteration.

Local teachability collapse. Liu et al. (2026b) identify a failure mode specific to strong-to-weak OPD that operates at a finer granularity than the flawed prefix trap. Even when the teacher’s feedback remains nominally available throughout a student-generated trajectory, later segments progressively lose the *local contrast* that makes dense supervision effective for learning. The teacher’s margin over the student’s top-K candidate set narrows to the point where token-level KL gradients carry negligible discriminative information, a condition the authors term local teachability collapse. This differs from the flawed prefix trap (where the teacher’s distribution is corrupted by out-of-distribution prefixes) in that the teacher remains well-calibrated but its signal becomes insufficiently informative. The proposed release rule aggregates the teacher’s discriminative margin across sentence segments and truncates OPD supervision at a BIC-style downward change point, concentrating learning

on trajectory regions where teacher feedback remains actionable. Across strong-to-weak distillation tasks using the Qwen3 model family, this truncation outperforms standard full-trajectory OPD on five in-domain benchmarks and better preserves out-of-domain capabilities. The finding refines FOPD’s prefix-truncation strategy (Section 6.3), which focuses on early tokens based on a fixed positional heuristic. Local teachability collapse suggests that the optimal truncation point is trajectory-specific and should be determined by the teacher’s local discriminative capacity rather than token position alone.

Self-play saturation (the Ouroboros problem). In self-distillation settings, the student optimizes against a target sharing its own inductive biases and architectural limitations, causing the hypothesis space to gradually collapse. If the model discovers a syntactic “hack” or a highly confident but flawed reasoning heuristic, no external signal penalizes it. The model self-reinforces this flawed trajectory, driving $p_\theta(y_{\text{flawed}}|x) \rightarrow 1$. Once entirely certain of its own hallucinations, gradients vanish, exploration ceases, and the policy is trapped. This is distinct from exposure bias (which arises from train-test distribution mismatch in off-policy settings). Saturation arises from the absence of distributional diversity in fully on-policy self-distillation, bearing structural similarities to GAN mode collapse, where a generator locks onto a narrow output manifold once the discriminator signal becomes non-informative.

The precision-recall tradeoff. Cha & Cho (2025) formalize a core tension in distillation. As the student concentrates mass on high-quality outputs (improving precision), it necessarily reduces coverage of the teacher’s full output distribution (reducing recall). In the OPD setting, this manifests as the “diversity collapse” problem. Aggressive Reverse KL distillation produces a student that scores highly on Pass@1 (precision) but catastrophically on Pass@ k (recall), because the student has collapsed onto a single reasoning strategy per prompt. On-policy generation can shift this tradeoff by allowing the student to explore its own distribution instead of being constrained to teacher-demonstrated paths, potentially improving recall without sacrificing precision. Temperature-scaled sampling during on-policy generation, as used in GKD and G-OPD, offers a direct control knob, with higher sampling temperatures encouraging distributional exploration (increasing recall) at the cost of noisier gradients (potentially reducing precision).

The calibration-capability gap. CaOPD (Zhang et al., 2026d) identifies a “Scaling Law of Miscalibration”. While OPD reliably improves task accuracy, the authors report that it tends to leave models in a state of severe overconfidence. They attribute the root cause to an information mismatch. Teacher supervision is formed under privileged context available during training (full reasoning traces, multiple samples), whereas the deployed model must report confidence using only its deployment-time context. By decoupling capability improvement from calibration degradation, CaOPD argues that the distilled model becomes more capable but less aware of its own uncertainty boundaries. This finding has potential implications for deployment. A distilled model that scores higher on benchmarks but does not reliably indicate when it might be wrong is arguably less safe than a weaker but better-calibrated baseline. The “illusion of certainty” connects to the saturation analysis above. Both indicate that OPD can produce students that lose metacognitive capabilities (uncertainty awareness, epistemic verbalization) even as task-level performance improves.

Agentic collapse in multi-turn OPD. The preceding failure modes all assume a single-turn response. Extending OPD to multi-turn tool-using agents surfaces a qualitatively different failure family, in which the interaction between teacher dynamics and trajectory structure produces pathologies unseen in single-turn settings. TT-OPD (Jeong, 2026) diagnoses three such pathologies on Healthcare AI Gym, and each maps onto a distinct structural mismatch that other agentic OPD methods also confront from complementary angles. *First, teacher-dynamics collapse.* Periodic hard-copy teacher resets cause a KL collapse (divergence drops abruptly from its accumulated value to near zero, e.g., $2.637 \rightarrow 0.343$ at a single copy event with period $T=30$), erasing the distillation gradient and driving accuracy monotonically downward. This is the multi-turn analogue of TCOD’s “Trajectory-Level KL Instability” (Section 6.2), but surfaced at the teacher-update granularity rather than the turn-depth granularity, and shows that stable teacher dynamics (EMA rather than hard copy) are a prerequisite for any other structural fix. *Second, trajectory-structure erosion.* Even with EMA teachers that prevent KL collapse, the multi-turn structure itself erodes, as the student learns

that verbose single-turn monologues maximize the sparse terminal reward more easily than coordinated tool-use sequences (turns collapse from 7.65 to 5.52 per episode under periodic resets, and still erode from 7.82 to 6.23 even with EMA teachers). This mirrors the reward-granularity mismatch that motivates Skill-SD’s skill-level decomposition and MAD-OPD’s step-level sampling (Section 8), suggesting that sparse terminal rewards alone cannot sustain multi-turn structure regardless of granularity. *Third, reward-hint runaway.* Adding outcome-conditioned teacher hints without length control triggers a response explosion, where positive hints reinforce ever-longer reasoning until responses saturate the context limit and accuracy collapses from 54.5% to 49.0%. This is the multi-turn extension of the epistemic-suppression phenomenon documented in Kim et al. (2026b), as both arise when informative supervision signals that are helpful *on average* are applied without regularizers on the shape of the student’s response.

Taken together, these pathologies imply a general principle for agentic OPD. Stable teacher dynamics, explicit regularizers on trajectory structure, and granularity-matched credit assignment appear jointly necessary based on current evidence, and each existing method (TCOD, Skill-SD, MAD-OPD, TT-OPD) targets a different subset of the three.

Off-metric capability erosion. Beyond task-specific failure modes, Wang (2026) argue that distillation evaluation itself is structurally incomplete. Current practice reports what students retain (task scores) without accounting for what they lose (off-metric teacher capabilities that make those scores reliable). Reframing distillation as a lossy projection, they synthesize existing evidence into a taxonomy of off-metric losses, including calibration degradation (the student scores well but does not know when it is wrong), reasoning trace quality erosion (shorter traces that lose intermediate verification steps), and distributional narrowing (mode collapse invisible to pass@1 metrics). This taxonomy cuts across the specific failure modes cataloged above and offers a unifying evaluation lens. The proposed Distillation Loss Statement (primary metric, critical off-metric capabilities, preservation targets, stress distributions, observed erosion) provides a structured template for reporting distillation outcomes that goes beyond headline accuracy, complementing the diagnostic checklist in Section 7.1 with a deployment-oriented evaluation protocol.

Exposure-bias recovery, multi-teacher interference, and gradient stability. Three further failure-mode analyses connect to the prefix-trap and saturation pathologies above but identify distinct structural causes. MOTAB (Wang et al., 2026a) catalogs *dual exposure biases* in reasoning distillation, the conventional training-inference prefix mismatch and a high-entropy-token exposure bias where the student rolls into states the teacher rarely visits. The proposed remedy monitors student on-policy trajectories and backtracks to the most recent safe state once a divergence threshold is exceeded, after which the teacher resumes corrective supervision. The construction frames exposure-bias mitigation as a recoverable operation rather than an avoidance constraint, complementing the prefix-truncation philosophies of FOPD and Prune-OPD (Section 6.3). Counteraction-Aware MOPD (Chen et al., 2026b) surfaces a multi-teacher analogue of the catastrophic-forgetting problem during domain specialization. Vanilla multi-teacher OPD silently degrades general capabilities when the prompt distribution mismatches each teacher’s expertise, because conflicting per-token signals from teachers with non-overlapping competence regions cancel rather than compose. A counteraction-aware router detects these conflicts and allocates supervision to avoid destructive interference, recovering general capabilities while preserving domain performance. Physics-Guided Self-Distillation (Wang et al., 2026f) attacks the gradient-stability side of the failure landscape by modulating the per-step update size with an information-theoretic signal. A mutual-information estimate between the student’s predictions and a feedback-conditioned teacher acts as a viscosity-like multiplier, bounded by an exponential cap, that contracts the step in regions of low MI (where the teacher signal is uninformative) and relaxes it in regions of high MI. The construction interprets the OPD update as a dissipative dynamical system and offers a principled alternative to fixed-step or learning-rate-only schedules. Each method addresses a different layer of the failure stack, namely trajectory backtracking (MOTAB), multi-teacher routing (Counteraction MOPD), and step-size modulation (Physics-Guided SD).

7.3 Unified Theoretical Perspectives

The success conditions and failure modes above share deeper mathematical roots than their diverse surface manifestations suggest.

The unified f -divergence framework of Wen et al. (2023) subsumes Forward KL, Reverse KL, JSD, and total variation distance as special cases parameterized by a single convex generator function f . Since all these divergences can be expressed as expectations under the student’s on-policy distribution, the choice of divergence is ultimately a regularization decision. The generator function f determines the implicit weighting of likelihood ratios $p_T(y)/p_\theta(y)$, with Forward KL up-weighting regions where the student underestimates the teacher (coverage) and Reverse KL up-weighting regions of overestimation (precision). Wu et al. (2025) refine this picture for discrete LLM distributions, showing that under practical epoch budgets the mode-seeking/covering dichotomy is better understood through head-tail focus. Forward KL concentrates learning on the head of the teacher distribution while Reverse KL emphasizes the tail, explaining why adaptive methods that switch between divergences based on local geometry tend to outperform any fixed choice in the settings tested. At a more foundational level, Cha & Cho (2025) show that KD in generative models induces a precision-recall tradeoff modulated by a single parameter, teacher entropy. As the teacher distribution becomes lower in entropy (more peaked), the distilled student concentrates probability mass on high-likelihood regions, improving sample precision at the cost of distributional coverage (recall). This tradeoff is validated in large-scale LLM experiments (SmolLM2 family), where lower-entropy teachers produce students that generate sharper, more fluent text but cover fewer modes of the output space. The connection to the “exploitable gap” identified empirically by Li et al. (2026j) is direct. The precision gain from distillation is largest when the teacher’s entropy is substantially lower than what the student would achieve from data alone, which is precisely when the teacher offers an informative compression of the output space that data-only training cannot replicate efficiently. A parallel unification comes from Li et al. (2026c), who show that GRPO and self-distillation tackle complementary failure modes and can be composed within a single framework via sample routing. Correct samples follow GRPO’s reward-aligned reinforcement path while failed samples receive dense logit-level correction from a feedback-conditioned self-teacher, evidence that the two paradigms target distinct learning signals instead of being redundant approaches to the same objective.

Beyond the divergence-level view, a parallel line of analysis asks how OPD reshapes a model in parameter space, and whether its dynamics differ from those of SFT and RLVR. Shen et al. (2026c) characterize the OPD update trajectory through a suite of parameter-space diagnostics and place it in a relaxed off-principal regime, with updates that touch fewer weights than SFT and avoid the principal singular directions more strongly, while remaining less tightly constrained than RLVR. Their central observation is subspace locking, where the cumulative updates rapidly enter a narrow low-dimensional channel that is by itself sufficient to recover OPD performance, whereas the same constraint degrades SFT. Yu et al. (2026a) reach a consistent picture across several language and vision-language model pairs, finding OPD updates that are small and coordinate-sparse yet numerically full-rank, concentrated in FFN modules and falling disproportionately on coordinates where the source weights are near zero. Training only the discovered subnetwork nearly recovers full-training performance, although the sparse support does not remove the need for adaptive optimization, since AdamW retains an edge over SGD that the dense teacher signal appears to require. Read together, these analyses suggest that dense teacher supervision does not collapse OPD into ordinary SFT-style parameter rewriting, and that the on-policy sampling leaves a geometric signature closer to RLVR than the dense objective alone would predict.

Length inflation (OPD students’ outputs growing progressively longer during training) is a persistent practical concern. Luo et al. (2026a) attribute the root cause to a feedback loop in the reverse-KL advantage. Once the student drifts toward long repetitive trajectories, the teacher’s log-probability increases faster than the student’s on those trajectories, producing large positive advantages that further encourage degenerate continuations. Their Stable-OPD framework adds a reference divergence term anchoring the student’s distribution to a pre-distillation baseline, combined with rollout mixing that blends on-policy student

generations with off-policy gold demonstrations in a decaying ratio. These stabilization mechanisms post +7.2% on average over vanilla OPD by breaking the self-amplification cycle.

G-OPD’s equivalence derivation (Section 4.3) supports the view that standard OPD can be cast as a special case of dense KL-constrained RL, bridging the distillation and RL communities and making it more likely that advances in either field transfer to the other. This connection extends to DPO-based methods. Substituting the teacher’s log-probability as DPO’s implicit reward (Section 4.3) produces the same geometric mixture targeted by GKD’s λ -mixing, relating token-level KL distillation, preference optimization, and reward-regularized RL as three optimization trajectories toward the same family of target distributions parameterized by divergence choice and supervision density.

The unified KD+RL framework of Li et al. (2025) makes this complementarity precise by decomposing the hybrid gradient into a dense KD component (stable, analytically computable) and a Monte Carlo RL component (high variance, reward-driven). KD dampens variance in policy gradient estimation while RL prevents collapse onto suboptimal teacher modes. Song et al. (2026) further show that textual feedback from a teacher can expand the effective capabilities of RL beyond what outcome-only rewards attain, offering theoretical grounding for the verbal feedback methods (Lion, OVD) surveyed in Section 5.1.2. Text feedback carries a denser reward signal than binary outcomes, allowing the student to learn from partial successes that binary verification would discard.

These perspectives suggest a unifying view. On-policy distillation, reinforcement learning, and preference optimization occupy a continuous spectrum parameterized by the density of supervision, the choice of divergence, and the source of the training signal, rather than constituting separate training paradigms. With this unified lens in place, the remaining practical question is when on-policy methods justify their additional compute cost over cheaper off-policy alternatives.

7.4 On-Policy vs. Off-Policy: A Decision Framework

The theoretical frameworks above indicate that on-policy training offers systematic advantages over off-policy alternatives in certain regimes. But theory alone cannot settle the engineering question of when this advantage justifies the additional infrastructure cost. This section distills the accumulated evidence into a practical decision framework.

The DeepSeek-R1 anomaly. DeepSeek-R1 (DeepSeek-AI et al., 2025) represents a prominent off-policy distillation success, training students from 1.5B to 70B parameters on $\sim 800K$ chain-of-thought traces via purely off-policy SFT. On AIME 2024 (pass@1), R1-Distill-Qwen-7B reaches 55.5% and R1-Distill-Qwen-32B reaches 72.6%, outperforming direct RL at the same scale (GRPO on Qwen2.5-32B-Base scores only 47.0%). Three factors plausibly explain this anomaly. An exceptionally strong 671B MoE teacher covers the problem space densely, diverse reasoning traces include backtracking and verification strategies, and relatively small student models maintain manageable effective conditional complexity under off-policy learning. The implication is that off-policy distillation often suffices when a massively capable teacher generates sufficiently diverse traces for a relatively small student. On-policy becomes more attractive as student capacity grows and diverges from the static training distribution.

The off-policy ceiling. DDT (Zhang et al., 2026f) supplies formal evidence for the generalization limits of off-policy approaches. A key factor driving on-policy generalization advantage is that generating training data from the current policy keeps learning aligned with the model’s evolving distribution. Off-policy performance is bounded by distributional mismatch that accumulates as the student’s policy drifts from the static training data, because fixed datasets are unlikely to cover the combinatorial space of valid reasoning paths. This gap widens with task complexity (more reasoning steps create more paths to cover), output diversity requirements (creative tasks vs. mathematics), and model capability (stronger students explore further from training data).

Distillation vs. direct RL. The choice between knowledge distillation and direct RL depends on three factors that interact non-trivially. Teacher quality relative to the reward signal determines whether dense per-token supervision outweighs sparse outcome rewards. Student scale modulates this tradeoff, with smaller students ($\leq 7B$) benefiting more from dense teacher supervision because their limited capacity cannot efficiently explore reward landscapes, while larger students ($\geq 32B$) have sufficient capacity for RL exploration. Target ceiling determines whether the goal is to match the teacher (distillation-bounded) or surpass it (RL-enabled). The hybrid approaches surveyed in Section 4.3 (G-OPD, KDRL, RLAD) resolve this tension by combining teacher logits as a dense variance-reducing baseline with sparse outcome rewards that drive exploration beyond the teacher’s ceiling. Empirically, this combination tends to outperform either approach alone in the settings tested so far (Yang et al., 2026e; Xu et al., 2025a; Zhang et al., 2026o).

The hybrid pipeline as standard recipe. Qwen3 (Yang et al., 2025a), DeepSeek-V4 (DeepSeek-AI, 2026), and MiMo-V2-Flash (Xiaomi LLM-Core Team et al., 2026) converge on a two-phase pattern in which off-policy SFT establishes a strong distributional foundation and on-policy OPD or RL closes the remaining gap between the student’s generation distribution and the optimal policy. The transition point depends on the student’s initial capability. Weaker students benefit from extended off-policy warmup to reduce the teacher-student pattern gap (Li et al., 2026j), while stronger students can transition earlier because their on-policy rollouts already occupy high-density regions of the teacher’s distribution.

Compute-quality tradeoff. The expected compute cost for N tokens under each regime:

$$\begin{aligned} C_{\text{off}} &\approx N \times (F_{\text{teacher}} + F_{\text{student}} + B_{\text{student}}) \\ C_{\text{on}} &\approx N \times (G_{\text{student}} + \rho F_{\text{teacher}} + F_{\text{student}} + B_{\text{student}}) \end{aligned}$$

where F, B denote forward and backward FLOPs, G is the autoregressive generation cost, and $\rho \in (0, 1]$ is the teacher supervision refresh rate (the fraction of steps requiring a fresh teacher forward pass). Because $G_{\text{student}} \gg F_{\text{student}}$, on-policy training incurs a substantial multiplier. For generic knowledge transfer where the off-policy ceiling is high, this overhead is rarely justified. For reasoning, code generation, and complex mathematics where the off-policy ceiling is low and marginal on-policy returns are high, the investment typically pays off. In practice, the most effective industrial pipelines tend to front-load cheap off-policy learning and reserve expensive on-policy compute for the final quality push (Section 8.3).

Summary decision rule. Use off-policy SFT alone when the teacher-student capacity ratio exceeds $10\times$ and the task admits bounded reasoning depth (factual QA, translation, summarization). Switch to on-policy OPD when any of three conditions holds: (1) the student exceeds $\sim 7B$ parameters and begins to explore regions uncovered by static teacher traces, (2) the target task involves multi-step reasoning where intermediate errors compound, or (3) the off-policy loss plateaus while held-out reward continues to improve under on-policy rollouts. In all other cases, the hybrid pipeline (off-policy warmup followed by on-policy refinement) tends to yield a favorable cost-quality tradeoff.

From theory to operational guidelines. The decision framework above rests on theoretical arguments and controlled experiments. Translating it into deployment practice introduces further constraints from infrastructure and student-scale interactions.

A key operational refinement concerns the choice between logit-level and reward-level supervision *within* the on-policy regime. Smaller students ($< 7B$) often benefit from dense logit supervision (GKD, DistiLLM) to bootstrap competence, while larger students ($\geq 7B$) with reliable verifiers tend to benefit from reward-guided OPD (KDRL (Xu et al., 2025a), G-OPD (Yang et al., 2026e)) that allows discovery beyond the teacher ceiling. The hybrid approach combining off-policy SFT warm-up with on-policy refinement performs well across the surveyed literature (Qwen3 (Yang et al., 2025a), MiMo-V2-Flash (Xiaomi LLM-Core Team et al., 2026)). At a deeper level, OPD and RLHF share much of their computational structure, including on-policy rollouts, per-trajectory signals, and policy gradient updates. The main distinction is the signal source (teacher distribution vs. reward model preference). Methods like G-OPD and REOPOLD (Ko et al., 2026) formalize this relationship, with practical consequence. For reward-guided OPD, much of the existing RLHF infrastructure applies directly (the rollout-score-update loop is the same), which helps explain why frameworks

like OpenRLHF (Hu et al., 2024) serve both communities. White-box logit-level OPD also requires teacher co-hosting and full-vocabulary logit transfer, an extension that existing frameworks support through their teacher-inference process groups.

8 Applications, Systems, and Emerging Domains

The preceding sections developed OPD from theory through methods to unified understanding. This section examines how these algorithmic advances translate into practice. The shift from research methods to industrial deployment indicates that the choice of OPD method is often shaped less by theoretical elegance than by practical constraints, available compute, teacher access level, multi-teacher coordination needs, system-level throughput bottlenecks, and domain-specific challenges such as cross-modal transfer and privacy. We organize industrial applications by their *deployment pattern* and emerging domains by the new challenges they introduce.

8.1 Industrial Deployment

Recent foundation models have increasingly incorporated OPD into their post-training pipelines, with four distinct deployment patterns that align with the theoretical progression surveyed above.

Two-phase distillation pipelines. The most common industrial pattern combines off-policy cold-start with on-policy refinement. Qwen3 (Yang et al., 2025a) employs strong-to-weak distillation where larger models (32B, 235B-A22B) supply logit-level supervision on student-generated on-policy sequences, with the student generating responses in either thinking or non-thinking mode. A key architectural choice is “Thinking Mode Fusion,” where the student learns to produce both short, direct answers (non-thinking mode) and extended chain-of-thought reasoning (thinking mode) by distilling from teachers operating in each mode independently. On-policy generation plays an important role here because the student’s mode-selection behavior (when to think vs. when to answer directly) is itself a learned policy that is difficult to capture with static datasets. Qwen3 reports that on-policy distillation outperforms direct reinforcement learning at roughly one-tenth of the GPU hours, and further improves pass@64 on AIME benchmarks where reinforcement learning from the same off-policy checkpoint does not. Gemma 2 (Gemma Team et al., 2024) embeds off-policy knowledge distillation directly into pre-training to bridge performance gaps across its model family. The 2B and 9B models are trained via logit-matching KD from a larger teacher (the 27B model itself is trained from scratch), with both small models distilled in parallel rather than via a 27B→9B→2B cascade. While this is not on-policy distillation per se (the student trains on fixed pre-training data rather than self-generated sequences), it represents an aggressive integration where distillation is a core training ingredient from the start, foreshadowing the post-training OPD patterns that later models adopt. MiMo-V2-Flash (Xiaomi LLM-Core Team et al., 2026) combines distillation from multiple domain-specialized teachers with Multi-Token Prediction (MTP) and hybrid attention, reaching strong performance on reasoning and agentic tasks with a 309B MoE model (15B active parameters).

OPD as model consolidation. A second pattern uses on-policy distillation for *merging* multiple specialized experts into a single deployable model. DeepSeek-V4 (DeepSeek-AI, 2026) entirely replaces the mixed RL stage of its predecessor (DeepSeek-V3) with pure multi-teacher OPD, consolidating more than ten domain-specific experts (mathematics, coding, agent, instruction-following) into a unified 1.6T-parameter model through full-vocabulary Reverse KL distillation. Computing the full $|V|$ -dimensional KL at each position (rather than the sampled-token approximation) ensures more stable gradients when merging experts with diverse specializations, while hidden-state caching with on-the-fly logit reconstruction makes this feasible across trillion-parameter teachers (Section 4.1). KAT-Coder-V2 (Li et al., 2026b) decomposes agentic coding into five expert domains, each undergoing independent SFT and RL training, and consolidates them via on-policy distillation, scoring 79.6% on SWE-bench Verified. Nemotron-Cascade 2 (Yang et al., 2026h) applies the same principle to a 30B MoE model with only 3B activated parameters, reaching Gold Medal-level performance on

IMO, IOI, and ICPC World Finals with $20\times$ fewer parameters than DeepSeek-R1. CoPD (Gu et al., 2026) identifies a key limitation of sequential expert-then-distill pipelines, that large behavioral pattern gaps between fully-trained teachers and untrained students prevent effective capability absorption. The root cause is that experts trained in isolation develop idiosyncratic reasoning strategies that are mutually incompatible when merged. CoPD’s solution is *co-evolution*, running parallel RLVR expert training with interleaved bidirectional OPD so that experts serve as mutual teachers during their own training rather than after completion. This maintains consistent behavioral patterns among experts while preserving complementary knowledge. The result is a single model integrating text, image, and video reasoning capabilities that surpasses both mixed RLVR and sequential multi-teacher OPD (MOPD). The deeper insight is architectural. Co-evolution imposes a soft constraint on teacher diversity, preventing the specialization-consolidation gap that plagues post-hoc distillation from independently trained experts.

OPD for multi-budget reasoning. ORBIT (Liang et al., 2026a) targets the deployment challenge of variable reasoning depth, where users need different compute budgets depending on query difficulty. It produces stage-wise expert policies through RLVR under successively tighter context budgets ($L_{k+1} = L_k/2$), then fuses them into a single model through mode-aware Reverse KL:

$$\mathcal{L}_{\text{fusion}}(\phi) = \mathbb{E}_{k \sim U(1,K)} \mathbb{E}_{\substack{q \sim \mathcal{D} \\ o \sim \pi_{\phi}(\cdot|q,p_k)}} \left[\sum_t \log \frac{\pi_{\phi}(o_t|q, p_k, o_{<t})}{\pi_{\theta_k}(o_t|q, p_k, o_{<t})} \right]$$

where p_k is the mode-specific prompt and $\{\pi_{\theta_k}\}$ are the frontier experts. Model merging initializes the student to mitigate the cold-start problem identified by Jang et al. (2026). Without initialization near the expert manifold, the student’s on-policy rollouts rarely reach the teachers’ high-density regions, starving the distillation objective of useful signal. Across DeepSeek-Distill-Qwen-1.5B, Qwen3-4B, and Openmath-Nemotron-7B, ORBIT produces models that smoothly interpolate between four reasoning modes (with budget tiers from 2K to 32K tokens depending on student size) with competitive performance at each budget level, matching the initial model’s accuracy at the highest budget while delivering significant token savings at lower budgets. This connects to the consolidation pattern above. Where KAT-Coder-V2 and Nemotron-Cascade 2 merge experts specialized by *domain*, ORBIT merges experts specialized by *compute budget*, a sign that on-policy distillation is emerging as a general-purpose multi-teacher fusion tool across diverse axes of model specialization.

Agentic distillation. Extending OPD from single-turn generation to multi-turn agentic tasks introduces a distinct challenge: *error compounding*. In single-turn settings, a flawed token affects only subsequent tokens within the same sequence. In multi-turn agent settings, a flawed action at step t shifts the entire environmental state, causing all subsequent teacher supervision to operate on a distribution far from what the teacher encountered during its own training. This error compounding problem motivates three complementary solutions that attack it at different granularities.

At the *trajectory level*, TCOD (Wang et al., 2026d) prevents compounding by controlling the temporal depth of teacher supervision. Its temporal curriculum progressively expands the horizon, initially supervising only short interaction prefixes (where compounding is minimal) and gradually extending to full trajectories as the student’s early-turn reliability improves. This directly counters the “Trajectory-Level KL Instability” that TCOD identifies empirically, gaining up to +18 points over vanilla multi-turn OPD on ALFWorld, WebShop, and ScienceWorld.

At the *step level*, MAD-OPD’s OPAD variant (Wang et al., 2026c) offers an orthogonal decomposition. Rather than limiting *which* steps receive supervision, it treats each agentic step as an independent distillation unit, breaking the error propagation chain by preventing gradients from flowing across step boundaries. This step-level isolation is complemented by the multi-teacher debate ensemble, which targets a *second* agentic failure mode, single-teacher unreliability at later trajectory steps. When individual teachers err on complex sequential decisions, the collective intelligence of debating teachers yields more reliable supervision than any single teacher can. Combined with task-adaptive divergence selection

(JSD for agentic stability, Reverse KL for code), OPAD lifts the agentic average by +2.4% over single-teacher OPD.

At the *skill level*, OpenClaw-RL (Wang et al., 2026i) uses hindsight-guided OPD with a Process Reward Model that delivers step-level credit assignment, training a Qwen3-4B agent across personal, terminal, GUI, and SWE environments. Skill-SD (Wang et al., 2026b) applies importance-weighted Reverse KL with dynamic teacher synchronization for skill-conditioned learning, posting +14.0% over GRPO on AppWorld and +10.9% on Sokoban.

At the *turn level*, Li et al. (2026a) revisit the classical DAgger algorithm (Ross et al., 2011) for multi-turn LLM agents and show that turn-level interpolation between student and teacher policies, where the student generates some turns and the teacher fills in others before the student trains on the full trajectory via supervised learning, effectively combines the distributional alignment of on-policy training with the dense supervision of imitation learning. The resulting method addresses the same dilemma that motivates TCOd (covariate shift from student-generated prefixes) and MAD-OPD (single-teacher unreliability at later steps) through a different mechanism, by mixing policy sources at the turn granularity rather than filtering or reweighting signals post-generation. On SWE-bench Verified with Qwen3-4B and Qwen3-8B students guided by GPT-4o, DAgger improves over pure SFT by +3.9 percentage points, the classical interactive imitation learning framework translates productively to LLM agents when the interaction unit is a conversational turn rather than a single token.

The ladder from trajectory-level control (TCOD) through step-level isolation (MAD-OPD/OPAD) to skill-level decomposition (Skill-SD) exposes a general design principle. Effective agentic distillation requires matching the granularity of the distillation unit to the granularity at which errors compound. Coarse-grained approaches (full-trajectory OPD) tend to underperform because compounding is local. Fine-grained approaches (token-level) waste compute on positions where the agent’s behavior is already correct. An emerging view suggests that an effective granularity is the *decision boundary*, the point at which the agent commits to an action whose consequences are difficult to reverse. This mirrors the role of option boundaries in hierarchical reinforcement learning, where the natural unit of credit assignment is the sub-policy rather than the individual action, and suggests that future agentic OPD methods may benefit from automated decision-boundary detection instead of hand-specified decompositions.

Unified safety infrastructure. Safactory (Chen et al., 2026d) represents a fifth deployment pattern, OPD as a component of a closed-loop safety evolution pipeline. Shanghai AI Lab’s platform integrates parallel simulation for trajectory generation, trustworthy data management for experience extraction, and an autonomous evolution module that combines asynchronous RL with on-policy distillation. The dedicated OPD stage (§5.3 in their report) distills safety-aligned behavior from specialized teacher agents into a unified student, producing models that maintain trustworthy decision-making across long-horizon agentic tasks. Safactory is notable as an early publicly documented infrastructure that embeds OPD not as a one-shot training step but as a recurring component of a continuous improvement loop for agent safety.

These five patterns reflect a broader shift in how industry treats OPD, namely as a quality refinement tool (improving an existing pipeline’s outputs), an architectural simplification tool (collapsing a multi-expert system into a single deployable model), an inference-time control mechanism (allowing a single model to serve multiple latency-accuracy operating points), an enabler for capability domains where traditional training falls short (agentic, multi-turn, multi-modal), and a safety consolidation mechanism for autonomous agent pipelines. All five patterns rely on on-policy generation rather than static dataset distillation, in line with the theoretical advantages of distributional alignment surveyed above.

Safety-tax reduction via privileged-context self-distillation. Orthogonal to the five deployment patterns above, a parallel line of safety-alignment work targets the *safety tax*, the empirical drop in general capability that reliably accompanies safety fine-tuning. Fu et al. (2026a) address this with a per-token reverse-KL distillation from a privileged-context teacher (the same model conditioned on a constitution-style safety preface) onto an un-

conditioned student that observes only the user prompt. The construction borrows the privileged-information template of OPSD (Section 5.2.1) and applies it to the safety domain, where the privileged context is a fixed constitution rather than a problem-specific oracle. Because the student trains on its own rollouts under the canonical prompt distribution while receiving dense per-token supervision from the constitution-conditioned teacher, the safety update remains local to safety-relevant tokens, leaving the general-capability surface largely untouched. The reported safety-tax reduction is consistent with the broader pattern that on-policy distillation can disentangle a target capability from its collateral effects when the privileged context isolates the capability of interest.

8.2 Emerging Domains

Multimodal OPD. Vision-language models face distinct exposure bias challenges because multimodal interleaved generation must maintain coherence across modalities, and high-quality visual reasoning data is scarce compared to text. This raises a question, whether on-policy distillation can transfer knowledge *across* modality boundaries, and if so, what structural conditions make this feasible?

VOLD (Bousselham et al., 2025) provides evidence that a text-only teacher can guide a VLM student through on-policy distillation without ever observing the visual input. The teacher supplies reasoning supervision on student-generated visual reasoning traces, using text-based reasoning structure as a modality-agnostic scaffold. The authors report that an SFT cold-start phase is important for cross-modal distributional alignment. Without it, the student’s on-policy rollouts in the visual reasoning space are too far from the teacher’s text-based distribution for distillation to converge. On MMMU-Pro, VOLD improves over the baseline model by a notable margin, consistent with the hypothesis that the reasoning structure itself (rather than perceptual grounding) carries much of the transferable knowledge.

This cross-modal principle manifests along two axes, intra-model modality self-alignment and external-teacher cross-modal transfer. CORD (Hu et al., 2026) is an early instance of intra-model self-distillation, where the stronger text mode of the *same* model serves as teacher for the weaker audio mode. CORD formalizes this with weighted Reverse KL between the audio and text modes of a single model:

$$\mathcal{L}_{\text{CORD}} = \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{\hat{y} \sim \pi_{\text{audio}}(\cdot|x)} \left[\sum_t w_t \cdot D_{\text{KL}}(\pi_{\text{audio}}(\cdot|x, \hat{y}_{<t}) \parallel \pi_{\text{text}}(\cdot|x, \hat{y}_{<t})) \right]$$

where the audio mode generates on-policy rollouts and the text mode supplies token-level targets, combined with GRPO for outcome reward maximization. CORD reports strong speech reasoning performance, indicating that cross-modal self-distillation can align weaker modality modes to stronger ones within a single model. Along the second axis, Video-OPD (Li et al., 2026f) tackles temporal video grounding, where the student must localize events within video sequences, adding a temporal dimension that static image reasoning lacks. X-OPD (Cao et al., 2026) distills a separate text-only teacher into a speech LLM student by generating on-policy rollouts and aligning them with the teacher’s evaluation through token-level feedback, effectively distilling the teacher’s capabilities into the student’s multimodal representations.

KEPO (Yang et al., 2026b) identifies a prerequisite common to these multimodal OPD methods. When initial solve rates approach zero (common in specialized visual domains such as medical VQA), sparse RLVR tends to fail due to exploration collapse before any meaningful on-policy data can be generated. By conditioning preference optimization on teacher knowledge, KEPO supplies the dense bootstrapping signal that allows subsequent on-policy training to converge. VISD (Lin et al., 2026a) extends the paradigm to video reasoning through structured self-distillation, where a video-aware judge decomposes reasoning quality into multiple diagnostic dimensions and a direction-magnitude decoupling mechanism combines RL stability with fine-grained OPD credit assignment, recording approximately $2\times$ faster convergence than RLVR baselines (see Section 5.2.1 for details). HyperEyes (Li et al., 2026d) applies OPD to multimodal search agents, where parallel tool

invocations across visual entities create a credit-assignment challenge that sparse outcome rewards alone cannot resolve. A dual-grained efficiency-aware RL framework combines a macro-level trajectory reward (TRACE, which monotonically tightens a reference tool-call budget during training) with micro-level OPD that injects dense token-level corrections from an external teacher on failed rollouts. HyperEyes-30B surpasses the strongest open-source multimodal search agent of comparable scale by 9.9% in accuracy while reducing tool-call rounds by $5.3\times$ on average, evidence that OPD and efficiency-oriented RL can be co-optimized in multimodal agentic settings. The arc from VOLD through CORD to KEPO, VISD, and HyperEyes thus maps a complete pipeline, from dense bootstrapping (KEPO) $\rightarrow$ cross-modal transfer from a text-only teacher (VOLD, X-OPD) $\rightarrow$ intra-model self-alignment (CORD) $\rightarrow$ structured self-distillation for temporal reasoning (VISD) $\rightarrow$ efficiency-aware agentic OPD for multimodal search (HyperEyes).

GUI agent distillation. While GUI-SD (Section 5.2.1) tackles self-distillation for GUI grounding via visual PI, distilling compact agents from a larger teacher for full GUI task completion poses distinct challenges. Standard GKD-style OPD struggles in this domain because student outputs are often invalid JSON structures that the teacher cannot meaningfully evaluate, and the multi-solution nature of GUI tasks (multiple valid action sequences per instruction) makes single-reference supervision lossy. LiteGUI (Wu et al., 2026b) addresses both through Guided On-policy Distillation, which provides the teacher with oracle reference trajectories from a dynamically retrieved valid action set, reducing hallucinations in teacher supervision while allowing the teacher to evaluate the student’s output in context. Building on this foundation, a Multi-solution Dual-level GRPO framework jointly aligns macro-level subtask planning with micro-level execution matching, improving exploration in long-horizon GUI agent scenarios. Using Qwen3-VL-32B as teacher and 2B–3B scale students, LiteGUI achieves strong reported performance among lightweight GUI agents on ScreenSpot-Pro, OS-World, and the newly introduced Lite-Bench, while remaining competitive with much larger models. LiteGUI represents a systematic integration of on-policy distillation into the GUI agent domain, complementing GUI-SD’s self-distillation approach with a full teacher-student transfer paradigm.

Embodied intelligence and physical reasoning. Moving beyond language-only domains, OPD is increasingly applied to settings where the output space is continuous and errors are physically irreversible. These applications form a spectrum from high-level spatial reasoning through planning to low-level motor control, each demanding increasingly tight distributional alignment.

At the reasoning end of this spectrum, HY-Embodied-0.5 (Tencent Robotics X et al., 2026) employs on-policy distillation as the final stage of a multi-stage training pipeline (RL $\rightarrow$ RFT $\rightarrow$ OPD) to compress a 32B expert VLM into a 2B embodied MoT model for spatial reasoning, outperforming similarly sized models on 16 out of 22 benchmarks. The task here is cognitive (understanding spatial relations), not physical, and OPD’s role is primarily model compression with minimal distributional mismatch.

Moving toward physical planning, OPD-AV (Afsharrad et al., 2026) applies GKD with $5\times$ compression for autonomous driving on nuScenes, distilling a Qwen3-8B SFT model into a 1.7B student. The challenge intensifies because planning errors compound through time. A suboptimal lane-change decision at t constrains all future maneuvers, making on-policy training important for exposing the student to the consequences of its own imperfect decisions. GUI-SD (Zhang et al., 2026i) tackles an analogous compounding problem in graphical interfaces, where the student generates click trajectories while a privileged version of itself (with access to bounding boxes and Gaussian soft masks) supplies token-level KL supervision. Its entropy-guided distillation selectively weights tokens based on digit positional significance and teacher confidence, concentrating optimization on the most impactful and reliable coordinate positions while filtering unreliable supervision, and reaches strong results across 6 GUI grounding benchmarks.

At the motor control end, VLA-OPD (Zhong et al., 2026b) applies Reverse KL on self-generated robot trajectories, using a strong teacher model to provide dense token-level supervision to a Vision-Language-Action student. Because the action space is continuous rather than discrete, distributional mismatch manifests as physically unsafe trajectories

(collisions, drops) instead of mere quality degradation. The student must learn precise motor control from the teacher’s dense token-level supervision on its own trajectories, reporting better sample efficiency than pure RL and better robustness than offline SFT in their evaluation.

A recurring observation across this spectrum is that on-policy generation appears to become *more* important as tasks shift from cognitive to physical domains. The diversity of valid outputs (multiple correct grasping trajectories, multiple valid temporal localizations, multiple acceptable spatial descriptions) makes off-policy data coverage difficult to achieve, and the consequence of distributional mismatch escalates from lower benchmark scores to physical failure. This gradient suggests that embodied OPD will increasingly require tighter integration with environment simulators that supply the dense rollout data needed for safe policy transfer.

Medical agents and tool-augmented clinical reasoning. Clinical decision making poses a different set of challenges as an emerging domain. The task is multi-turn by nature (history taking, test ordering, result interpretation, treatment selection), tool libraries are large and heterogeneous (medical knowledge bases, imaging interpreters, drug interaction checkers), and the terminal reward is sparse relative to the episode length. Healthcare AI Gym (Jeong, 2026) offers a gymnasium-compatible environment at this scale, covering 10 clinical domains, 3.6K+ tasks, 135 domain-specific tools, and 828K medical passages, and uses it to systematically study how OPD behaves in tool-augmented clinical reasoning. Beyond benchmarking, the study exposes an agentic-textual transfer gap. Multi-turn agentic evaluation introduces systematic overhead on knowledge-recall benchmarks (MMLU-Med drops from 83.8% under logprob evaluation to 60–66% under agentic evaluation), indicating that multi-turn tool use trades parametric precision for retrieval-augmented reasoning. The proposed TT-OPD method (discussed in Sections 5.2.1 and 7.2) reaches the best score on 10 of 18 benchmarks, with an average +3.9 percentage-point improvement over a non-RL baseline, while sustaining stable multi-turn structure throughout training. Taken together, Healthcare AI Gym suggests that high-stakes specialist domains are not merely *scaled-up* agentic environments but require distillation objectives that explicitly decouple procedural competence from parametric recall, a distinction absent in current single-domain agentic benchmarks.

Protein sequence design. ProteinOPD (Zhang et al., 2026m) extends on-policy distillation to protein language models (PLMs), where the goal is to align generated protein sequences with multiple competing property objectives (stability, activity, solubility) without destroying the pre-trained designability. Standard RL-based alignment suffers from catastrophic forgetting of basic protein structure, while SFT’s mode-covering objective constrains novelty. ProteinOPD adapts a pre-trained PLM into preference-specific teachers (one per objective) and distills their knowledge into a shared student via token-level OPD on the student’s own trajectories, aligning to a normalized geometric consensus of weighted teacher distributions with bounded optimization under gradient conflicts. While multi-teacher OPD has been explored in LLM settings (DeepSeek-V4’s domain-expert consolidation, MAD-OPD’s debate ensemble, ORBIT’s mode fusion), ProteinOPD’s contribution is formal convergence guarantees under conflicting teacher gradients, a property that multi-objective biological optimization benefits from because the teachers represent not merely diverse viewpoints on the same task but genuinely antagonistic objectives. ProteinOPD records clear gains on target preference objectives without compromising designability at an $8\times$ training speedup over RL-based competitors. OPD’s mode-seeking property transfers to biological sequence generation.

Hardware description generation. RWOPD (Zou et al., 2026) extends OPD into formal hardware verification, where the student generates SystemVerilog Assertions (SVA) from natural-language specifications and an open property-equivalence verifier scores each generation for functional correctness. The verifier reward weights the forward-KL distillation objective so that rollouts closer to formal equivalence receive stronger teacher supervision, creating a domain where the precision of the reward signal (binary formal correctness) enables clean integration of verifier feedback with dense distillation. This application illustrates a broader pattern. Domains with cheap, reliable verifiers (formal verification, compiler checking, unit testing) are well-suited to reward-weighted OPD because the reward signal

is largely deterministic and the sparse-reward problem common in RL for these domains can be mitigated by dense teacher supervision on high-reward rollouts.

Test-time on-policy distillation for speculative decoding. Test-Time Speculation (TTS) (Kumar et al., 2026) addresses a degradation phenomenon in speculative decoding, where acceptance lengths of even state-of-the-art speculators decay toward 1 (no speedup) within a few thousand output tokens because draft models are trained offline on short sequences but must match the target model on much longer outputs at inference time. TTS treats this as an online on-policy distillation problem, continuously adapting the draft model (student) to the target model (teacher) during generation. Since the token verification step already invokes the target model for each draft token, the training signal is available at no additional cost. Over successive speculation rounds, TTS tightens the draft-target alignment on the *actual* generation context. Across Qwen-3, Qwen-3.5, and Llama-3.1 families, TTS improves acceptance lengths by up to 72% (41% on average) over state-of-the-art speculators, with gains scaling with generation length. While DistillSpec (Zhou et al., 2024) pioneered using offline OPD to improve draft models for speculative decoding (Section 6.3), TTS takes this further by performing distillation *online during generation*, adapting to the actual inference-time distribution rather than a static training corpus. This represents an application of on-policy distillation at test time rather than training time, the distributional alignment benefits of OPD extend naturally to inference-time adaptation where the student-teacher gap evolves continuously.

Privacy-preserving OPD. DP-OPD (Khadem et al., 2026) injects differential privacy noise (DP-SGD) into the student’s gradients during on-policy training while keeping the teacher frozen. The study finds that on-policy querying can extract useful utility while providing formal differential-privacy guarantees on the teacher’s training data, tackling a growing concern as distillation from proprietary models becomes standard practice. Without such guarantees, on-policy querying carries a risk of memorizing and reproducing private training examples from the teacher.

Across these emerging domains, a common pattern recurs. The on-policy mechanism itself transfers unchanged, but the definition of what constitutes “privileged information,” “reward,” and “trajectory” must be renegotiated for each modality and task structure. The diversity of successful adaptations suggests that the core OPD framework is modality-agnostic, with domain-specific design choices concentrated in the signal interface rather than the optimization loop.

8.3 System-Level Integration

The compute optimization methods surveyed in Section 6.3 target algorithmic efficiency, but industrial-scale OPD also requires systems-level solutions that span the full hardware-software stack.

Production OPD pipelines must orchestrate three concurrent workloads, namely student rollout generation (autoregressive, memory-bound), teacher scoring (compute-bound forward pass), and student gradient update (compute-bound backward pass). Frameworks like OpenRLHF (Hu et al., 2024) and veRL (Sheng et al., 2024) decompose these into separate process groups with distinct parallelism strategies. Teacher inference uses tensor parallelism across 4–8 GPUs for low latency, student generation uses pipeline parallelism for throughput, and the optimizer uses ZeRO-3 for memory efficiency. The main engineering challenge is synchronization. Since the student’s policy changes after each gradient step, all pending rollouts become stale. Asynchronous architectures that tolerate slight policy staleness gain higher throughput at the cost of noisier gradients. High-throughput serving engines (vLLM (Kwon et al., 2023), TensorRT-LLM (NVIDIA, 2024)) serve the generation phase, with continuous batching and PagedAttention allowing efficient rollout generation across large prompt pools. White-box OPD also requires exposing the full vocabulary logits rather than just the sampled token. The communication cost is substantial. A 70B teacher with 7B student on 8×H100 must transmit logit data of approximately 16 GB per batch ($[B, T, |V|] \times 2$ bytes in BF16, e.g., $B=16, T=4096, |V|=128K$), feasible within a single

node via NVLink (900 GB/s) but requiring careful tensor slicing or top- k sparsification for multi-node setups.

DeepSeek-V4 (DeepSeek-AI, 2026) offers a detailed public account of full-vocabulary OPD systems engineering at trillion-parameter scale, tackling the specific challenges of distilling from 10+ teachers each potentially comprising trillions of parameters. Rather than materializing the full $|V| > 100K$ logit tensor for each teacher (which is prohibitive when distilling from 10+ teachers simultaneously), the framework caches only the last-layer hidden states in a centralized buffer. It reconstructs logits on-the-fly through the corresponding prediction head at training time, virtually eliminating memory pressure while incurring negligible recomputation cost. Training samples are ordered by teacher identity within each mini-batch, so that at most one teacher’s prediction head resides in GPU memory at any time. All weight loading and offloading proceed asynchronously without blocking the critical path. The framework further integrates FP4 quantization-aware training during the OPD phase itself, quantizing MoE expert weights and attention QK activations so that the distilled model is natively adapted to deployment-time precision from the start rather than requiring post-hoc quantization.

OPD-specific infrastructure demands. Compared to pure SFT or RL pipelines, OPD imposes three additional systems requirements that current RLHF frameworks only partially address, namely teacher co-hosting, logit-tensor transfer, and staleness tolerance. *Teacher co-hosting* requires the teacher to remain resident throughout training rather than being used for one-shot data generation, inflating GPU memory requirements by a factor that scales with teacher size. *Logit-tensor transfer* arises because white-box OPD requires transmitting full or top- k vocabulary distributions rather than scalar rewards, increasing inter-node bandwidth pressure by orders of magnitude. *Staleness tolerance* matters because the teacher’s scoring of student rollouts must remain consistent with the current student policy, which constrains the asynchrony window that reward-model-based RLHF more easily accommodates. These constraints explain why the frameworks most widely adopted for OPD (OpenRLHF, verL, SLIME, verl-pipeline) share an architectural pattern that separates rollout generation, teacher scoring, and gradient updates into distinct process groups with explicit synchronization barriers. An emerging view is that OPD at trillion-parameter scale requires framework-level support for hidden-state caching, logit reconstruction, and teacher-aware batch scheduling, rather than treating the teacher as a passive data-generating service external to the training loop.

Compute budget allocation. Beyond infrastructure architecture, the allocation of compute across training stages is itself a systems-level decision. Based on the compute analysis in Section 6.3 and the staged pipelines described in recent industrial reports (DeepSeek-AI et al., 2025; Yang et al., 2025a), a recurring pattern is observable in compute budget allocation. Several prominent industrial pipelines approximate a three-stage structure:

1. **Off-policy warm-up** (majority of budget): Standard SFT on teacher-generated data supplies cheap, high-bandwidth learning that establishes the student’s basic capabilities and aligns its distribution close enough to the teacher for effective on-policy training.
2. **On-policy logit distillation** (largest on-policy share): Token-level OPD with adaptive divergences (GKD, DistiLLM, EOPD) closes the distribution gap on the student’s own generation trajectories, mitigating exposure bias where it matters most.
3. **Reward-guided refinement** (final smaller share): RL-augmented objectives (G-OPD, KDRL, RLAD) push the student beyond the teacher’s capability ceiling on reasoning tasks, using sparse outcome rewards to discover new solution paths.

This staged approach front-loads cheap, high-bandwidth learning and reserves expensive on-policy compute for the final quality push. Qwen3 (Yang et al., 2025a) exemplifies a related pattern, with its four-stage post-training pipeline (cold-start SFT $\rightarrow$ reasoning RL $\rightarrow$ thinking-mode fusion SFT $\rightarrow$ general RL) interleaving supervised and on-policy stages, while distillation from larger models (32B, 235B-A22B) serves as the primary transfer mechanism for smaller variants, achieving comparable performance to RL at roughly 1/10

the GPU cost. On-policy training tends to be most useful after off-policy warm-up has brought the student close to the teacher’s distribution, since this reduces the wasted rollouts that occur when the student-teacher gap is too large (Section 7.1).

9 Open Problems and Future Directions

The methods surveyed above follow a recognizable trajectory from fixed objectives to adaptive ones, from external teachers to self-generated signals, and from single-turn generation to multi-turn trajectories. The open problems below are organized along this direction.

Distillation scaling laws. While parameter-scaling laws for pre-training are well-established (Kaplan et al., 2020; Hoffmann et al., 2022), the study of distillation-specific scaling has only recently begun. Busbridge et al. (2025) proposed the first distillation scaling laws under an off-policy setting, exposing a non-monotonic relationship between compute budget and optimal teacher size. As compute increases, the optimal teacher first grows until it slightly exceeds the student, then plateaus, and eventually *decreases* because inference cost with large teachers dominates the compute budget. Meanwhile, optimal student distillation tokens continue to scale as a power law. DeepSeek-R1 (DeepSeek-AI et al., 2025) offers complementary off-policy evidence on the student-size axis (Section 7.4): on AIME 2024, performance scales as 28.9% $\rightarrow$ 55.5% $\rightarrow$ 69.7% $\rightarrow$ 72.6% for student sizes 1.5B $\rightarrow$ 7B $\rightarrow$ 14B $\rightarrow$ 32B, with the steepest gain between 1.5B and 7B and sharply diminishing returns beyond 14B. Student capacity follows a power-law saturation pattern even under off-policy distillation.

Yet no equivalent framework exists for on-policy distillation, where the generation cost of student rollouts introduces an additional compute axis absent from off-policy settings. By analogy with these off-policy findings, one plausible conjectured functional form for the on-policy distillation loss is a joint power-law:

$$L(N_S, N_T, D_{\text{on}}) = E + \frac{A}{N_S^\alpha} + \frac{B}{N_T^\beta} + \frac{C}{D_{\text{on}}^\gamma} + f(N_S, N_T)$$

where E is the irreducible task entropy, the first three terms capture independent scaling of student size, teacher size, and on-policy rollout budget, and $f(N_S, N_T)$ models the capacity-gap interference between student and teacher. This form is speculative and has not been validated empirically for on-policy settings. Whether a separable power-law with an additive interaction term captures the true scaling behavior remains an open question, and the actual relationship may be substantially more complex. Nevertheless, even as a heuristic scaffold, this conjectured form helps organize the constraints that existing off-policy evidence imposes. The non-monotone teacher-size finding implies that B/N_T^β alone is insufficient and the interaction term f must capture the regime where teacher inference cost overwhelms quality gains. Qwen3 (Yang et al., 2025a) further shows that distilling from multiple specialized teachers yields better scaling than from a single monolithic teacher of equivalent total parameters, indicating that f must also account for teacher diversity. A key open question is whether the on-policy rollout term C/D_{on}^γ interacts synergistically or antagonistically with the teacher-size term, that is, whether more on-policy data can compensate for a weaker teacher, or whether teacher quality sets a hard ceiling regardless of rollout volume. Future work could conduct controlled grid searches that independently vary N_S , N_T , and D_{on} under on-policy training to disentangle these effects, addressing allocation questions such as “Given a fixed GPU budget, is it more effective to invest in a larger teacher or in more student rollouts?” Establishing such scaling laws could move OPD compute allocation closer to the predictable footing that pre-training scaling enjoys.

Uncertainty-aware feedback. Teachers currently supply point-estimate probabilities. When the teacher itself is uncertain about a token, its logits carry little useful information, yet the student treats all teacher signals equally. A promising next step is for OPD frameworks to incorporate the teacher’s *epistemic uncertainty*, allowing the student to discount feedback on queries where the teacher is guessing. Process reward models (Lightman et al., 2024) offer a potential path by scoring each reasoning step independently, allowing identification of precisely where teacher supervision is reliable versus speculative.

The connection to the failure-mode analysis of Section 7.2 is direct. The flawed prefix trap is at its core a problem of unrecognized teacher uncertainty, and explicitly modeling this uncertainty would yield a targeted solution. The idea is to suppress the distillation gradient at positions where the teacher’s output entropy approaches its maximum, effectively treating such positions as uninformative. TIP (Xu et al., 2026c) offers empirical support for this direction, showing that teacher-student divergence-based filtering outperforms entropy-only schemes, *relative* uncertainty (how surprised the teacher is by the student’s continuation) matters more than the teacher’s absolute confidence.

The next frontier is *predictive* uncertainty. Rather than measuring the teacher’s entropy on the student’s current prefix, the goal would be to estimate what the teacher’s reliability would be several tokens ahead, allowing the student to avoid committing to reasoning paths that will lead to regions of teacher incompetence. Systematically avoiding these paths before the student wastes compute exploring them could yield meaningful efficiency gains.

Agent-level, continual, and lifelong distillation. Extending to multi-turn agentic settings requires new credit assignment mechanisms, as the student must learn long-horizon tool use, environmental state tracking, and error recovery across entire interaction trajectories. OpenClaw-RL (Wang et al., 2026i), Skill-SD (Wang et al., 2026b), TCO (Wang et al., 2026d), and MAD-OPD’s OPAD variant (Wang et al., 2026c) represent early steps, but the systematic degradation of teacher supervision with trajectory depth (Li et al., 2026j) remains a challenging open problem. MAD-OPD’s step-level sampling strategy offers a partial mitigation by treating each agentic step as an independent distillation unit instead of propagating errors across the full trajectory, while TCO’s temporal curriculum controls how much of the trajectory receives teacher supervision at each training stage. Together, these methods suggest that the multi-turn distillation frontier benefits from *both* temporal decomposition (breaking trajectories into manageable units) and progressive exposure (gradually increasing horizon complexity).

First, environment non-stationarity. Agentic environments change in response to actions, making direct trajectory comparison meaningless unless the teacher delivers counterfactual evaluations conditioned on the student’s actual state. *Second, tool-use combinatorics.* Modern agent frameworks expose dozens of tools with structured argument schemas, suggesting distillation may need to operate at the tool-call level instead of individual tokens, with credit assignment mechanisms attributing trajectory success to specific tool selections. *Third, safety-critical exploration.* A single incorrect file write in coding agents or misguided click in web-browsing agents can trigger irreversible actions, requiring safety constraints that prevent exploring catastrophically dangerous action sequences.

On a separate axis, production models undergo repeated fine-tuning cycles (safety patches, capability updates), raising the question of how distillation objectives should adapt when the teacher itself evolves. SDFT (Shenfeld et al., 2026) and OEL (Ye et al., 2026a) offer initial answers, but *lifelong distillation*, attaining re-alignment without catastrophic forgetting, without retraining from scratch, and without accumulating distributional drift, connects to continual learning with the added complication that both teacher and student distributions are non-stationary.

Efficiency frontiers. The considerable compute overhead of on-policy training remains a barrier. While FOPD (Zhang et al., 2026b), Lightning-OPD (Wu et al., 2026a), and SKD (Xu et al., 2025c) have made progress, the theoretical minimum cost of on-policy distillation (the information-theoretic lower bound on how much on-policy data is needed to reach a given quality level) remains unknown. A promising underexplored direction is *selective teacher inference*. MiniPLM (Gu et al., 2025) shows that selecting training instances based on the log-probability discrepancy between teacher and a reference model (“Difference Sampling”) substantially improves data efficiency under off-policy KD. Transplanting this principle to on-policy rollouts, using teacher-reference divergence to prioritize which student-generated sequences deserve dense teacher feedback, could yield significant compute savings by skipping expensive teacher inference on uninformative rollouts where the student has already converged.

Latent-space and cross-modal distillation. Current OPD operates exclusively in vocabulary space, where cross-architecture alignment requires dedicated mechanisms (DSKD’s dual-space projections, ULD’s probability-space transport). A more radical approach would distill in a *shared latent space*, bypassing the vocabulary bottleneck entirely and allowing distillation between architecturally dissimilar models that share no structural overlap. The extension to multimodal settings raises additional open questions. How does exposure bias manifest when multiple modalities must maintain coherence? What are appropriate divergence choices when teacher and student operate in different modality spaces? Can the unified f -divergence framework be extended to handle continuous signal spaces (images, audio) alongside discrete tokens? VOLD (Bousselham et al., 2025), X-OPD (Cao et al., 2026), and CORD (Hu et al., 2026) offer empirical starting points, but the theoretical foundations for cross-modal OPD remain to be established.

Privacy and evaluation methodology. DP-OPD (Khadem et al., 2026) offers an early principled framework for differentially private distillation, with an asymmetric design applying DP-SGD only to the student while keeping the teacher frozen. The on-policy formulation is well-suited to the DP setting because training on self-generated prefixes reduces compounding errors from DP noise. Open questions include scaling DP-OPD to reasoning-capable LLMs and integrating DP with self-distillation. On the evaluation front, Fu et al. (2026b) demonstrate that token-level OPD can appear successful on standard benchmarks while failing on distribution-shifted prompts, highlighting a gap between benchmark performance and true generalization. The precision-recall framework of Cha & Cho (2025) shows distillation induces a deep trade-off between sample precision and distributional coverage, evaluation protocols therefore need to probe whether the student has learned underlying causal mechanisms rather than superficial correlations.

Diagnostic tools for failure modes. As the scale and complexity of OPD increase, the failure modes diagnosed in Section 7.2 demand more systematic diagnostic tools. While current literature identifies isolated traps, such as the flawed prefix trap and epistemic suppression (Kim et al., 2026b), the field lacks a unified framework to detect and quantify these pathologies during training. Real-time diagnostic probes that track gradient signal-to-noise ratios, representation collapse, and teacher-student divergence dynamics without incurring the full forward-pass cost would fill this gap. Drawing inspiration from gradient profiling in pre-training, such tools could automatically surface when a student enters the Ouroboros self-distillation saturation or when a multi-turn agent begins to suffer from trajectory-structure erosion (Jeong, 2026). Such probes would shift OPD debugging from post-hoc benchmark failure analysis to proactive mid-training intervention.

Cross-architecture scalability. As explored in Section 5.1.1, mapping between heterogeneous model families introduces significant friction. While DSKD (Zhang et al., 2025b) and ULD (Boizard et al., 2025) offer initial solutions via dual-space projection and optimal transport, these methods have primarily been validated on relatively small scale differences. The viability of these mappings when the capacity gap spans orders of magnitude, such as distilling a 400B MoE teacher into a 1B dense student, remains largely unproven. At such extremes, the representational bottleneck may preclude simple linear projections or vocabulary-level alignments. Developing non-linear, hierarchical alignment mechanisms that bridge massive architectural divides without prohibitive compute overhead remains an open challenge with significant practical implications.

Unifying distillation and reinforcement learning. The convergence of distillation and RL has three facets best understood together. The first is temporal scheduling. Joint objectives such as G-OPD (Yang et al., 2026e), KDRL (Xu et al., 2025a), and RLAD (Zhang et al., 2026e) establish the value of combining dense teacher guidance with sparse outcome rewards, yet how to allocate compute between imitating a teacher and exploring beyond it across the training lifecycle is still heuristic. A rigorous account of this allocation would condition on the student’s relative competence, so the appropriate schedule is likely a dynamic policy that shifts from distillation when the capability gap is large toward RL when the student matches the teacher, rather than a static hyperparameter. The second facet is the structure of the loop itself. Static distillation saturates once the student exhausts the teacher’s knowledge, while pure RL destabilizes when rewards are sparse, and KDRL (Xu et al., 2025a) and REOPOLD (Ko et al., 2026) mitigate this by interleaving dense KD gradients

for early-stage stability with reward signals for late-stage exploration. CoPD (Gu et al., 2026) pushes furthest, replacing the train-then-distill pipeline with co-evolutionary training in which RLVR and bidirectional OPD run together (Section 8.1), maintaining the behavioral proximity that lets knowledge be absorbed continuously rather than transferred post hoc. The third facet is the limit this process approaches. π -Play (Zhang et al., 2026l) offers an early existence proof of a system that improves through self-generated curricula and self-filtered data, but the saturation analysis of Section 7.2 marks the boundary, that pure self-distillation converges to the model’s own prior unless grounded by external verification or internally generated privileged information. Whether this ceiling can be raised without a stronger oracle, achieving capability *creation* rather than mere *redistribution*, is the deepest open question at the intersection of distillation, RL, and learning theory, and progress likely depends on verification mechanisms that scale without human labeling.

Connections to game-theoretic self-play. SPIN (Chen et al., 2024) trains the updated model to distinguish its previous iteration’s generations from human references, with a fixed-point guarantee at the data distribution, and IRIS (Liao et al., 2026) generalizes this via interpolative Rényi divergence, unifying several self-play variants under a single divergence-parameterized framework. These methods target convergence to a Nash equilibrium over the reference distribution, a structurally different objective from the distributional alignment central to on-policy distillation. Their game-theoretic structure, however, offers complementary strengths. The convergence guarantees that accompany Nash-equilibrium targeting could potentially stabilize the iterative self-distillation loops discussed above, while OPD’s dense token-level supervision could mitigate the sample inefficiency of adversarial training. Investigating hybrid frameworks that combine game-theoretic convergence properties with distributional distillation signals is a promising direction.

10 Conclusion

This survey has systematically examined on-policy distillation for large language models, synthesizing contributions across foundational formulations, algorithmic innovations, theoretical analyses, failure-mode diagnostics, and industrial deployments. We organize the literature along three core design axes rather than surface-level categories. *Objective function* design progresses from fixed divergences through adaptive routing to RL-augmented objectives that lift the teacher ceiling. *Signal source* architecture spans from dense white-box logits through black-box API access to self-distillation that removes the external teacher entirely. *Training dynamics* encompasses token weighting, curriculum adaptation, and compute optimization.

Key Findings. Three patterns recur across the methods surveyed. *First*, shifting from off-policy to on-policy training distributions, from $y \sim p_{\text{data}}$ to $y \sim p_{\theta}$, has been associated with consistent accuracy gains across mathematical reasoning, code generation, and instruction following (Tables 3–9 and 2). The improvement is commonly attributed to reduced exposure bias, as the student learns to recover from its own errors instead of memorizing idealized trajectories, with the performance gap widening for longer reasoning chains where error compounding is more severe (Gu et al., 2024; Agarwal et al., 2024). *Second*, divergence selection appears to benefit from being adapted to both task and token position rather than fixed globally. Reverse KL tends to excel on reasoning tasks where mode-seeking can prevent hallucination, while Forward KL preserves the diversity often needed for open-ended generation. Adaptive methods such as EOPD (Jin et al., 2026b) and AOPD (Jia et al., 2026), which select divergences based on local distributional geometry, represent the current frontier. *Third*, the boundary between distillation and reinforcement learning continues to narrow. G-OPD (Yang et al., 2026e), RLAD (Zhang et al., 2026o), and REOPOLD (Ko et al., 2026) suggest that combining dense teacher supervision with sparse outcome rewards can be effective, with KD providing gradient stability while RL permits exploration beyond the teacher’s capability. The convergence is not only conceptual but also yields a shared toolkit (trust regions, advantage estimation, policy gradient variance reduction) that both communities can draw on.

Practical Takeaways. With white-box teacher access, token-level methods (GKD (Agarwal et al., 2024), DistiLLM (Ko et al., 2024)) tend to offer a favorable quality-compute tradeoff for general instruction following. For reasoning-intensive tasks, sequence-level methods (MiniLLM (Gu et al., 2024)) or hybrid approaches (FOPD (Zhang et al., 2026b), PAGED (Xu et al., 2026b), Lightning-OPD (Wu et al., 2026a)) are often preferred despite higher gradient variance inherent to REINFORCE-style estimators (Gu et al., 2024). With only API access, GAD (Ye et al., 2025) and OVD (Xiong et al., 2026) can support on-policy learning through adversarial and verbal formulations, respectively. Self-distillation methods (OPSD (Zhao et al., 2026b), SD-ZERO (He et al., 2026)) remove the need for a separate teacher entirely, which makes them attractive under resource constraints. The substantial compute overhead of on-policy generation can often be reduced via speculative decoding (Xu et al., 2025c), prefix truncation (Zhang et al., 2026b), or offline precomputation (Wu et al., 2026a).

Looking Ahead. The shift from off-policy imitation to on-policy self-correction echoes a broader pattern in machine learning. Classical supervised learning assumes a fixed target distribution, whereas on-policy distillation treats the learner’s own distribution as part of the training objective. The methods surveyed here progressively relax the assumptions under which distillation was originally formulated as a compression technique. Fixed objectives are replaced by adaptive ones, external teachers by self-generated signals, and single-turn optimization by multi-turn trajectory learning. A plausible continuation of these trends is a training regime in which models generate, evaluate, and update their own behavior more continuously rather than only during a one-shot training run. The absence of distillation scaling laws (Section 9), the open problem of teacher uncertainty quantification, and the challenge of lifelong adaptation without catastrophic forgetting (Kirkpatrick et al., 2017) appear to be among the main barriers between the current state of the art and such a regime. As LLMs evolve from static text generators into more interactive reasoning agents, the principles of on-policy alignment surveyed here are likely to become a standard component of the LLM training stack.

Limitations. This survey focuses exclusively on methods that involve on-policy generation from the student during training. We deliberately exclude general knowledge distillation (e.g., feature-based, attention transfer), model compression techniques orthogonal to the training regime (pruning, quantization), and methods that use student-generated data only at inference time (self-consistency, majority voting). Our coverage is current through June 2026. Given the field’s rapid pace (more than ten new OPD papers per month through mid-2026), some concurrent work may be absent. The method selection considerations in Section 3.3 reflect current empirical evidence and hardware constraints. As compute costs decrease and training frameworks mature, the cost-benefit calculus will shift toward more aggressive on-policy methods.

Reproducibility challenges. A significant concern across the OPD literature is the difficulty of fair comparison. Methods are evaluated on different base models, different training compute budgets, different benchmark versions (e.g., MATH-500 vs. full MATH), and with different numbers of rollouts per prompt. Few papers control for all these variables simultaneously, making cross-method comparison unreliable from published numbers alone. Our Tables 3–9 and 2 report original-paper results rather than controlled reproductions, and readers should interpret cross-row comparisons with appropriate caution. We advocate for standardized OPD benchmarking protocols that fix the base model, compute budget, and evaluation suite, analogous to what HELM (Liang et al., 2023) offers for general LLM evaluation.

Broader impact. OPD carries immediate implications for AI accessibility, safety, and environmental sustainability.

Regarding accessibility, effective distillation allows smaller, more efficient models that run on consumer hardware, democratizing advanced reasoning capabilities that would otherwise require expensive API access to frontier models. The self-distillation methods surveyed here (OPSD, SD-ZERO) are particularly relevant, as they permit capability improvement

without any external teacher, allowing resource-constrained researchers and developers to iterate rapidly.

Regarding safety, OPD introduces both risks and mitigations. The risk is that effective distillation of frontier capabilities lowers barriers to misuse. The mitigation is that the teacher retains a supervisory role during training, offering a natural point for injecting safety constraints, alignment objectives, and capability boundaries. MSD (Qin et al., 2026) concretely instantiates safety-aware OPD, transferring safety alignment from high-resource to low-resource languages through on-policy self-distillation without requiring target-language safety data. By adaptively upweighting safety-critical tokens where teacher confidence and student disagreement are both high, MSD generalizes to unseen languages and more challenging jailbreak attacks, showing that OPD can serve as a scalable mechanism for multilingual safety alignment. At the systems level, Safactory (Chen et al., 2026d) suggests that OPD can serve as a recurring component of a continuous safety evolution loop rather than a one-shot training step, embedding distillation within a closed-loop pipeline of simulation, experience extraction, and autonomous evolution. As the field matures, we expect safety-aware OPD, where the teacher not only transfers knowledge but also enforces alignment properties during the interactive training loop, to be an important research direction.

Regarding environmental sustainability, OPD’s compute overhead ($4\text{--}5\times$ over off-policy SFT, Section 6.3) raises concerns at industrial scale. Yet the resulting smaller, more capable students can substantially reduce downstream inference costs, and the net carbon footprint over a model’s lifetime can be lower when amortized across billions of inference calls. The broader convergence of distillation, reinforcement learning, and self-play into a more unified on-policy training picture suggests that future LLMs may be trained less as one-pass artifacts and more as systems that are refined periodically through their interactions with users and environments.

References

- Amirhossein Afsharrad, Amirhesam Abedsoltan, Ahmadreza Moradipari, and Sanjay Lall. On-Policy Distillation of Language Models for Autonomous Vehicle Motion Planning. *arXiv preprint arXiv:2604.07944*, 2026. URL <https://arxiv.org/abs/2604.07944>.
- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes. *Proceedings of ICLR*, 2024. URL <https://arxiv.org/abs/2306.13649>.
- Rishabh Agrawal, Jacob Fein-Ashley, and Paria Rashidinejad. Reinforcement Learning from Rich Feedback with Distributional DAGger. *arXiv preprint arXiv:2606.05152*, 2026. URL <https://arxiv.org/abs/2606.05152>.
- Mohammadreza Armandpour, Fatih Ilhan, David Harrison, Ajay Jaiswal, Duc N. M Hoang, Fartash Faghri, Yizhe Zhang, Minsik Cho, and Mehrdad Farajtabar. Unmasking on-policy distillation: Where it helps, where it hurts, and why. *arXiv preprint arXiv:2605.10889v1*, 2026.
- Nicolas Boizard, Kevin El Haddad, Céline Hudelot, and Pierre Colombo. Towards Cross-Tokenizer Distillation: the Universal Logit Distillation Loss for LLMs. *Transactions on Machine Learning Research*, 2025. URL <https://arxiv.org/abs/2402.12030>.
- Walid Bousselham, Hilde Kuehne, and Cordelia Schmid. VOLD: Reasoning Transfer from LLMs to Vision-Language Models via On-Policy Distillation. *arXiv preprint arXiv:2510.23497*, 2025. URL <https://arxiv.org/abs/2510.23497>.
- Dan Busbridge, Amitis Shidani, Floris Weers, Jason Ramapuram, Etai Littwin, and Russ Webb. Distillation Scaling Laws. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2025. URL <https://arxiv.org/abs/2502.08606>.
- Woohyeon Byeon, Jiwon Jeon, Jeonghye Kim, and Youngchul Sung. Be My Tutor: On-Policy Co-Distillation for Mutual LLM Improvement via Peer Feedback. *arXiv preprint arXiv:2606.14368*, 2026. URL <https://arxiv.org/abs/2606.14368>.

- Yishuo Cai, Jiahui Liu, Yuanxin Liu, Haobo Deng, Linli Yao, Yuhao Zheng, Kun Ouyang, Zhimo Li, Ziyue Wang, Xu Sun, Haoli Bai, and Xiaohui Li. Thinking Without Images: Internalizing Visual Manipulation with On-Policy Self-Distillation. *arXiv preprint arXiv:2606.08719*, 2026a. URL <https://arxiv.org/abs/2606.08719>.
- Yuang Cai and Yuyu Yuan. $\mathcal{X}$ -KD: General Experiential Knowledge Distillation for Large Language Models. *arXiv preprint arXiv:2602.12674*, 2026. URL <https://arxiv.org/abs/2602.12674>.
- Yuchen Cai, Ding Cao, Liang Lin, Chunxi Luo, Xin Xu, Kai Yang, Weijie Liu, Saiyong Yang, Tianxiang Zhao, Guangzhong Sun, Guiquan Liu, and Junfeng Fang. Learning to foresee: Unveiling the unlocking efficiency of on-policy distillation. *arXiv preprint arXiv:2605.11739v1*, 2026b.
- Di Cao, Dongjie Fu, Hai Yu, Siqi Zheng, Xu Tan, and Tao Jin. X-OPD: Cross-Modal On-Policy Distillation for Capability Alignment in Speech LLMs. *arXiv preprint arXiv:2603.24596*, 2026. URL <https://arxiv.org/abs/2603.24596>.
- Sungmin Cha and Kyunghyun Cho. Why Knowledge Distillation Works in Generative Models: A Minimal Working Explanation. *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. URL <https://arxiv.org/abs/2505.13111>.
- Jingwen Chen, Wenkai Yang, Shengda Fan, Wenbo Nie, Chenxing Sun, Shaodong Zheng, Yangen Hu, Lu Pan, Ke Zeng, and Yankai Lin. Rethinking Continual Experience Internalization for Self-Evolving LLM Agents. *arXiv preprint arXiv:2606.04703*, 2026a. URL <https://arxiv.org/abs/2606.04703>.
- Tianlei Chen, Jiao Ou, Ziyuan Liu, Ruiming Tang, Jian Liang, and Han Li. Counteraction-Aware Multi-Teacher On-Policy Distillation for General Capability Recovery with Domain Preservation. *arXiv preprint arXiv:2605.27115*, 2026b. URL <https://arxiv.org/abs/2605.27115>.
- Xianwei Chen, Shimin Zhang, and Jibin Wu. f -OPD: Stabilizing Long-Horizon On-Policy Distillation with Freshness-Aware Control. *arXiv preprint arXiv:2605.17862*, 2026c. URL <https://arxiv.org/abs/2605.17862>.
- Xinquan Chen, Zhenyun Yin, Shan He, Bin Huang, Shanzhe Lei, Pengcheng Shi, Kun Cai, Bei Chen, Bangwei Liu, Zeyu Kang, Chao Huang, Yang Zhang, Wenjie Li, Ruijun Ge, Yajie Wang, Tianshun Fang, Tianyang Xu, Yiwen Cong, Meng Jin, Gaolei Li, Xuansheng Wu, Linhan Liu, Zijing He, An Li, Yan Teng, Xin Tan, Dongrui Liu, Jing Shao, ChaoChao Lu, Ji He, Jie Li, Chunfeng Song, Jinya Xu, Fan Song, Shujie Wang, Jianmin Qian, Jie Hou, Xuhong Wang, Yingchun Wang, Hui Wang, and Xia Hu. Safactory: A scalable agentic infrastructure for training trustworthy autonomous intelligence. *arXiv preprint arXiv:2605.06230v2*, 2026d. URL <https://arxiv.org/abs/2605.06230>.
- Zixiang Chen, Yihe Deng, Huizhuo Yuan, Kaixuan Ji, and Quanquan Gu. Self-Play Fine-Tuning Converts Weak Language Models to Strong Language Models. *Proceedings of ICML*, 2024. URL <https://arxiv.org/abs/2401.01335>.
- Zizhao Chen, Yuying Li, Siting Lin, and Lianxi Wang. Efficient LLM Reasoning via Variational Posterior Guidance with Efficiency Awareness. *arXiv preprint arXiv:2605.11019*, 2026e. URL <https://arxiv.org/abs/2605.11019>.
- DeepSeek-AI. DeepSeek-V4: Towards highly efficient million-token context intelligence. Technical report, DeepSeek, 2026. URL <https://huggingface.co/deepseek-ai/DeepSeek-V4-Pro/blob/main/DeepSeek-V4.pdf>.
- DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, Aixin Liu, Bing Xue, Bingxuan Wang, Bochao Wu, Bei Feng, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong

- Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Han Bao, Hanwei Xu, Haocheng Wang, Honghui Ding, Huajian Xin, Huazuo Gao, Hui Qu, Hui Li, Jianzhong Guo, Jiashi Li, Jiawei Wang, Jingchang Chen, Jingyang Yuan, Junjie Qiu, Junlong Li, J. L. Cai, Jiaqi Ni, Jian Liang, Jin Chen, Kai Dong, Kai Hu, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Liang Zhao, Litong Wang, Liyue Zhang, Lei Xu, Leyi Xia, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Meng Li, Miaojun Wang, Mingming Li, Ning Tian, Panpan Huang, Peng Zhang, Qiancheng Wang, Qinyu Chen, Qiushi Du, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, R. J. Chen, R. L. Jin, Ruyi Chen, Shanghao Lu, Shangyan Zhou, Shanhuang Chen, Shengfeng Ye, Shiyu Wang, Shuiping Yu, Shunfeng Zhou, Shuting Pan, S. S. Li, Shuang Zhou, Shaoqing Wu, Shengfeng Ye, Tao Yun, Tian Pei, Tianyu Sun, T. Wang, Wangding Zeng, Wanxia Zhao, Wen Liu, Wenfeng Liang, Wenjun Gao, Wenqin Yu, Wentao Zhang, W. L. Xiao, Wei An, Xiaodong Liu, Xiaohan Wang, Xiaokang Chen, Xiaotao Nie, Xin Cheng, Xin Liu, Xin Xie, Xingchao Liu, Xinyu Yang, Xinyuan Li, Xuecheng Su, Xuheng Lin, X. Q. Li, Xiangyue Jin, Xiaojin Shen, Xiaosha Chen, Xiaowen Sun, Xiaoxiang Wang, Xinnan Song, Xinyi Zhou, Xianzu Wang, Xinxia Shan, Y. K. Li, Y. Q. Wang, Y. X. Wei, Yang Zhang, Yanhong Xu, Yao Li, Yao Zhao, Yaofeng Sun, Yaohui Wang, Yi Yu, Yichao Zhang, Yifan Shi, Yiliang Xiong, Ying He, Yishi Piao, Yisong Wang, Yixuan Tan, Yiyang Ma, Yiyuan Liu, Yongqiang Guo, Yuan Ou, Yuduan Wang, Yue Gong, Yuheng Zou, Yujia He, Yunfan Xiong, Yuxiang Luo, Yuxiang You, Yuxuan Liu, Yuyang Zhou, Y. X. Zhu, Yanhong Xu, Yanping Huang, Yaohui Li, Yi Zheng, Yuchen Zhu, Yunxian Ma, Ying Tang, Yukun Zha, Yuting Yan, Z. Z. Ren, Zehui Ren, Zhangli Sha, Zhe Fu, Zhean Xu, Zhenda Xie, Zhengyan Zhang, Zhewen Hao, Zhicheng Ma, Zhigang Yan, Zhiyu Wu, Zihui Gu, Zijia Zhu, Zijun Liu, Zilin Li, Ziwei Xie, Ziyang Song, Zizheng Pan, Zhen Huang, Zhipeng Xu, Zhongyu Zhang, and Zhen Zhang. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. *arXiv preprint arXiv:2501.12948*, 2025. URL <https://arxiv.org/abs/2501.12948>.
- Ken Ding. HDPO: Hybrid Distillation Policy Optimization via Privileged Self-Distillation. *arXiv preprint arXiv:2603.23871*, 2026. URL <https://arxiv.org/abs/2603.23871>.
- Haoyu Dong. Self-Distillation Policy Optimization via Visual Feedback: Bridging Code and Visual Artifacts. *arXiv preprint arXiv:2606.10334*, 2026. URL <https://arxiv.org/abs/2606.10334>.
- Jiabin Fan, Guoqing Luo, Michael Bowling, and Lili Mou. KETCHUP: K-Step Return Estimation for Sequential Knowledge Distillation. *arXiv preprint arXiv:2504.19024*, 2025. URL <https://arxiv.org/abs/2504.19024>.
- Junfeng Fang, Zhepei Hong, Mao Zheng, Mingyang Song, Gengsheng Li, Houcheng Jiang, Dan Zhang, Haiyun Guo, Xiang Wang, and Tat-Seng Chua. Rubric-based On-policy Distillation. *arXiv preprint arXiv:2605.07396*, 2026.
- Yu Fu, Longxuan Yu, Haz Sameen Shahgir, Zhipeng Wei, Hui Liu, N. Benjamin Erichson, and Yue Dong. Reducing the Safety Tax in LLM Safety Alignment with On-Policy Self-Distillation. *arXiv preprint arXiv:2605.15239*, 2026a. URL <https://arxiv.org/abs/2605.15239>.
- Yuqian Fu, Haohuan Huang, Kaiwen Jiang, Jiakai Liu, Zhuo Jiang, Yuanheng Zhu, and Dongbin Zhao. Revisiting On-Policy Distillation: Empirical Failure Modes and Simple Fixes. *arXiv preprint arXiv:2603.25562*, 2026b. URL <https://arxiv.org/abs/2603.25562>.
- Gemma Team, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, Cassidy Hardin, Surya Bhupatiraju, Léonard Hussenot, Thomas Mesnard, Bobak Shahriari, Alexandre Ramé, Johan Ferret, Peter Liu, Pouya Tafti, Abe Friesen, Michelle Casbon, Sabela Ramos, Ravin Kumar, Charline Le Lan, Sammy Jerome, Anton Tsitsulin, Nino Vieillard, Piotr Stanczyk, Sertan Girgin, Nikola Momchev, Matt Hoffman, Shantanu Thakoor, Jean-Bastien Grill, Behnam Neyshabur, Olivier Bachem, Alanna Walton, Aliaksei Severyn, Alicia Parrish, Aliya Ahmad, Allen Hutchison, Alvin Abdagic, Amanda Carl, Amy Shen, Andy Brock, Andy Coenen, Anthony Laforge, Antonia Paterson, Ben Bastian, Bilal Piot, Bo Wu, Brandon Royal, Charlie Chen, Chintu Kumar, Chris Perry, Chris Welty, Christopher A. Choquette-Choo, Danila Sinopalnikov, David Weinberger, Dimple Vijaykumar,

- Dominika Rogozińska, Dustin Herbison, Elisa Bandy, Emma Wang, Eric Noland, Erica Moreira, Evan Senter, Evgenii Eltyshev, Francesco Visin, Gabriel Rasskin, Gary Wei, Glenn Cameron, Gus Martins, Hadi Hashemi, Hanna Klimczak-Plucińska, Harleen Batra, Harsh Dhand, Ivan Nardini, Jacinda Mein, Jack Zhou, James Svensson, Jeff Stanway, Jetha Chan, Jin Peng Zhou, Joana Carrasqueira, Joana Iljazi, Jocelyn Becker, Joe Fernandez, Joost van Amersfoort, Josh Gordon, Josh Lipschultz, Josh Newlan, Ju yeong Ji, Kareem Mohamed, Kartikeya Badola, Kat Black, Katie Millican, Keelin McDonell, Kelvin Nguyen, Kiranbir Sodhia, Kish Greene, Lars Lowe Sjoesund, Lauren Usui, Laurent Sifre, Lena Heuermann, Leticia Lago, Lilly McNealus, Livio Baldini Soares, Logan Kilpatrick, Lucas Dixon, Luciano Martins, Machel Reid, Manvinder Singh, Mark Iverson, Martin Görner, Mat Velloso, Mateo Wirth, Matt Davidow, Matt Miller, Matthew Rahtz, Matthew Watson, Meg Risdal, Mehran Kazemi, Michael Moynihan, Ming Zhang, Minsuk Kahng, Minwoo Park, Mofi Rahman, Mohit Khatwani, Natalie Dao, Nenshad Bardoliwalla, Nesh Devanathan, Neta Dumai, Nilay Chauhan, Oscar Wahltinez, Pankil Botarda, Parker Barnes, Paul Barham, Paul Michel, Pengchong Jin, Petko Georgiev, Phil Culliton, Pradeep Kuppala, Ramona Comanescu, Ramona Merhej, Reena Jana, Reza Ardeshtir Rokni, Rishabh Agarwal, Ryan Mullins, Samaneh Saadat, Sara Mc Carthy, Sarah Cogan, Sarah Perrin, Sébastien M. R. Arnold, Sebastian Krause, Shengyang Dai, Shruti Garg, Shruti Sheth, Sue Ronstrom, Susan Chan, Timothy Jordan, Ting Yu, Tom Eccles, Tom Hennigan, Tomas Kocisky, Tulsee Doshi, Vihan Jain, Vikas Yadav, Vilobh Meshram, Vishal Dharmadhikari, Warren Barkley, Wei Wei, Wenming Ye, Woohyun Han, Woosuk Kwon, Xiang Xu, Zhe Shen, Zhitao Gong, Zichuan Wei, Victor Cotruta, Phoebe Kirk, Anand Rao, Minh Giang, Ludovic Peran, Tris Warkentin, Eli Collins, Joelle Barral, Zoubin Ghahramani, Raia Hadsell, D. Sculley, Jeanine Banks, Anca Dragan, Slav Petrov, Oriol Vinyals, Jeff Dean, Demis Hassabis, Koray Kavukcuoglu, Clement Farabet, Elena Buchatskaya, Sebastian Borgeaud, Noah Fiedel, Armand Joulin, Kathleen Kenealy, Robert Dadashi, and Alek Andreev. Gemma 2: Improving Open Language Models at a Practical Size. *arXiv preprint arXiv:2408.00118*, 2024. URL <https://arxiv.org/abs/2408.00118>.
- Naibin Gu, Chenxu Yang, Qingyi Si, Chuanyu Qin, Dingyu Yao, Peng Fu, Zheng Lin, Weiping Wang, Nan Duan, and Jiaqi Wang. Co-Evolving Policy Distillation. *arXiv preprint arXiv:2604.27083*, 2026. URL <https://arxiv.org/abs/2604.27083>.
- Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge Distillation of Large Language Models. *Proceedings of ICLR*, 2024. URL <https://arxiv.org/abs/2306.08543>.
- Yuxian Gu, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. MiniPLM: Knowledge Distillation for Pre-Training Language Models. *Proceedings of ICLR*, 2025. URL <https://arxiv.org/abs/2410.17215>.
- Arnav Gudibande, Eric Wallace, Charlie Snell, Xinyang Geng, Hao Liu, Pieter Abbeel, Sergey Levine, and Dawn Song. The False Promise of Imitating Proprietary LLMs. *arXiv preprint arXiv:2305.15717*, 2023. URL <https://arxiv.org/abs/2305.15717>.
- Haowei Guo, Baolong Bi, Ruicheng Zhang, Bingqian Sun, and Wentao Zhang. When Should the Teacher Move? Temporal Coupling and Stability in Self On-Policy Distillation. *arXiv preprint arXiv:2606.03532*, 2026. URL <https://arxiv.org/abs/2606.03532>.
- Andy Han, Kristina Fujimoto, Avidan Shah, Kiet Nguyen, Kai Xu, Chen Yueh-Han, Ilia Sucholutsky, and Rico Angell. On-Policy Consistency Training Improves LLM Safety with Minimal Capability Degradation. *arXiv preprint arXiv:2605.21834*, 2026a. URL <https://arxiv.org/abs/2605.21834>.
- Zihao Han, Tiangang Zhang, Huaibin Wang, and Yilun Sun. Adaptive teacher exposure for self-distillation in llm reasoning. *arXiv preprint arXiv:2605.11458v1*, 2026b.
- Dongze Hao, Zhiwei Jin, Chen Chen, and Haonan Lu. Stabilizing On-Policy Distillation for MLLM Reasoning with Global Normalization. *arXiv preprint arXiv:2606.09091*, 2026. URL <https://arxiv.org/abs/2606.09091>.

- Yinghui He, Simran Kaur, Adithya Bhaskar, Yongjin Yang, Jiarui Liu, Narutatsu Ri, Liam Fowl, Abhishek Panigrahi, Danqi Chen, and Sanjeev Arora. Self-Distillation Zero: Self-Revision Turns Binary Rewards into Dense Supervision. *arXiv preprint arXiv:2604.12002*, 2026. URL <https://arxiv.org/abs/2604.12002>.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katie Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Jack W. Rae, Oriol Vinyals, and Laurent Sifre. Training Compute-Optimal Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2203.15556>.
- Wenjin Hou, Shangpin Peng, Weinong Wang, Zheng Ruan, Yue Zhang, Zhenglin Zhou, Mingqi Gao, Yifei Chen, Kaiqi Wang, Hongming Yang, Chengquan Zhang, Zhuotao Tian, Han Hu, Yi Yang, Fei Wu, and Hehe Fan. Uni-opd: Unifying on-policy distillation with a dual-perspective recipe. *arXiv preprint arXiv:2605.03677*, 2026.
- Cheng-Yu Hsieh, Chun-Liang Li, Chih-Kuan Yeh, Hootan Nakhost, Yasuhisa Fujii, Alexander Ratner, Ranjay Krishna, Chen-Yu Lee, and Tomas Pfister. Distilling Step-by-Step! Outperforming Larger Language Models with Less Training Data and Smaller Model Sizes. *Findings of ACL*, 2023. URL <https://arxiv.org/abs/2305.02301>.
- Jian Hu, Xibin Tao, Weixun Hou, Quanzhi Liu, Zhengyu Feng, Yang Xu, Qing Ye, Qi Li, Zenan Li, and Shuai Zhang. OpenRLHF: An easy-to-use, scalable and high-performance RLHF framework. *arXiv preprint arXiv:2405.11143*, 2024.
- Jing Hu, Danxiang Zhu, Xianlong Luo, Dan Zhang, Shuwei He, Yishu Lei, Haitao Zheng, Shikun Feng, Jingzhou He, Yu Sun, Hua Wu, and Haifeng Wang. CORD: Bridging the Audio-Text Reasoning Gap via Weighted On-policy Cross-modal Distillation. *arXiv preprint arXiv:2601.16547*, 2026. URL <https://arxiv.org/abs/2601.16547>.
- Haiduo Huang, Jiangcheng Song, Yadong Zhang, and Pengju Ren. SelecTKD: Selective Token-Weighted Knowledge Distillation for LLMs. *arXiv preprint arXiv:2510.24021*, 2025. URL <https://arxiv.org/abs/2510.24021>.
- Jiazhen Huang, Xiao Chen, Xiao Luo, Yong Dai, Senkang Hu, and Yuzhi Zhao. Skill-Conditioned Gated Self-Distillation for LLM Reasoning. *arXiv preprint arXiv:2605.28791*, 2026. URL <https://arxiv.org/abs/2605.28791>.
- Jonas Hübner, Frederike Lübeck, Lejs Behric, Anton Baumann, Marco Bagatella, Daniel Marta, Ido Hakimi, Idan Shenfeld, Thomas Kleine Buening, Carlos Guestrin, and Andreas Krause. Reinforcement Learning via Self-Distillation. *arXiv preprint arXiv:2601.20802*, 2026. URL <https://arxiv.org/abs/2601.20802>.
- Ijun Jang, Jewon Yeom, Juan Yeo, Hyunggu Lim, and Taesup Kim. Stable On-Policy Distillation through Adaptive Target Reformulation. In *Findings of ACL*, 2026. URL <https://arxiv.org/abs/2601.07155>.
- Minbyul Jeong. Healthcare AI GYM for medical agents. *arXiv preprint arXiv:2605.02943*, 2026.
- Nan Jia, Haojin Yang, Xing Ma, Jiesong Lian, Shuailiang Zhang, Weipeng Zhang, Ke Zeng, Xunliang Cai, and Zequn Sun. Asymmetric On-Policy Distillation: Bridging Exploitation and Imitation at the Token Level. *arXiv preprint arXiv:2605.06387*, 2026. URL <https://arxiv.org/abs/2605.06387>.
- Li Jiang, Haoran Xu, Yichuan Ding, and Amy Zhang. Trajectory-Refined Distillation. *arXiv preprint arXiv:2606.08432*, 2026a. URL <https://arxiv.org/abs/2606.08432>.

- Yuxin Jiang, Chunkit Chan, Mingyang Chen, and Wei Wang. Lion: Adversarial Distillation of Proprietary Large Language Models. *Proceedings of EMNLP*, 2023. URL <https://arxiv.org/abs/2305.12870>.
- Yuxuan Jiang and Francis Ferraro. Bridging Reasoning Trajectories in On-Policy Distillation via Near-Future Guidance. *arXiv preprint arXiv:2606.00305*, 2026. URL <https://arxiv.org/abs/2606.00305>.
- Yuxuan Jiang, Runchao Li, Shubhashis Roy Dipta, Dawei Li, and Zhao Yang. Cornerstones or stumbling blocks? deciphering the rock tokens in on-policy distillation. *arXiv preprint arXiv:2605.09253v1*, 2026b.
- Can Jin, Jiakang Li, Rui Wu, Eddy Zhang, and Dimitris N. Metaxas. Weak Critics Make Strong Learners: On-Policy Critique Distillation for Scalable Oversight. *arXiv preprint arXiv:2606.00424*, 2026a. URL <https://arxiv.org/abs/2606.00424>.
- Woogyoul Jin, Taywon Min, Yongjin Yang, Swanand Ravindra Kadhe, Yi Zhou, Dennis Wei, Nathalie Baracaldo, and Kimin Lee. Entropy-Aware On-Policy Distillation of Language Models. *arXiv preprint arXiv:2603.07079*, 2026b. URL <https://arxiv.org/abs/2603.07079>.
- Yiqiao Jin, Yiyang Wang, Lucheng Fu, Yijia Xiao, Yinyi Luo, Haoxin Liu, B. Aditya Prakash, Josiah Hester, Jindong Wang, and Srikanth Kumar. UniSD: Towards a Unified Self-Distillation Framework for Large Language Models. *arXiv preprint arXiv:2605.06597*, 2026c.
- Jaehun Jung, Hyunwoo Kim, Brandon Cui, Ximing Lu, David Acuna, Prithviraj Ammanabrolu, and Yejin Choi. DeltaPrompts: Escaping the Zero-Delta Trap in Multimodal Distillation. *arXiv preprint arXiv:2605.15532*, 2026. URL <https://arxiv.org/abs/2605.15532>.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.
- Junlong Ke, Yuxin Huang, Jingdi Lei, Zefan Cai, Xiangyu Peng, and Baobao Chang. Respecting self-uncertainty in on-policy self-distillation for efficient LLM reasoning. *arXiv preprint arXiv:2605.13255*, 2026.
- Fatemeh Khadem, Sajad Mousavi, Yi Fang, and Yuhong Liu. DP-OPD: Differentially Private On-Policy Distillation for Language Models. *arXiv preprint arXiv:2604.04461*, 2026. URL <https://arxiv.org/abs/2604.04461>.
- Gyeongman Kim, Doohyuk Jang, and Eunho Yang. PromptKD: Distilling Student-Friendly Knowledge for Generative Language Models via Prompt Tuning. In *Findings of EMNLP*, 2024. URL <https://arxiv.org/abs/2402.12842>.
- Jaehoon Kim and Dongha Lee. OPSD Compresses What RLVR Teaches: A Post-RL Compaction Stage for Reasoning Models. *arXiv preprint arXiv:2605.06188*, 2026. URL <https://arxiv.org/abs/2605.06188>.
- Jeonghye Kim, Jiwon Jeon, Dongsheng Li, and Yuqing Yang. Rebellious Student: Reversing Teacher Signals for Reasoning Exploration with Self-Distilled RLVR. *arXiv preprint arXiv:2605.10781*, 2026a.
- Jeonghye Kim, Xufang Luo, Minbeom Kim, Sangmook Lee, Dohyung Kim, Jiwon Jeon, Dongsheng Li, and Yuqing Yang. Why Does Self-Distillation (Sometimes) Degrade the Reasoning Capability of LLMs? *arXiv preprint arXiv:2603.24472*, 2026b. URL <https://arxiv.org/abs/2603.24472>.
- Minsang Kim and Seung Jun Baek. Explain in Your Own Words: Improving Reasoning via Token-Selective Dual Knowledge Distillation. *Proceedings of ICLR*, 2026. URL <https://arxiv.org/abs/2603.13260>.

- Yeongmin Kim, Donghyeok Shin, Mina Kang, Byeonghu Na, and Il-Chul Moon. Distillation of Large Language Models via Concrete Score Matching. In *International Conference on Learning Representations (ICLR)*, 2026c. URL <https://arxiv.org/abs/2509.25837>.
- Yoon Kim and Alexander M. Rush. Sequence-level knowledge distillation. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pp. 1317–1327, 2016.
- John Kirchenbauer, Abhimanyu Hans, Brian Bartoldson, Micah Goldblum, Ashwinee Panda, and Tom Goldstein. Multi-Token Prediction via Self-Distillation. *arXiv preprint arXiv:2602.06019*, 2026. URL <https://arxiv.org/abs/2602.06019>.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. DistiLLM: Towards Streamlined Distillation for Large Language Models. *Proceedings of ICML*, 2024. URL <https://arxiv.org/abs/2402.03898>.
- Jongwoo Ko, Tianyi Chen, Sungnyun Kim, Tianyu Ding, Luming Liang, Ilya Zharkov, and Se-Young Yun. DistiLLM-2: A Contrastive Approach Boosts the Distillation of LLMs. In *Proceedings of ICML*, 2025. URL <https://arxiv.org/abs/2503.07067>.
- Jongwoo Ko, Sara Abdali, Young Jin Kim, Tianyi Chen, and Pashmina Cameron. Scaling Reasoning Efficiently via Relaxed On-Policy Distillation. *arXiv preprint arXiv:2603.11137*, 2026. URL <https://arxiv.org/abs/2603.11137>.
- Avinash Kumar, Sujay Sanghavi, and Poulami Das. Test-time speculation. *arXiv preprint arXiv:2605.09329v1*, 2026. URL <https://arxiv.org/abs/2605.09329>.
- Soo Min Kwon, Ziteng Sun, Ananda Theertha Suresh, Himanshu Jain, and Sanjiv Kumar. Codistill-grpo: A co-distillation recipe for efficient group relative policy optimization. *arXiv preprint arXiv:2605.08873v1*, 2026.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles (SOSP)*, 2023.
- Changhao Li, Bohao Yang, Zexian Li, Chenyang Zhao, Jiahao Li, Hongjin Su, Qian Liu, and Tao Yu. Revisiting DAGger in the era of LLM-agents. *arXiv preprint arXiv:2605.12913*, 2026a.
- Fengxiang Li, Han Zhang, Haoyang Huang, Jinghui Wang, Jinhua Hao, Kun Yuan, Mengtong Li, Minglei Zhang, Pengcheng Xu, Wenhao Zhuang, Yizhen Shao, Zongxian Feng, Can Tang, Chao Wang, Chengxiao Tong, Fan Yang, Gang Xiong, Haixuan Gao, Han Gao, Hao Wang, Haochen Liu, Hongliang Sun, Jiabao Li, Jingwen Chang, Jun Du, Junyi Peng, Leizhen Cui, Meimei Jing, Mingqi Wu, Shangpeng Yan, Shaotong Qi, Suzhe Xu, Wenxuan Zhao, Xianda Sun, Xuan Xie, Yanbo Wang, Yao Xia, Yinghan Cui, Yingpeng Chen, Yong Wang, Yuze Shi, Zhiwei Shen, Ziyu Wang, Ming Sun, Lin Ye, and Bin Chen. KAT-Coder-V2 Technical Report. *arXiv preprint arXiv:2603.27703*, 2026b. URL <https://arxiv.org/abs/2603.27703>.
- Gengsheng Li, Tianyu Yang, Junfeng Fang, Mingyang Song, Mao Zheng, Haiyun Guo, Dan Zhang, Jinqiao Wang, and Tat-Seng Chua. Unifying Group-Relative and Self-Distillation Policy Optimization via Sample Routing. *arXiv preprint arXiv:2604.02288*, 2026c. URL <https://arxiv.org/abs/2604.02288>.
- Guankai Li, Jiabin Chen, Yi Xu, Xichen Zhang, and Yuan Lu. Hypereyes: Dual-grained efficiency-aware reinforcement learning for parallel multimodal search agents. *arXiv preprint arXiv:2605.07177v2*, 2026d. URL <https://arxiv.org/abs/2605.07177>.

- Hao Li, Jingkun An, Zijun Song, Pengyu Zhu, Rui Li, Hao Wang, Wendi Feng, Yesheng Liu, Lijun Li, Jin-Ge Yao, and Lei Sha. SafeSteer: Localized On-Policy Distillation for Efficient Safety Alignment. *arXiv preprint arXiv:2606.02530*, 2026e. URL <https://arxiv.org/abs/2606.02530>.
- Jiaze Li, Hao Yin, Haoran Xu, Boshen Xu, Wenhui Tan, Zewen He, Jianzhong Ju, Zhenbo Luo, and Jian Luan. Video-OPD: Efficient Post-Training of Multimodal Large Language Models for Temporal Video Grounding via On-Policy Distillation. *arXiv preprint arXiv:2602.02994*, 2026f. URL <https://arxiv.org/abs/2602.02994>.
- Sijia Li, Jiazuo Yu, Zhoujun Li, Buyun Zhang, Reza Azad, Salman Khan, and Fahad Shahbaz Khan. GEAR: Granularity-adaptive advantage reweighting for LLM agents via self-distillation. *arXiv preprint arXiv:2605.11853*, 2026g.
- Xin Li, Hao Jiang, Annan Wang, Yichi Zhang, and Chau Yuen. The extrapolation cliff in on-policy distillation of near-deterministic structured outputs. *arXiv preprint arXiv:2605.08737v1*, 2026h.
- Yang Li, Erik Nijkamp, Semih Yavuz, and Shafiq Rayhan Joty. Learning from language feedback via variational policy distillation. *arXiv preprint arXiv:2605.15113*, 2026i. URL <https://arxiv.org/abs/2605.15113>.
- Yaxuan Li, Yuxin Zuo, Bingxiang He, Jinqian Zhang, Chaojun Xiao, Cheng Qian, Tianyu Yu, Huan ang Gao, Wenkai Yang, Zhiyuan Liu, and Ning Ding. Rethinking On-Policy Distillation of Large Language Models: Phenomenology, Mechanism, and Recipe. *arXiv preprint arXiv:2604.13016*, 2026j. URL <https://arxiv.org/abs/2604.13016>.
- Yingru Li, Ziniu Li, and Jiakai Liu. A Note on Hybrid Online Reinforcement and Imitation Learning for LLMs: Formulations and Algorithms. *arXiv preprint arXiv:2512.23097*, 2025. URL <https://arxiv.org/abs/2512.23097>.
- Yu Li, Rui Miao, Tian Lan, and Zhengling Qi. OPPO: Bayesian Value Recursion for Token-Level Credit Assignment in LLM Reasoning. *arXiv preprint arXiv:2605.21851*, 2026k. URL <https://arxiv.org/abs/2605.21851>.
- Yuduo Li, Xiaofeng Shi, Qian Kou, Longbin Yu, and Hua Zhou. RAFT: Data Refinement and Adaptive Distillation for Domain Fine-Tuning with Alleviated Forgetting. *arXiv preprint arXiv:2606.00147*, 2026l. URL <https://arxiv.org/abs/2606.00147>.
- Yuying Li, Leqi Zheng, Yongzi Yu, Wenrui Zhou, Xuchang Zhong, Xing Hu, Jing Jin, Hangjie Yuan, and Tao Feng. Filter, Then Reweight: Rethinking Optimization Granularity in On-Policy Distillation. *arXiv preprint arXiv:2606.02684*, 2026m. URL <https://arxiv.org/abs/2606.02684>.
- Kun Liang, Clive Bai, Xin Xu, Chenming Tang, Sanwoo Lee, Weijie Liu, Saiyong Yang, and Yunfang Wu. ORBIT: On-policy Exploration-Exploitation for Controllable Multi-Budget Reasoning. *arXiv preprint arXiv:2601.08310*, 2026a. URL <https://arxiv.org/abs/2601.08310>.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Christian Cosgrove, Christopher D. Manning, Christopher Ré, Diana Acosta-Navas, Drew A. Hudson, Eric Zelikman, Esin Durmus, Faisal Ladhak, Frieda Rong, Hongyu Ren, Huaxiu Yao, Jue Wang, Keshav Santhanam, Laurel Orr, Lucia Zheng, Mert Yuksekgonul, Mirac Suzgun, Nathan Kim, Neel Guha, Niladri Chatterji, Omar Khattab, Peter Henderson, Qian Huang, Ryan Chi, Sang Michael Xie, Shibani Santurkar, Surya Ganguli, Tatsunori Hashimoto, Thomas Icard, Tianyi Zhang, Vishrav Chaudhary, William Wang, Xuechen Li, Yifan Mai, Yuhui Zhang, and Yuta Koreeda. Holistic evaluation of language models. *Transactions on Machine Learning Research*, 2023.

- Zihan Liang, Yufei Ma, Ben Chen, Zhipeng Qian, Xuxin Zhang, Huangyu Dai, and Lingtao Mao. Search-E1: Self-Distillation Drives Self-Evolution in Search-Augmented Reasoning. *arXiv preprint arXiv:2605.22511*, 2026b. URL <https://arxiv.org/abs/2605.22511>.
- Wenjie Liao, Like Wu, Liangjie Zhao, Shihui Xu, and Shigeru Fujimura. IRIS: Interpolative Rényi Iterative Self-play for Large Language Model Fine-Tuning. *arXiv preprint arXiv:2604.20933*, 2026. URL <https://arxiv.org/abs/2604.20933>.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. *Proceedings of ICLR*, 2024. URL <https://arxiv.org/abs/2305.20050>.
- Hao Lin, Kunyang Lv, Xu Jiang, Jingqi Tian, Zhongjing Du, Jiayu Ding, Qiaoman Zhang, and Hongbo Jin. VISD: Enhancing Video Reasoning via Structured Self-Distillation. *arXiv preprint arXiv:2605.06094*, 2026a. URL <https://arxiv.org/abs/2605.06094>.
- Wenye Lin and Kai Han. Surgical Post-Training: Proximal On-Policy Distillation for Reasoning with Knowledge Retention. *arXiv preprint arXiv:2603.01683*, 2026. URL <https://arxiv.org/abs/2603.01683>.
- Zizhuo Lin, Quanling Liu, Jinsheng Quan, Chao Zhang, Yifan Zhu, Xing Shi, Jingtao Xu, Zhihui Li, and Yawei Luo. Same Evidence, Different Answers: Canonical-Context On-Policy Distillation for Multi-Turn Language Models. *arXiv preprint arXiv:2605.30251*, 2026b. URL <https://arxiv.org/abs/2605.30251>.
- Haoran Liu, Yuwei Zhang, Xiyao Li, Bohan Lyu, and Jingbo Shang. HERO: Hindsight-Enhanced Reflection from Environment Observations for Agentic Self-Distillation. *arXiv preprint arXiv:2606.11559*, 2026a. URL <https://arxiv.org/abs/2606.11559>.
- Kaiyuan Liu, Ruijie Xu, Yifan Zhang, Silu Mao, Lizhen Qu, Shuming Shi, and Zhi-Hong Deng. Prefix teach, suffix fade: Local teachability collapse in strong-to-weak on-policy distillation. *arXiv preprint arXiv:2605.13643*, 2026b.
- Ruiqi Liu, Xiaolei Lv, Gengsheng Li, Ximo Zhu, Zhiheng Wang, Zhengbo Zhang, Junkai Chen, Zhiheng Li, Bo Li, Jun Gao, and Shu Wu. Visual-Advantage On-Policy Distillation for Vision-Language Models. *arXiv preprint arXiv:2605.21924*, 2026c. URL <https://arxiv.org/abs/2605.21924>.
- Xiaogeng Liu, Xinyan Wang, Yingzi Ma, Yechao Zhang, and Chaowei Xiao. When Are Teacher Tokens Reliable? Position-Weighted On-Policy Self-Distillation for Reasoning. *arXiv preprint arXiv:2605.21606*, 2026d. URL <https://arxiv.org/abs/2605.21606>.
- Xinyu Liu, Kechen Jiao, Chunyang Xiao, Runsong Zhao, Junhao Ruan, Bei Li, Jiahao Liu, Qifan Wang, Xin Chen, Jingang Wang, Tong Xiao, and Jingbo Zhu. Teacher-guided policy optimization for LLM distillation. *arXiv preprint arXiv:2605.13230*, 2026e.
- Yanjiang Liu, Jie Lou, Xinyan Guan, Yuqiu Ji, Hongyu Lin, Ben He, Xianpei Han, Le Sun, Xing Yu, and Yaojie Lu. Your Teacher Can’t Help You Here: Combating Supervision Fidelity Decay in On-Policy Distillation. *arXiv preprint arXiv:2605.30833*, 2026f. URL <https://arxiv.org/abs/2605.30833>.
- Yifeng Liu, Shiyuan Zhang, Yifan Zhang, and Quanquan Gu. Self-Distilled Policy Gradient. *arXiv preprint arXiv:2606.04036*, 2026g. URL <https://arxiv.org/abs/2606.04036>.
- Yihong Liu, Raoyuan Zhao, Michael A. Hedderich, and Hinrich Schütze. Crosslingual on-policy self-distillation for multilingual reasoning. *arXiv preprint arXiv:2605.09548v1*, 2026h.
- Youwei Liu, Jian Wang, Hanlin Wang, and Wenjie Li. COMAP: Co-Evolving World Models and Agent Policies for LLM Agents. *arXiv preprint arXiv:2606.02372*, 2026i. URL <https://arxiv.org/abs/2606.02372>.

- Kevin Lu and Thinking Machines Lab. On-policy distillation. Thinking Machines Lab: Connectionism, 2025. URL <https://www.thinkingmachines.ai/blog/on-policy-distillation>. doi: 10.64434/tml.20251026.
- Zhengxi Lu, Zhiyuan Yao, Zhuowen Han, Zi-Han Wang, Jinyang Wu, Qi Gu, Xunliang Cai, Weiming Lu, Jun Xiao, Yueting Zhuang, and Yongliang Shen. Self-distilled agentic reinforcement learning. *arXiv preprint arXiv:2605.15155*, 2026.
- Feng Luo, Yu-Neng Chuang, Guanchu Wang, Zicheng Xu, Xiaotian Han, Tianyi Zhang, and Vladimir Braverman. Demystifying OPD: Length Inflation and Stabilization Strategies for Large Language Models. *arXiv preprint arXiv:2604.08527*, 2026a. URL <https://arxiv.org/abs/2604.08527>.
- Haocheng Luo, Jiahui Liu, Ruicheng Zhang, Zhizhou Zhong, Jiaqi Huang, Zunnan Xu, Quan Shi, Jun Zhou, and Xiu Li. Learning Visual Spatial Planning from Symbolic State via Modality-Gap-Aware Self-Distillation. *arXiv preprint arXiv:2606.06076*, 2026b. URL <https://arxiv.org/abs/2606.06076>.
- Yuanjie Lyu, Chengyu Wang, Jun Huang, and Tong Xu. From Correction to Mastery: Reinforced Distillation of Large Language Model Agents. *arXiv preprint arXiv:2509.14257*, 2025. URL <https://arxiv.org/abs/2509.14257>.
- Yufei Ma, Zihan Liang, Ben Chen, Zhipeng Qian, Huangyu Dai, Lingtao Mao, Xuxin Zhang, Chenyi Lei, and Wenwu Ou. SD-Search: On-Policy Hindsight Self-Distillation for Search-Augmented Reasoning. *arXiv preprint arXiv:2605.18299*, 2026. URL <https://arxiv.org/abs/2605.18299>.
- Duy Nguyen, Hanqi Xiao, Archiki Prasad, Zaid Khan, Anirban Das, Austin Zhang, Sambit Sahu, Hyunji Lee, Elias Stengel-Eskin, and Mohit Bansal. AVSD: Adaptive-View Self-Distillation by Balancing Consensus and Teacher-Specific Privileged Signals. *arXiv preprint arXiv:2605.20643*, 2026. URL <https://arxiv.org/abs/2605.20643>.
- Yifan Niu, Han Xiao, Dongyi Liu, Zelong Wang, Dihong Gong, Yasheng Wang, and Jia Li. Breaking the Tokenizer Barrier: On-Policy Distillation across Model Families. *arXiv preprint arXiv:2606.09456*, 2026. URL <https://arxiv.org/abs/2606.09456>.
- NVIDIA. TensorRT-LLM: High-performance inference framework for large language models. <https://github.com/NVIDIA/TensorRT-LLM>, 2024.
- Minjae Oh, Sangjun Song, Gyubin Choi, Yunho Choi, and Yohan Jo. KL for a KL: On-Policy Distillation with Control Variate Baseline. *arXiv preprint arXiv:2605.07865*, 2026.
- Leyi Pan, Shuchang Tao, Yunpeng Zhai, Lingzhe Zhang, Zhaoyang Liu, Bolin Ding, Aiwei Liu, and Lijie Wen. RLCSD: Reinforcement Learning with Contrastive On-Policy Self-Distillation. *arXiv preprint arXiv:2606.11709*, 2026. URL <https://arxiv.org/abs/2606.11709>.
- Sangwoo Park, Woongyeong Yeo, Seanie Lee, Yumin Choi, Hyomin Lee, Kangsan Kim, Jinheon Baek, Seong Joon Oh, and Sung Ju Hwang. It Takes Two: Complementary Self-Distillation for Contextual Integrity in LLMs. *arXiv preprint arXiv:2605.20258*, 2026. URL <https://arxiv.org/abs/2605.20258>.
- Patomporn Payoungkhamdee, Tinnakit Udsa, Jian Gang Ngui, Sarana Nutanong, Alham Fikri Aji, and Peerat Limkonchotiwat. DuDi: Dual-Signal Distillation with Cross-Lingual Verbalizer. *arXiv preprint arXiv:2606.04694*, 2026. URL <https://arxiv.org/abs/2606.04694>.
- Emiliano Penalosa, Dheeraj Vattikonda, Nicolas Gontier, Alexandre Lacoste, Laurent Charlin, and Massimo Caccia. Privileged Information Distillation for Language Models. *arXiv preprint arXiv:2602.04942*, 2026. URL <https://arxiv.org/abs/2602.04942>.

- Jingyu Peng, Maolin Wang, Hengyi Cai, Yuchen Li, Kai Zhang, Shuaiqiang Wang, Dawei Yin, and Xiangyu Zhao. AdaSwitch: Balancing Exploration and Guidance in Knowledge Distillation via Adaptive Switching. *arXiv preprint arXiv:2510.07842*, 2025. URL <https://arxiv.org/abs/2510.07842>.
- Daniil Plyusov, Alexey Gorbatsovski, Alexey Malakhov, Nikita Balagansky, Boris Shaposhnikov, Daria Korotyshova, and Daniil Gavrilov. Trust-Region Behavior Blending for On-Policy Distillation. *arXiv preprint arXiv:2605.31159*, 2026. URL <https://arxiv.org/abs/2605.31159>.
- Ruiyang Qin, Qingzhuo Wang, Dongrui Liu, Qiang Li, Zhihua Wei, and Wen Shen. Multilingual Safety Alignment via Self-Distillation. *arXiv preprint arXiv:2605.02971*, 2026. URL <https://arxiv.org/abs/2605.02971>.
- Miguel Moura Ramos, Duarte M. Alves, and André F. T. Martins. Combining on-policy optimization and distillation for long-context reasoning in large language models. *arXiv preprint arXiv:2605.12227v1*, 2026.
- Miao Rang, Zhenni Bi, Hang Zhou, Kai Han, Xuechun Wang, An Xiao, Xinghao Chen, Yunhe Wang, and Hanting Chen. Near-Policy: Accelerating On-Policy Distillation via Asynchronous Generation and Selective Packing. *arXiv preprint arXiv:2605.05940*, 2026. URL <https://arxiv.org/abs/2605.05940>.
- MohammadHossein Rezaei, Anas Mahmoud, Zihao Wang, Utkarsh Tyagi, Advait Gosai, Razvan-Gabriel Dumitru, Aakash Sabharwal, Bing Liu, and Yunzhong He. Rubric-Guided Self-Distillation: Post-Training Without Rubric Verifiers. *arXiv preprint arXiv:2606.12507*, 2026. URL <https://arxiv.org/abs/2606.12507>.
- Stéphane Ross, Geoffrey J. Gordon, and J. Andrew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics*, pp. 627–635, 2011.
- Hejian Sang, Yuanda Xu, Zhengze Zhou, Ran He, Zhipeng Wang, and Jiachen Sun. CRISP: Compressed Reasoning via Iterative Self-Policy Distillation. *arXiv preprint arXiv:2603.05433*, 2026. URL <https://arxiv.org/abs/2603.05433>.
- Guobin Shen, Xiang Cheng, Chenxiao Zhao, Lei Huang, Jindong Li, Dongcheng Zhao, and Xing Yu. Anti-self-distillation for reasoning rl via pointwise mutual information. *arXiv preprint arXiv:2605.11609v1*, 2026a.
- Guobin Shen, Lei Huang, Xiang Cheng, Chenxiao Zhao, Jindong Li, Dongcheng Zhao, and Xing Yu. From generic correlation to input-specific credit in on-policy self distillation. *arXiv preprint arXiv:2605.11613v1*, 2026b.
- Zhennan Shen, Yanshu Li, Qingyu Yin, Chak Tou Leong, Zhilin Wang, Yanxu Chen, Rongduo Han, Sunbowen Lee, and Yi R. Fung. On the Geometry of On-Policy Distillation. *arXiv preprint arXiv:2606.07082*, 2026c. URL <https://arxiv.org/abs/2606.07082>.
- Idan Shenfeld, Mehul Damani, Jonas Hübner, and Pulkit Agrawal. Self-Distillation Enables Continual Learning. *arXiv preprint arXiv:2601.19897*, 2026. URL <https://arxiv.org/abs/2601.19897>.
- Guangming Sheng, Chi Cao, Zilingfeng Gao, et al. HybridFlow: A flexible and efficient RLHF framework. *arXiv preprint arXiv:2409.19256*, 2024.
- Aasheesh Singh, Vishal Vaddina, and Dagnachew Birru. ORPO-Distill: Mixed-Policy Preference Optimization for Cross-Architecture LLM Distillation. *arXiv preprint arXiv:2509.25100*, 2025. URL <https://arxiv.org/abs/2509.25100>.
- Yuda Song, Lili Chen, Fahim Tajwar, Remi Munos, Deepak Pathak, J. Andrew Bagnell, Aarti Singh, and Andrea Zanette. Expanding the Capabilities of Reinforcement Learning via Text Feedback. *arXiv preprint arXiv:2602.02482*, 2026. URL <https://arxiv.org/abs/2602.02482>.

- Alex Stein, Furong Huang, and Tom Goldstein. GATES: Self-Distillation under Privileged Context with Consensus Gating. *arXiv preprint arXiv:2602.20574*, 2026. URL <https://arxiv.org/abs/2602.20574>.
- Xingyu Su, Jacob Helwig, Shubham Parashar, Atharv Chagi, Lakshmi Jotsna, Degui Zhi, James Caverlee, Dileep Kalathil, and Shuiwang Ji. Data-Efficient Autoregressive-to-Diffusion Language Models via On-Policy Distillation. *arXiv preprint arXiv:2606.06712*, 2026. URL <https://arxiv.org/abs/2606.06712>.
- Jie Sun, Mao Zheng, Mingyang Song, Qiyong Zhong, Yilin Cheng, Bichuan Feng, Pengfei Liu, Junfeng Fang, and Xiang Wang. SimCT: Recovering Lost Supervision for Cross-Tokenizer On-Policy Distillation. *arXiv preprint arXiv:2605.07711*, 2026.
- Wenhui Tan, Minghao Li, Xiaoqian Ma, Siqi Fan, Xiusheng Huang, Liujie Zhang, Ruihua Song, and Weihang Chen. Pair-In, Pair-Out: Latent Multi-Token Prediction for Efficient LLMs. *arXiv preprint arXiv:2605.27255*, 2026. URL <https://arxiv.org/abs/2605.27255>.
- Zhiquan Tan and Yinrong Hong. PAINT: Partial-Solution Adaptive Interpolated Training for Self-Distilled Reasoners. *arXiv preprint arXiv:2604.26573*, 2026a. URL <https://arxiv.org/abs/2604.26573>.
- Zhiquan Tan and Yinrong Hong. Self-Supervised On-Policy Distillation for Reasoning Language Models. *arXiv preprint arXiv:2605.17497*, 2026b. URL <https://arxiv.org/abs/2605.17497>.
- Tencent Robotics X, HY Vision Team, Xumin Yu, Zuyan Liu, Ziyi Wang, He Zhang, Yongming Rao, Fangfu Liu, Yani Zhang, Ruowen Zhao, Oran Wang, Yves Liang, Haitao Lin, Minghui Wang, Yubo Dong, Kevin Cheng, Bolin Ni, Rui Huang, Han Hu, Zhengyou Zhang, Linus, and Shunyu Yao. HY-Embodied-0.5: Embodied Foundation Models for Real-World Agents. *arXiv preprint arXiv:2604.07430*, 2026. URL <https://arxiv.org/abs/2604.07430>.
- Yang Tian, Rui Wang, Xumeng Wen, Junjie Li, Shizhao Sun, Lei Song, Jiang Bian, and Bo Zhao. PBSO: Privileged Bayesian Self-Distillation for Long-Horizon Credit Assignment. *arXiv preprint arXiv:2606.09348*, 2026. URL <https://arxiv.org/abs/2606.09348>.
- Bing Wang, Shaotian Yan, Chen Shen, kaiyuan liu, Sinan Fan, Ximing Li, Rui Miao, Xiaosong Yuan, Zhanming Shen, and Jieping Ye. Backtracking When It Strays: Mitigating Dual Exposure Biases in LLM Reasoning Distillation. *arXiv preprint arXiv:2605.19433*, 2026a. URL <https://arxiv.org/abs/2605.19433>.
- Hao Wang, Guozhi Wang, Han Xiao, Yufeng Zhou, Yue Pan, Jichao Wang, Ke Xu, Yafei Wen, Xiaohu Ruan, Xiaoxin Chen, and Honggang Qi. Skill-SD: Skill-Conditioned Self-Distillation for Multi-turn LLM Agents. *arXiv preprint arXiv:2604.10674*, 2026b. URL <https://arxiv.org/abs/2604.10674>.
- Jianze Wang, Ying Liu, Jinlong Chen, Xuchun Hu, Qilong Zhang, Yu Cao, Jun Wang, Hua Yang, Yong Xie, and Qianglong Chen. MAD-OPD: Breaking the Ceiling in On-Policy Distillation via Multi-Agent Debate. *arXiv preprint arXiv:2605.01347*, 2026c. URL <https://arxiv.org/abs/2605.01347>.
- Jiaqi Wang, Wenhao Zhang, Weijie Shi, Yaliang Li, and James Cheng. TCOO: Exploring Temporal Curriculum in On-Policy Distillation for Multi-turn Autonomous Agents. *arXiv preprint arXiv:2604.24005*, 2026d. URL <https://arxiv.org/abs/2604.24005>.
- Jiaxuan Wang, Xuan Ouyang, Zhiyu Chen, Yulan Hu, Zheng Pan, Xin Li, and Lan-Zhe Guo. Trace: Distilling where it matters via token-routed self on-policy alignment. *arXiv preprint arXiv:2605.10194v1*, 2026e. URL <https://arxiv.org/abs/2605.10194>.
- Ke Wang, Yuning Wu, Haoran Liu, Chaoqun Jia, Devin Chen, and Kai Wei. Physics-Guided Policy Optimization with Self-Distillation. *arXiv preprint arXiv:2606.03620*, 2026f. URL <https://arxiv.org/abs/2606.03620>.

- Sudong Wang, Weiquan Huang, Xiaomin Yu, Zuhao Yang, Hehai Lin, Keming Wu, Chaojun Xiao, Chen Chen, Wenxuan Wang, Beier Zhu, Yunjian Zhang, and Chengwei Qin. Beyond SFT-to-RL: Pre-alignment via Black-Box On-Policy Distillation for Multimodal RL. *arXiv preprint arXiv:2604.28123*, 2026g. URL <https://arxiv.org/abs/2604.28123>.
- Wenshuo Wang. Knowledge Distillation Must Account for What It Loses. *arXiv preprint arXiv:2604.25110*, 2026.
- Xun Wang, Ruishuo Chen, Zhuoran Li, Yu Chen, and Longbo Huang. When Context Returns: Toward Robust Internalization in On-Policy Distillation. *arXiv preprint arXiv:2606.11627*, 2026h. URL <https://arxiv.org/abs/2606.11627>.
- Yinjie Wang, Xuyang Chen, Xiaolong Jin, Mengdi Wang, and Ling Yang. OpenClaw-RL: Train Any Agent Simply by Talking. *arXiv preprint arXiv:2603.10165*, 2026i. URL <https://arxiv.org/abs/2603.10165>.
- Yuanyi Wang, Su Lu, Yanggan Gu, Pengkai Wang, Yifan Yang, Zhaoyi Yan, Congkai Xie, Jianmin Wu, and Hongxia Yang. Not All Disagreement Is Learnable: Token Teachability in On-Policy Distillation. *arXiv preprint arXiv:2605.26844*, 2026j. URL <https://arxiv.org/abs/2605.26844>.
- Zhenlin Wei, Pu Jian, Yingzhuo Deng, Xiaohan Wang, Jiajun Chai, Zhexin Hu, Wei Lin, Shanbin Zhang, and Guojun Yin. AMR-SD: Asymmetric Meta-Reflective Self-Distillation for Token-Level Credit Assignment. *arXiv preprint arXiv:2605.18529*, 2026. URL <https://arxiv.org/abs/2605.18529>.
- Ming Wen, Yuxuan Liu, Kun Yang, Yunhao Feng, Zhuoer Xu, Yuhao Sun, Shiwen Cui, Xiang Zheng, Xingjun Ma, and Yu-Gang Jiang. Constitutional On-Policy Safe Distillation. *arXiv preprint arXiv:2606.03089*, 2026. URL <https://arxiv.org/abs/2606.03089>.
- Yuqiao Wen, Zichao Li, Wenyu Du, and Lili Mou. f-Divergence Minimization for Sequence-Level Knowledge Distillation. *Proceedings of ACL*, 2023. URL <https://arxiv.org/abs/2307.15190>.
- Taiqiang Wu, Chaofan Tao, Jiahao Wang, Runming Yang, Zhe Zhao, and Ngai Wong. Rethinking Kullback-Leibler Divergence in Knowledge Distillation for Large Language Models. *Proceedings of COLING*, 2025. URL <https://arxiv.org/abs/2404.02657>.
- Yecheng Wu, Song Han, and Han Cai. Lightning OPD: Efficient Post-Training for Large Reasoning Models with Offline On-Policy Distillation. *arXiv preprint arXiv:2604.13010*, 2026a. URL <https://arxiv.org/abs/2604.13010>.
- Yubin Wu, Zicheng Cai, Liping Ning, Hua Wang, Zhi Chen, Yaohua Tang, and Hao Chen. LiteGUI: Distilling Compact GUI Agents with Reinforcement Learning. *arXiv preprint arXiv:2605.07505*, 2026b.
- Shizhe Xiang, Ke An, Wenlong Yu, Yue Liu, Jian Luan, Pei Fu, and Qilong Wang. Teaching the Way, Not the Answer: Privileged Tutoring Distillation for Multimodal Policy Optimization. *arXiv preprint arXiv:2606.07000*, 2026. URL <https://arxiv.org/abs/2606.07000>.
- Xiaomi LLM-Core Team, Bangjun Xiao, Bingquan Xia, Bo Yang, Bofei Gao, Bowen Shen, Chen Zhang, Chenhong He, Chiheng Lou, Fuli Luo, Gang Wang, Gang Xie, Hailin Zhang, Hanglong Lv, Hanyu Li, Heyu Chen, Hongshen Xu, Houbin Zhang, Huaqiu Liu, Jiangshan Duo, Jianyu Wei, Jiebao Xiao, Jinhao Dong, Jun Shi, Junhao Hu, Kainan Bao, Kang Zhou, Lei Li, Liang Zhao, Linghao Zhang, Peidian Li, Qianli Chen, Shaohui Liu, Shihua Yu, Shijie Cao, Shimao Chen, Shouqiu Yu, Shuo Liu, Tianling Zhou, Weijiang Su, Weikun Wang, Wenhan Ma, Xiangwei Deng, Bohan Mao, Bowen Ye, Can Cai, Chenghua Wang, Chengxuan Zhu, Chong Ma, Chun Chen, Chunan Li, Dawei Zhu, Deshan Xiao, Dong Zhang, Duo Zhang, Fangyue Liu, Feiyu Yang, Fengyuan Shi, Guoan Wang, Hao Tian, Hao Wu, Heng Qu, Hongfei Yi, Hongxu An, Hongyi Guan, Xing Zhang, Yifan Song, Yihan Yan, Yihao Zhao, Yingchun Lai, Yizhao Gao, Yu Cheng, Yuanyuan Tian, Yudong Wang, Zhen Tang, Zhengju Tang, Zhengtao Wen, Zhichao Song, Zhixian Zheng, Zihan

- Jiang, Jian Wen, Jiarui Sun, Jiawei Li, Jinlong Xue, Jun Xia, Kai Fang, Menghang Zhu, Nuo Chen, Qian Tu, Qihao Zhang, Qiyang Wang, Rang Li, Rui Ma, Shaolei Zhang, Shengfan Wang, Shicheng Li, Shuhao Gu, Shuhuai Ren, Sirui Deng, Tao Guo, Tianyang Lu, Wei Ji Zhuang, Weikang Zhang, Weimin Xiong, Wenshan Huang, Wenyu Yang, Xin Zhang, Xing Yong, Xu Wang, Xueyang Xie, Yilin Jiang, Yixin Yang, Yongzhe He, Yu Tu, Yuanliang Dong, Yuchen Liu, Yue Ma, Yue Yu, Yuxing Xiang, Zhaojun Huang, Zhenru Lin, Zhipeng Xu, Zhiyang Chen, Zhonghua Deng, Zihan Zhang, and Zihao Yue. MiMo-V2-Flash Technical Report. *arXiv preprint arXiv:2601.02780*, 2026. URL <https://arxiv.org/abs/2601.02780>.
- Haoran Xin, Anhao Zhao, Ying Sun, Jin Li, Xiaoyu Shen, and Hui Xiong. Escaping the KL Agreement Trap in On-Policy Distillation. *arXiv preprint arXiv:2606.09471*, 2026. URL <https://arxiv.org/abs/2606.09471>.
- Xingrun Xing, Haoqing Wang, Boyan Gao, Ziheng Li, and Yehui Tang. Trust Region On-Policy Distillation. *arXiv preprint arXiv:2606.01249*, 2026. URL <https://arxiv.org/abs/2606.01249>.
- Jing Xiong, Hui Shen, Shansan Gong, Yuxin Cheng, Jianghan Shen, Chaofan Tao, Haochen Tan, Haoli Bai, Lifeng Shang, and Ngai Wong. OVD: On-policy Verbal Distillation. *arXiv preprint arXiv:2601.21968*, 2026. URL <https://arxiv.org/abs/2601.21968>.
- Haoran Xu, Hongyu Wang, Yifei Gao, Jiaze Li, Xiaofeng Zhang, and Xiaosong Yuan. SG-OPD: Sign-Gated On-Policy Distillation via Sign-Consistency Gating and Phased Teacher Sampling. *arXiv preprint arXiv:2606.09304*, 2026a. URL <https://arxiv.org/abs/2606.09304>.
- Hongling Xu, Qi Zhu, Heyuan Deng, Jinpeng Li, Lu Hou, Yasheng Wang, Lifeng Shang, Ruifeng Xu, and Fei Mi. KDRL: Post-Training Reasoning LLMs via Unified Knowledge Distillation and Reinforcement Learning. *arXiv preprint arXiv:2506.02208*, 2025a. URL <https://arxiv.org/abs/2506.02208>.
- Shicheng Xu, Liang Pang, Yunchang Zhu, Jia Gu, Zihao Wei, Jingcheng Deng, Feiyang Pan, Huawei Shen, and Xueqi Cheng. RLKD: Distilling LLMs’ Reasoning via Reinforcement Learning. *arXiv preprint arXiv:2505.16142*, 2025b. URL <https://arxiv.org/abs/2505.16142>.
- Wenda Xu, Rujun Han, Zifeng Wang, Long T. Le, Dhruv Madeka, Lei Li, William Yang Wang, Rishabh Agarwal, Chen-Yu Lee, and Tomas Pfister. Speculative Knowledge Distillation: Bridging the Teacher-Student Gap Through Interleaved Sampling. *Proceedings of ICLR*, 2025c. URL <https://arxiv.org/abs/2410.11325>.
- Xiaohan Xu, Ming Li, Chongyang Tao, Tao Shen, Reynold Cheng, Jinyang Li, Can Xu, Dacheng Tao, and Tianyi Zhou. A Survey on Knowledge Distillation of Large Language Models. *arXiv preprint arXiv:2402.13116*, 2024. URL <https://arxiv.org/abs/2402.13116>.
- Yuanda Xu, Hejian Sang, Zhengze Zhou, Ran He, and Zhipeng Wang. PACED: Distillation and On-Policy Self-Distillation at the Frontier of Student Competence. *arXiv preprint arXiv:2603.11178*, 2026b. URL <https://arxiv.org/abs/2603.11178>.
- Yuanda Xu, Hejian Sang, Zhengze Zhou, Ran He, Zhipeng Wang, and Alborz Geramifard. TIP: Token Importance in On-Policy Distillation. *arXiv preprint arXiv:2604.14084*, 2026c. URL <https://arxiv.org/abs/2604.14084>.
- Yuanda Xu, Hejian Sang, Zhengze Zhou, Ran He, Zhipeng Wang, and Alborz Geramifard. Beyond grpo and on-policy distillation: An empirical sparse-to-dense reward principle for language-model post-training. *arXiv preprint arXiv:2605.12483v1*, 2026d.
- Jianhao Yan, Yafu Li, Zican Hu, Zhi Wang, Ganqu Cui, Xiaoye Qu, Yu Cheng, and Yue Zhang. Learning to Reason under Off-Policy Guidance. *arXiv preprint arXiv:2504.14945*, 2025. URL <https://arxiv.org/abs/2504.14945>.

- Shaotian Yan, Kaiyuan Liu, Chen Shen, Bing Wang, Sinan Fan, Jun Zhang, Yue Wu, Zheng Wang, and Jieping Ye. Distribution-Aligned Sequence Distillation for Superior Long-CoT Reasoning. *arXiv preprint arXiv:2601.09088*, 2026. URL <https://arxiv.org/abs/2601.09088>.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jing Zhou, Jingren Zhou, Junyang Lin, Kai Dang, Keqin Bao, Kexin Yang, Le Yu, Lianghao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang Ren, Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yinger Zhang, Yu Wan, Yuqiong Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan Qiu. Qwen3 Technical Report. *arXiv preprint arXiv:2505.09388*, 2025a. URL <https://arxiv.org/abs/2505.09388>.
- Chenxu Yang, Chuanyu Qin, Qingyi Si, Minghui Chen, Naibin Gu, Dingyu Yao, Zheng Lin, Weiping Wang, Jiaqi Wang, and Nan Duan. Self-Distilled RLVR. *arXiv preprint arXiv:2604.03128*, 2026a. URL <https://arxiv.org/abs/2604.03128>.
- Fan Yang, Rui Meng, Trudi Di Qi, Ali Ezzati, and Yuxin Wen. KEPO: Knowledge-Enhanced Preference Optimization for Multimodal Reasoning with Applications to Medical VQA. *arXiv preprint arXiv:2602.00400*, 2026b. URL <https://arxiv.org/abs/2602.00400>.
- Han Yang, Mingyan Wu, Bailan He, Zeyu Cao, Sikuan Yan, Kevin Qinghong Lin, and Zifeng Ding. Reasoning compression with mixed-policy distillation. *arXiv preprint arXiv:2605.08776*, 2026c.
- Ling Yang, Zhaochen Yu, Tianjun Zhang, Minkai Xu, Joseph E. Gonzalez, Bin Cui, and Shuicheng Yan. SuperCorrect: Advancing Small LLM Reasoning with Thought Template Distillation and Self-Correction. *Proceedings of ICLR*, 2025b. URL <https://arxiv.org/abs/2410.09008>.
- Shenzhi Yang, Guangcheng Zhu, Bowen Song, Haobo Wang, Mingxuan Xia, Xing Zheng, Yingfan Ma, Zhongqi Chen, Weiqiang Wang, and Gang Chen. OPRD: On-Policy Representation Distillation. *arXiv preprint arXiv:2606.06021*, 2026d. URL <https://arxiv.org/abs/2606.06021>.
- Wenkai Yang, Weijie Liu, Ruobing Xie, Kai Yang, Saiyong Yang, and Yankai Lin. Learning beyond Teacher: Generalized On-Policy Distillation with Reward Extrapolation. *arXiv preprint arXiv:2602.12125*, 2026e. URL <https://arxiv.org/abs/2602.12125>.
- Yuxiao Yang, Xiaoyun Wang, and Weitong Zhang. Ogls-sd: On-policy self-distillation with outcome-guided logit steering for llm reasoning. *arXiv preprint arXiv:2605.12400v1*, 2026f.
- Zhicheng Yang, Zhijiang Guo, Yifan Song, Minrui Xu, Yongxin Wang, Yiwei Wang, Xiaodan Liang, and Jing Tang. Prune-OPD: Efficient and Reliable On-Policy Distillation for Long-Horizon Reasoning. *arXiv preprint arXiv:2605.07804*, 2026g.
- Zhuolin Yang, Zihan Liu, Yang Chen, Wenliang Dai, Boxin Wang, Sheng-Chieh Lin, Chankyu Lee, Yangyi Chen, Dongfu Jiang, Jiafan He, Renjie Pi, Grace Lam, Nayeon Lee, Alexander Bukharin, Mohammad Shoeybi, Bryan Catanzaro, and Wei Ping. Nemotron-Cascade 2: Post-Training LLMs with Cascade RL and Multi-Domain On-Policy Distillation. *arXiv preprint arXiv:2603.19220*, 2026h. URL <https://arxiv.org/abs/2603.19220>.
- Tianzhu Ye, Li Dong, Zewen Chi, Xun Wu, Shaohan Huang, and Furu Wei. Black-Box On-Policy Distillation of Large Language Models. *arXiv preprint arXiv:2511.10643*, 2025. URL <https://arxiv.org/abs/2511.10643>.
- Tianzhu Ye, Li Dong, Qingxiu Dong, Xun Wu, Shaohan Huang, and Furu Wei. Online Experiential Learning for Language Models. *arXiv preprint arXiv:2603.16856*, 2026a. URL <https://arxiv.org/abs/2603.16856>.

- Tianzhu Ye, Li Dong, Xun Wu, Shaohan Huang, and Furu Wei. On-Policy Context Distillation for Language Models. *arXiv preprint arXiv:2602.12275*, 2026b. URL <https://arxiv.org/abs/2602.12275>.
- Woongyeng Yeo, Yumin Choi, Taekyung Ki, and Sung Ju Hwang. HINT-SD: Targeted Hindsight Self-Distillation for Long-Horizon Agents. *arXiv preprint arXiv:2605.17873*, 2026. URL <https://arxiv.org/abs/2605.17873>.
- Hee Suk Yoon, Eunseop Yoon, Jaehyun Jang, SooHwan Eom, Ji Woo Hong, Mark Hasegawa-Johnson, Qi Dai, Chong Luo, and Chang D. Yoo. Decomposed On-Policy Distillation for Vision-Language Reasoning: Steering Gradients for Visual Grounding. *arXiv preprint arXiv:2606.00564*, 2026. URL <https://arxiv.org/abs/2606.00564>.
- Guo Yu, Wenlin Liu, Yulan Hu, Hao-Xuan Ma, Jun-Peng Jiang, and Han-Jia Ye. Dense Supervision, Sparse Updates: On the Sparsity and Geometry of On-Policy Distillation. *arXiv preprint arXiv:2606.13657*, 2026a. URL <https://arxiv.org/abs/2606.13657>.
- Weichen Yu, Xiaomin Li, Yizhou Zhao, Xiaozhe Liu, Ruowang Zhang, Haixin Wang, Yinyi Luo, Chen Henry Wu, Gaurav Mittal, Matt Fredrikson, and Yu Hu. Multi-rollout on-policy distillation via peer successes and failures. *arXiv preprint arXiv:2605.12652*, 2026b.
- Xin Yu, Liuchen Liao, Yiwen Zhang, Yingchen Yu, Lingzhou Xue, and Qinzhen Guo. Preference-based self-distillation: Beyond KL matching via reward regularization. *arXiv preprint arXiv:2605.05040*, 2026c.
- Qianhao Yuan, Jie Lou, Xing Yu, Hongyu Lin, Le Sun, Xianpei Han, and Yaojie Lu. Vision-OPD: Learning to See Fine Details for Multimodal LLMs via On-Policy Self-Distillation. *arXiv preprint arXiv:2605.18740*, 2026. URL <https://arxiv.org/abs/2605.18740>.
- Azure Zhang, Ning Gao, Yuqin Dai, Ruiyuan Wu, Jinpeng Wang, Rena Wei Gao, Bingdong Tan, Shuzheng Gao, Zongjie Li, and Chaozheng Wang. Unlocking Proactivity in Task-Oriented Dialogue. *arXiv preprint arXiv:2605.22240*, 2026a. URL <https://arxiv.org/abs/2605.22240>.
- Dongxu Zhang, Zhichao Yang, Sepehr Janghorbani, Jun Han, Andrew Ressler, Qian Qian, Gregory D. Lyng, Sanjit Singh Batra, and Robert E. Tillman. Fast and Effective On-policy Distillation from Reasoning Prefixes. *arXiv preprint arXiv:2602.15260*, 2026b. URL <https://arxiv.org/abs/2602.15260>.
- Hongbin Zhang, Chaozheng Wang, Kehai Chen, Youcheng Pan, Yang Xiang, Jinpeng Wang, and Min Zhang. Tailoring Teaching to Aptitude: Direction-Adaptive Self-Distillation for LLM Reasoning. *arXiv preprint arXiv:2605.22263*, 2026c. URL <https://arxiv.org/abs/2605.22263>.
- Jiaxin Zhang, Xiangyu Peng, Qinglin Chen, Qinyuan Ye, Caiming Xiong, and Chien-Sheng Wu. The Illusion of Certainty: Decoupling Capability and Calibration in On-Policy Distillation. *arXiv preprint arXiv:2604.16830*, 2026d. URL <https://arxiv.org/abs/2604.16830>.
- Ke Zhang, Yunjie Tian, DongDi Zhao, Yijiang Li, Yuanye Liu, Vishal M Patel, and Di Fu. On-policy distillation with best-of-n teacher rollout selection. *arXiv preprint arXiv:2605.09725*, 2026e.
- Miaosen Zhang, Yishan Liu, Shuxia Lin, Xu Yang, Qi Dai, Chong Luo, Weihao Jiang, Peng Hou, Anxiang Zeng, Xin Geng, and Baining Guo. Towards On-Policy SFT: Distribution Discriminant Theory and its Applications in LLM Training. *arXiv preprint arXiv:2602.12222*, 2026f. URL <https://arxiv.org/abs/2602.12222>.
- Songming Zhang, Xue Zhang, Tong Zhang, Bojie Hu, Yufeng Chen, and Jinan Xu. AlignDis: Token-Level Language Model Alignment as Adaptive Policy Distillation. In *Proceedings of ACL*, 2025a. URL <https://arxiv.org/abs/2503.02832>.

- Wenhao Zhang. Beyond Absolute Imitation: Anchored Residual Guidance for Privileged On-Policy Distillation. *arXiv preprint arXiv:2606.10385*, 2026. URL <https://arxiv.org/abs/2606.10385>.
- Xinsen Zhang, Zhenkai Ding, Tianjun Pan, Run Yang, Chun Kang, Xue Xiong, and Jingnan Gu. OPSDL: On-Policy Self-Distillation for Long-Context Language Models. *arXiv preprint arXiv:2604.17535*, 2026g. URL <https://arxiv.org/abs/2604.17535>.
- XiuYu Zhang, Yi Shan, Junfeng Fang, and Zhenkai Liang. Self-Evaluation Is Already There: Eliciting Latent Judge Calibration in Base LLMs with Minimal Data. *arXiv preprint arXiv:2606.05122*, 2026h. URL <https://arxiv.org/abs/2606.05122>.
- Xue Zhang, Songming Zhang, Yunlong Liang, Fandong Meng, Yufeng Chen, Jinan Xu, and Jie Zhou. A Dual-Space Framework for General Knowledge Distillation of Large Language Models. *arXiv preprint arXiv:2504.11426*, 2025b. URL <https://arxiv.org/abs/2504.11426>.
- Yan Zhang, Daiqing Wu, Huawen Shen, Yu Zhou, and Can Ma. Learn where to click from yourself: On-policy self-distillation for gui grounding. *arXiv preprint arXiv:2605.00642*, 2026i.
- Yanfei Zhang, Xu Lin, and Chenglin Wu. StepOPSD: Step-Aware Online Preference Distillation for Agent Reinforcement Learning. *arXiv preprint arXiv:2605.27140*, 2026j. URL <https://arxiv.org/abs/2605.27140>.
- Yaocheng Zhang, Jiajun Chai, Yuqian Fu, Songjun Tu, Xiaohan Wang, Wei Lin, Guojun Yin, Qichao Zhang, Yuanheng Zhu, and Dongbin Zhao. Are Full Rollouts Necessary for On-Policy Distillation? *arXiv preprint arXiv:2605.31490*, 2026k. URL <https://arxiv.org/abs/2605.31490>.
- Yaocheng Zhang, Yuanheng Zhu, Wenyue Chong, Songjun Tu, Qichao Zhang, Jiajun Chai, Xiaohan Wang, Wei Lin, Guojun Yin, and Dongbin Zhao. π -Play: Multi-Agent Self-Play via Privileged Self-Distillation without External Data. *arXiv preprint arXiv:2604.14054*, 2026l. URL <https://arxiv.org/abs/2604.14054>.
- Yulin Zhang, He Cao, Zihao Jiang, Chenyi Zi, Zhipeng Zhou, Zijing Liu, Yu Li, Jia Li, and Ziqi Gao. Proteinopd: Towards effective and efficient preference alignment for protein design. *arXiv preprint arXiv:2605.10189v1*, 2026m. URL <https://arxiv.org/abs/2605.10189>.
- Yuwei Zhang, Sha Li, Changlong Yu, Qin Lu, Shuowei Jin, Chengyu Dong, Haoran Liu, Ilgee Hong, Xintong Li, Zhenyu Shi, Bing Yin, and Jingbo Shang. Learning with rare success but rich feedback via reflection-enhanced self-distillation. *arXiv preprint arXiv:2605.12741*, 2026n.
- Zhaoyang Zhang, Shuli Jiang, Yantao Shen, Yuting Zhang, Dhananjay Ram, Shuo Yang, Zhuowen Tu, Wei Xia, and Stefano Soatto. Reinforcement-aware Knowledge Distillation for LLM Reasoning. *arXiv preprint arXiv:2602.22495*, 2026o. URL <https://arxiv.org/abs/2602.22495>.
- Hanyang Zhao, Haoxian Chen, Han Lin, Genta Indra Winata, David Yao, and Wenpin Tang. OPD+: Rethinking the Advantage Design for On-Policy Distillation. *arXiv preprint arXiv:2606.01039*, 2026a. URL <https://arxiv.org/abs/2606.01039>.
- Siyan Zhao, Zhihui Xie, Mengchen Liu, Jing Huang, Guan Pang, Feiyu Chen, and Aditya Grover. Self-Distilled Reasoner: On-Policy Self-Distillation for Large Language Models. *arXiv preprint arXiv:2601.18734*, 2026b. URL <https://arxiv.org/abs/2601.18734>.
- Zhengyang Zhao, Lu Ma, and Wentao Zhang. Training with harnesses: On-policy harness self-distillation for complex reasoning. *arXiv preprint arXiv:2605.08741v1*, 2026c.
- Ziqi Zhao, Xinyu Ma, Liu Yang, Yujie Feng, Daiting Shi, Jingzhou He, Xin Xin, Zhaochun Ren, and Xiao-Ming Wu. ROSD: Reflective On-Policy Self-Distillation for Language Model Reasoning across Domains. *arXiv preprint arXiv:2605.28014*, 2026d. URL <https://arxiv.org/abs/2605.28014>.

- Binbin Zheng, Xing Ma, Yiheng Liang, Jingqing Ruan, Xiaoliang Fu, Kepeng Lin, Benchang Zhu, Ke Zeng, and Xunliang Cai. SCOPE: Signal-Calibrated On-Policy Distillation Enhancement with Dual-Path Adaptive Weighting. *arXiv preprint arXiv:2604.10688*, 2026a. URL <https://arxiv.org/abs/2604.10688>.
- Haoyu Zheng, Yun Zhu, Shu Yuan, Shangming Chen, Qing Wang, Wenqiao Zhang, Jun Xiao, and Yueting Zhuang. MAIGO: Mitigating Lost-in-Conversation with History-Cleaned On-Policy Self-Distillation. *arXiv preprint arXiv:2605.27186*, 2026b. URL <https://arxiv.org/abs/2605.27186>.
- Qihuang Zhong, Liang Ding, Li Shen, Juhua Liu, Bo Du, and Dacheng Tao. Revisiting Knowledge Distillation for Autoregressive Language Models. *Proceedings of ACL*, 2024. URL <https://arxiv.org/abs/2402.11890>.
- Qiyong Zhong, Mao Zheng, Mingyang Song, Xin Lin, Jie Sun, Houcheng Jiang, Xiang Wang, and Junfeng Fang. SOD: Step-wise On-policy Distillation for Small Language Model Agents. *arXiv preprint arXiv:2605.07725*, 2026a.
- Zhide Zhong, Haodong Yan, Junfeng Li, Junjie He, Tianran Zhang, and Haoang Li. VLA-OPD: Bridging Offline SFT and Online RL for Vision-Language-Action Models via On-Policy Distillation. *arXiv preprint arXiv:2603.26666*, 2026b. URL <https://arxiv.org/abs/2603.26666>.
- Yongchao Zhou, Kaifeng Lyu, Ankit Singh Rawat, Aditya Krishna Menon, Afshin Rostamizadeh, Sanjiv Kumar, Jean-François Kagy, and Rishabh Agarwal. DistillSpec: Improving Speculative Decoding via Knowledge Distillation. *Proceedings of ICLR*, 2024. URL <https://arxiv.org/abs/2310.08461>.
- Yucheng Zhou, Wei Tao, Yiwen Guo, and Jianbing Shen. World Models Meet Language Models: On the Complementarity of Concrete and Abstract Reasoning. *arXiv preprint arXiv:2606.03603*, 2026a. URL <https://arxiv.org/abs/2606.03603>.
- Yuhang Zhou, Lizhu Zhang, Yifan Wu, Mingyi Wang, Peng Bo, Jiayi Liu, Xiangjun Fan, and Zhuokai Zhao. OmniOPD: Logit-Free On-Policy Distillation via Speculative Verification. *arXiv preprint arXiv:2606.01476*, 2026b. URL <https://arxiv.org/abs/2606.01476>.
- Siqi Zhu, Xuyan Ye, Hongyu Lu, Weiye Shi, and Ge Liu. The many faces of on-policy distillation: Pitfalls, mechanisms, and fixes. *arXiv preprint arXiv:2605.11182v1*, 2026.
- Zhou Ziheng, Jiaqi Li, Huacong Tang, Ying Nian Wu, and Demetri Terzopoulos. Less is More: Early Stopping Rollout for On-Policy Distillation. *arXiv preprint arXiv:2605.27028*, 2026. URL <https://arxiv.org/abs/2605.27028>.
- Matthieu Zimmer, Xiaotong Ji, Tu Nguyen, and Haitham Bou Ammar. Rethinking Large Language Model Distillation: A Constrained Markov Decision Process Perspective. *arXiv preprint arXiv:2509.22921*, 2025. URL <https://arxiv.org/abs/2509.22921>.
- Qingyun Zou, Yifan Zhang, Bin Lin, and Wei Chen. Reward-weighted on-policy distillation with an open property-equivalence verifier for NL-to-SVA generation. *arXiv preprint arXiv:2605.13501*, 2026.